



Open Client™ Embedded SQL™/C Programmer's Guide

**Embedded SQL/C
12.5**

DOCUMENT ID: 37695-01-1250-01

LAST REVISED: May 2001

Copyright © 1989-2001 by Sybase, Inc. All rights reserved.

This publication pertains to Sybase database management software and to any subsequent release until otherwise indicated in new editions or technical notes. Information in this document is subject to change without notice. The software described herein is furnished under a license agreement, and it may be used or copied only in accordance with the terms of that agreement.

To order additional documents, U.S. and Canadian customers should call Customer Fulfillment at (800) 685-8225, fax (617) 229-9845.

Customers in other countries with a U.S. license agreement may contact Customer Fulfillment via the above fax number. All other international customers should contact their Sybase subsidiary or local distributor. Upgrades are provided only at regularly scheduled software release dates. No part of this publication may be reproduced, transmitted, or translated in any form or by any means, electronic, mechanical, manual, optical, or otherwise, without the prior written permission of Sybase, Inc.

Sybase, the Sybase logo, ADA Workbench, Adaptable Windowing Environment, Adaptive Component Architecture, Adaptive Server, Adaptive Server Anywhere, Adaptive Server Enterprise, Adaptive Server Enterprise Monitor, Adaptive Server Enterprise Replication, Adaptive Server Everywhere, Adaptive Server IQ, Adaptive Warehouse, AnswerBase, Anywhere Studio, Application Manager, AppModeler, APT Workbench, APT-Build, APT-Edit, APT-Execute, APT-FORMS, APT-Translator, APT-Library, Backup Server, ClearConnect, Client-Library, Client Services, Data Pipeline, Data Workbench, DataArchitect, Database Analyzer, DataExpress, DataServer, DataWindow, DB-Library, dbQueue, Developers Workbench, Direct Connect Anywhere, DirectConnect, Distribution Director, E-Anywhere, E-Whatever, Embedded SQL, EMS, Enterprise Application Studio, Enterprise Client/Server, Enterprise Connect, Enterprise Data Studio, Enterprise Manager, Enterprise SQL Server Manager, Enterprise Work Architecture, Enterprise Work Designer, Enterprise Work Modeler, EWA, Financial Fusion, Financial Fusion Server, Gateway Manager, ImpactNow, InfoMaker, Information Anywhere, Information Everywhere, InformationConnect, InternetBuilder, iScript, Jaguar CTS, jConnect for JDBC, KnowledgeBase, MainframeConnect, Maintenance Express, MAP, MDI Access Server, MDI Database Gateway, media.splash, MetaWorks, MySupport, Net-Gateway, Net-Library, ObjectConnect, ObjectCycle, OmniConnect, OmniSQL Access Module, OmniSQL Toolkit, Open Client, Open ClientConnect, Open Client/Server, Open Client/Server Interfaces, Open Gateway, Open Server, Open ServerConnect, Open Solutions, Optima++, PB-Gen, PC APT Execute, PC DB-Net, PC Net Library, Power++, power.stop, PowerAMC, PowerBuilder, PowerBuilder Foundation Class Library, PowerDesigner, PowerDimensions, PowerDynamo, PowerJ, PowerScript, PowerSite, PowerSocket, Powersoft, PowerStage, PowerStudio, PowerTips, Powersoft Portfolio, Powersoft Professional, PowerWare Desktop, PowerWare Enterprise, ProcessAnalyst, Report Workbench, Report-Execute, Replication Agent, Replication Driver, Replication Server, Replication Server Manager, Replication Toolkit, Resource Manager, RW-DisplayLib, RW-Library, S-Designor, SDF, Secure SQL Server, Secure SQL Toolset, Security Guardian, SKILS, smart.partners, smart.parts, smart.script, SQL Advantage, SQL Anywhere, SQL Anywhere Studio, SQL Code Checker, SQL Debug, SQL Edit, SQL Edit/TPU, SQL Everywhere, SQL Modeler, SQL Remote, SQL Server, SQL Server Manager, SQL SMART, SQL Toolset, SQL Server/CFT, SQL Server/DBM, SQL Server SNMP SubAgent, SQL Station, SQLJ, STEP, SupportNow, Sybase Central, Sybase Client/Server Interfaces, Sybase Financial Server, Sybase Gateways, Sybase MPP, Sybase SQL Desktop, Sybase SQL Lifecycle, Sybase SQL Workgroup, Sybase User Workbench, SybaseWare, Syber Financial, SyberAssist, SyBooks, System 10, System 11, System XI (logo), SystemTools, Tabular Data Stream, Transact-SQL, Translation Toolkit, UNIBOM, Unilib, Uninull, Unisep, Unistring, URK Runtime Kit for UniCode, Viewer, Visual Components, VisualSpeller, VisualWriter, VQL, WarehouseArchitect, Warehouse Control Center, Warehouse Studio, Warehouse WORKS, Watcom, Watcom SQL, Watcom SQL Server, Web Deployment Kit, Web.PB, Web.SQL, WebSights, WebViewer, WorkGroup SQL Server, XA-Library, XA-Server and XP Server are trademarks of Sybase, Inc. 3/01

Unicode and the Unicode Logo are registered trademarks of Unicode, Inc.

All other company and product names used herein may be trademarks or registered trademarks of their respective companies.

Use, duplication, or disclosure by the government is subject to the restrictions set forth in subparagraph (c)(1)(ii) of DFARS 52.227-7013 for the DOD and as set forth in FAR 52.227-19(a)-(d) for civilian agencies.

Sybase, Inc., 6475 Christie Avenue, Emeryville, CA 94608.

Contents

| | |
|-----------------------|----|
| About This Book | ix |
|-----------------------|----|

| | | |
|------------------|---|----------|
| CHAPTER 1 | Introduction | 1 |
| | Embedded SQL overview | 1 |
| | Embedded SQL features | 2 |
| | New features and enhancements | 2 |
| | Client-Library Runtime library | 2 |
| | Localization | 3 |
| | FIPS flagger | 3 |
| | Transact-SQL support in Embedded SQL | 3 |
| | Getting started | 4 |
| | Using the examples | 4 |
| | Backward compatibility | 5 |
| | Creating and running an Embedded SQL program | 5 |
| | How the precompiler processes your applications | 6 |
| | Multiple Embedded SQL source files | 7 |
| | Precompiler compatibility | 7 |
| | Precompiler-generated files | 7 |

| | | |
|------------------|---|----------|
| CHAPTER 2 | General Information | 9 |
| | Five tasks of an Embedded SQL program | 9 |
| | Simplified Embedded SQL program | 10 |
| | General rules for Embedded SQL | 11 |
| | Statement placement | 11 |
| | Comments | 12 |
| | Identifiers | 12 |
| | Quotation marks | 12 |
| | Reserved words | 12 |
| | Variable naming conventions | 13 |
| | Scoping rules | 13 |
| | Statement batches | 13 |

| | | |
|------------------|---|-----------|
| CHAPTER 3 | Communicating with Adaptive Server | 15 |
|------------------|---|-----------|

| | |
|---|----|
| Scoping rules: SQLCA, SQLCODE, and SQLSTATE | 16 |
| Declaring SQLCA | 16 |
| Multiple SQLCAs | 16 |
| SQLCA variables | 17 |
| Accessing SQLCA variables | 17 |
| SQLCODE within SQLCA | 18 |
| Declaring SQLCODE as a standalone area | 18 |
| Using SQLSTATE | 20 |
| Obtaining SQLSTATE Codes and Error Messages | 20 |
| Summary | 20 |

| | | |
|------------------|---|-----------|
| CHAPTER 4 | Using Variables | 21 |
| | Declaring variables | 21 |
| | Using datatypes | 22 |
| | Using type definitions | 24 |
| | Using #define | 25 |
| | Declaring an array | 25 |
| | Declaring unions and structures | 27 |
| | Using host variables | 28 |
| | Host input variables | 29 |
| | Host result variables | 29 |
| | Host status variables | 30 |
| | Host output variables | 30 |
| | Using indicator variables | 31 |
| | Indicator variables and server restrictions | 31 |
| | Using host variables with indicator variables | 31 |
| | Host variable conventions | 33 |
| | Using arrays | 34 |
| | Multiple arrays | 35 |
| | Scoping rules | 35 |
| | Datatypes and Adaptive Server | 35 |
| | Converting datatypes | 36 |

| | | |
|------------------|--|-----------|
| CHAPTER 5 | Connecting to Adaptive Server | 39 |
| | Connecting to a server | 39 |
| | user | 40 |
| | password | 40 |
| | connection_name | 40 |
| | server | 40 |
| | connect example | 41 |
| | Changing the current connection | 41 |
| | Establishing multiple connections | 41 |
| | Naming a connection | 42 |

| | | |
|------------------|--|-----------|
| | Using Adaptive Server connections | 43 |
| | Disconnecting from a server | 44 |
| CHAPTER 6 | Using Transact-SQL Statements | 47 |
| | Transact-SQL statements in Embedded SQL | 47 |
| | exec sql syntax | 47 |
| | Invalid statements | 48 |
| | Transact-SQL statements that differ in Embedded SQL | 48 |
| | Selecting rows | 48 |
| | Selecting one row | 49 |
| | Selecting multiple rows via arrays | 49 |
| | Selecting multiple rows via cursors | 51 |
| | Using stored procedures | 58 |
| | Grouping statements | 61 |
| | Grouping statements by batches | 61 |
| | Grouping statements by transactions | 62 |
| CHAPTER 7 | Using Dynamic SQL | 65 |
| | Dynamic SQL overview | 65 |
| | Dynamic SQL protocol | 66 |
| | Method 1: Using execute immediate | 67 |
| | Method 1: Examples | 68 |
| | Method 2: Using prepare and execute | 69 |
| | prepare | 70 |
| | execute | 70 |
| | Method 2: Example | 71 |
| | Method 3: Using prepare and fetch with a cursor | 72 |
| | prepare | 72 |
| | declare | 72 |
| | open | 73 |
| | fetch and close | 74 |
| | Method 3: Example | 74 |
| | Method 4: Using prepare and fetch with dynamic descriptors | 75 |
| | Method 4: Dynamic descriptors | 76 |
| | Dynamic descriptor statements | 77 |
| | Method 4: Example Using SQL descriptors | 78 |
| | About SQLDAs | 80 |
| | Method 4: Example using SQLDAs | 82 |
| | Summary | 84 |
| CHAPTER 8 | Handling Errors | 85 |
| | Testing for errors | 86 |

| | |
|---|----|
| Using SQLCODE..... | 86 |
| Testing for warning conditions | 86 |
| Trapping errors with whenever..... | 87 |
| whenever testing conditions | 88 |
| whenever actions | 89 |
| Using get diagnostics | 89 |
| Writing routines to handle warnings and errors..... | 90 |
| Precompiler-detected errors..... | 91 |

| | | |
|------------------|--|-----------|
| CHAPTER 9 | Improving Performance with Persistent Binding | 93 |
| | About persistent binding..... | 94 |
| | When binding occurs..... | 95 |
| | Programs that can benefit from persistent binding..... | 96 |
| | Scope of persistent bindings | 97 |
| | Precompiler options for persistent binding | 97 |
| | About the -p option | 97 |
| | About the -b option | 98 |
| | Which option to use: -p, -b, or both | 98 |
| | Scope of the -p and -b precompiler options | 98 |
| | Overview of rules for persistent binding | 98 |
| | Statements that can use persistent binding | 99 |
| | Persistent binding in statements without a cursor..... | 99 |
| | Persistent binding in statements with a cursor..... | 100 |
| | Guidelines for using persistent binding | 105 |
| | Notes on the binding of host variables | 106 |
| | Subscripted arrays | 106 |
| | Scope of host variables | 108 |

| | | |
|-------------------|---|------------|
| CHAPTER 10 | Embedded SQL Statements: Reference Pages | 111 |
| | allocate descriptor | 112 |
| | begin declare section | 114 |
| | begin transaction..... | 115 |
| | close..... | 116 |
| | commit..... | 118 |
| | connect..... | 119 |
| | deallocate cursor | 121 |
| | deallocate descriptor | 123 |
| | deallocate prepare | 124 |
| | declare cursor (dynamic)..... | 125 |
| | declare cursor (static)..... | 126 |
| | declare cursor (stored procedure) | 128 |
| | delete (positioned cursor)..... | 129 |
| | delete (searched) | 131 |

| | |
|---------------------------------------|-----|
| describe input (SQL descriptor)..... | 133 |
| describe input (SQLDA)..... | 134 |
| describe output (SQL descriptor)..... | 135 |
| describe output (SQLDA) | 136 |
| disconnect | 137 |
| exec | 139 |
| exec sql | 142 |
| execute | 144 |
| execute immediate | 146 |
| exit | 147 |
| fetch | 147 |
| get descriptor | 150 |
| get diagnostics..... | 152 |
| include "filename" | 153 |
| include sqlca | 155 |
| include sqlda | 156 |
| initialize_application..... | 157 |
| open (dynamic cursor) | 158 |
| open (static cursor) | 159 |
| prepare | 161 |
| rollback | 163 |
| select | 164 |
| set connection | 165 |
| set descriptor | 166 |
| update..... | 167 |
| whenever | 169 |

CHAPTER 11**Open Client/Server**

| | |
|--|------------|
| Configuration File | 175 |
| Purpose of the open Client/Server configuration file | 175 |
| Accessing the configuration functionality..... | 175 |
| Default settings..... | 176 |
| Syntax for the Open Client/Server configuration file..... | 177 |
| Syntax..... | 177 |
| Sample programs | 179 |
| Embedded SQL/C sample makefile on Windows NT | 179 |
| Embedded SQL/C sample programs..... | 180 |
| Embedded SQL program version for use with the -x option .. | 180 |
| Same Embedded SQL program with the -e option | 182 |
| Summary | 185 |

APPENDIX A

| | |
|--|------------|
| Precompiler Warning and Error Messages..... | 187 |
|--|------------|

| | | |
|------------|--|------------|
| APPENDIX B | Type Definitions and Limits | 199 |
| | Implementation limits..... | 199 |
| APPENDIX C | Embedded SQL Constructs | 201 |
| | Glossary..... | 203 |
| | Index..... | 211 |

About This Book

The *Open Client Embedded SQL/C Programmer's Manual* explains how to use Embedded SQL™ and the Embedded SQL precompiler with C applications. Sybase® Embedded SQL is a superset of Transact-SQL® that lets you place Transact-SQL statements in application programs written in languages such as C and COBOL.

The information in this guide is platform-independent. For platform-specific instructions on using Embedded SQL, see the *Open Client/Server Programmer's Supplement* for your platform.

Audience

This guide is intended for application developers and others interested in Embedded SQL concepts and uses. To use this guide, you should:

- Be familiar with the information presented in the *Adaptive Server Enterprise Reference Manual*
- Have C programming experience

How to use this book

The first two chapters of this guide are introductory. If you are an experienced Embedded SQL user, you may go directly to Chapter 3, “Communicating with Adaptive Server”. The guide is organized as follows:

- Chapter 1, “Introduction” presents a brief overview of Embedded SQL and describes its advantages and capabilities.
- Chapter 2, “General Information” describes the parts of an Embedded SQL program and provides general rules for programming with Embedded SQL.
- Chapter 3, “Communicating with Adaptive Server” describes how to establish and use a communication area with SQLCA, SQLCODE, and SQLSTATE. This chapter also describes the system variables used in the communication area.
- Chapter 4, “Using Variables” explains how to declare and use host and indicator variables in Embedded SQL. This chapter also describes arrays and explains datatype conversions.

-
- Chapter 5, “Connecting to Adaptive Server” explains how to use Embedded SQL to connect an application program to Sybase Adaptive Server Enterprise™ and data servers, in general.
 - Chapter 6, “Using Transact-SQL Statements” describes how to use Transact-SQL in an Embedded SQL application program. This chapter describes how to select rows using arrays and batches, and how to group Transact-SQL statements.
 - Chapter 7, “Using Dynamic SQL” describes how to create Embedded SQL statements that your application’s users can enter interactively at run time.
 - Chapter 8, “Handling Errors” describes return codes and the Embedded SQL precompiler’s facilities for detecting and handling errors.
 - Chapter 9, “Improving Performance with Persistent Binding” describes how performance might benefit from using persistent binding and how to implement it.
 - Chapter 10, “Embedded SQL Statements: Reference Pages” provides reference pages for each Embedded SQL statement.
 - Chapter 11, “Open Client/Server Configuration File” explains how to use the external configuration file with Embedded SQL.
 - Appendix A, “Precompiler Warning and Error Messages” lists precompiler and runtime messages.
 - Appendix B, “Type Definitions and Limits” lists Embedded SQL limits and typedefs.
 - Appendix C, “Embedded SQL Constructs” lists available Embedded SQL statements.
 - The Glossary defines many of the terms used in this book.

Related documents

This guide is one of several manuals you will need to have a complete understanding of Embedded SQL. The following illustration shows the other manuals you may need to consult.

- Sybase Adaptive Server Enterprise Reference Manual
- Open Client Client-Library Reference Manual
- Open Client/Server Installation Guide
- Open Client Embedded SQL Reference Manual
- Open Client/Server Programmer’s Supplement

Other sources of information

Use the Sybase Technical Library CD and the Technical Library Product Manuals Web site to learn more about your product:

- Technical Library CD contains product manuals and is included with your software. The DynaText browser (downloadable from Product Manuals at <http://www.sybase.com/detail/1,3693,1010661,00.html>) allows you to access technical information about your product in an easy-to-use format.

Refer to the *Technical Library Installation Guide* in your documentation package for instructions on installing and starting the Technical Library.

- Technical Library Product Manuals Web site is an HTML version of the Technical Library CD that you can access using a standard Web browser. In addition to product manuals, you will find links to the Technical Documents Web site (formerly known as Tech Info Library), the Solved Cases page, and Sybase/Powersoft newsgroups.

To access the Technical Library Product Manuals Web site, go to Product Manuals at <http://www.sybase.com/support/manuals/>.

Sybase certifications on the Web

Technical documentation at the Sybase Web site is updated frequently.

❖ **For the latest information on product certifications**

- 1 Point your Web browser to Technical Documents at <http://www.sybase.com/support/techdocs/>.
- 2 Select Products from the navigation bar on the left.
- 3 Select a product name from the product list.
- 4 Select the Certification Report filter, specify a time frame, and click Go.
- 5 Click a Certification Report title to display the report.

❖ **For the latest information on EBFs and Updates**

- 1 Point your Web browser to Technical Documents at <http://www.sybase.com/support/techdocs/>.
- 2 Select EBFs/Updates. Enter user name and password information, if prompted (for existing Web accounts) or create a new account (a free service).
- 3 Specify a time frame and click Go.
- 4 Select a product.
- 5 Click an EBF/Update title to display the report.

❖ **To create a personalized view of the Sybase Web site (including support pages)**

Set up a MySybase profile. MySybase is a free service that allows you to create a personalized view of Sybase Web pages.

- 1 Point your Web browser to Technical Documents at <http://www.sybase.com/support/techdocs/>
- 2 Click MySybase and create a MySybase profile.

Conventions

This section describes font style and naming conventions used in this book.

- Bold type indicates words that you type in exactly as shown, for command names and keywords. For example, in the following sentence:

With **execute immediate**, the user can enter all or part of a Transact-SQL statement.

- In all examples and syntax statements, each clause of a statement begins on a new line. Clauses that have more than one part extend to additional lines, which are indented.
- *Italic* type indicates syntax elements that you supply. In the following example, *cursor* is a keyword, and *cursor_name* represents a user-supplied identifier:

```
exec sql declare cursor_name cursor for  
           select_statement;
```

- Embedded SQL keywords are not case sensitive. You can enter them in uppercase, lowercase, or mixed case. This guide lists Embedded SQL keywords in lowercase.
- Square brackets indicate that a word or phrase is optional. In the following example, *at connection_name* is optional:

```
exec sql [at connection_name]
```

- Ellipses (...) indicate that you can repeat the item as many times as necessary. In the following example, one or more columns and one or more host variables can be listed:

```
exec sql select column [, column] ...  
           into host_variable [, host_variable] ...;
```

- Curly braces and vertical bars indicate a choice you must make. You can choose only one item in the braces. The syntax for the *whenever* statement, for example, gives a choice of three conditions and four actions:

```
exec sql whenever {sqlerror | sqlwarning | not found}  
           {continue | goto label |
```

```
call function_name([param [, param]...]) | stop};
```

If you need help

Each Sybase installation that has purchased a support contract has one or more designated people who are authorized to contact Sybase Technical Support. If you cannot resolve a problem using the manuals or online help, please have the designated person contact Sybase Technical Support or the Sybase subsidiary in your area.

Introduction

This chapter includes the following topics to introduce Embedded SQL and the Embedded SQL precompiler.

| Topic | Page |
|--|------|
| Embedded SQL overview | 1 |
| Embedded SQL features | 2 |
| New features and enhancements | 2 |
| Transact-SQL support in Embedded SQL | 3 |
| Getting started | 4 |
| Creating and running an Embedded SQL program | 5 |
| Multiple Embedded SQL source files | 7 |

Embedded SQL overview

Embedded SQL is a superset of Transact-SQL that lets you place Transact-SQL statements in application programs written in languages such as C and COBOL.

Open Client™ Embedded SQL enables you to create programs that access and update Adaptive Server data. Embedded SQL programmers write SQL statements directly into an application program written in a conventional programming language such as C or COBOL. A preprocessing program—the Embedded SQL precompiler—processes the completed application program, resulting in a program that the host language compiler can compile. The program is linked with Open Client Client-Library before it is executed.

Embedded SQL is one of the two programming methods Sybase provides for accessing Adaptive Server. The other programming method is the call-level interface. With the call-level interface, you place Client-Library calls directly into an application program, then link with Client-Library.

You can place Embedded SQL statements anywhere in a host program and mix them with host language statements. All Embedded SQL statements must begin with the keywords `exec sql` and end with a semicolon (`;`).

You can use *host variables* in Embedded SQL statements to store data retrieved from Adaptive Server and as parameters in Embedded SQL statements, such as in the `where` clause of a `select` statement. In dynamic SQL, host variables can also contain text for Embedded SQL statements.

After you write an Embedded SQL program, run it through the precompiler, which translates the Embedded SQL statements into Client-Library function calls.

Embedded SQL features

Embedded SQL provides several advantages over a call-level interface:

- Embedded SQL is easy to use because it is simply Transact-SQL with some added features that facilitate using it in an application.
- It is an ANSI/ISO-standard programming language.
- It requires less coding to achieve the same results as a call-level approach.
- Embedded SQL is essentially identical across different host languages. Programming conventions and syntax change very little. Therefore, to write applications in different languages, you need not learn new syntax.
- The precompiler can optimize execution time by generating stored procedures for the Embedded SQL statements.

New features and enhancements

Client-Library Runtime library

System 11 and later uses Open Client Client-Library as the runtime library. All function calls in Embedded SQL generated code are documented in the *Open Client Client-Library/C Reference Manual*.

Localization

An Embedded SQL application program can make calls to CS-Library routines that specify:

- A language, **character set**, and collating sequence
- How to represent dates, times, and numeric and monetary values in character format

Note See the *Open Client/Server Programmer's Supplement* for your platform for details about localization.

FIPS flagger

The FIPS (Federal Information Processing Standards) implementation upholds the goal of SQL standardization by issuing a warning when it encounters Sybase extensions to SQL statements. FIPS uses SQL-89 as the standard.

When you set the FIPS flagger, you can still use Transact-SQL extensions such as triggers and stored procedures. FIPS flags these non-ANSI statements and issues warning messages, but the application still compiles and executes.

Transact-SQL support in Embedded SQL

Transact-SQL is the set of SQL commands described in the *Adaptive Server Enterprise Reference Manual*. With the exception of `print`, `readtext`, and `writetext`, all Transact-SQL statements, functions, and control-of-flow language are valid in Embedded SQL. You can develop an interactive prototype of your Embedded SQL application in Transact-SQL to facilitate debugging your application, then easily incorporate it into your application.

Most Adaptive Server datatypes have an equivalent in Embedded SQL. Also, you can use host language datatypes in Embedded SQL. Many datatype conversions occur automatically when a host language datatype does not exactly match an Adaptive Server datatype.

You can place host language variables in Embedded SQL statements wherever literal quotes are valid in Transact-SQL. Enclose the literal with either single (') or double (") quotation marks. For information on delimiting literals that contain quotation marks, see the *Adaptive Server Enterprise Reference Manual*.

Embedded SQL has several features that Transact-SQL does not have:

- *Automatic datatype conversion* occurs between host language types and Adaptive Server types.
- *Dynamic SQL* lets you define SQL statements at run time.
- *SQLCA*, *SQLCODE*, and *SQLSTATE* let you communicate between Adaptive Server and the application program. The three entities contain error, warning, and informational message codes that Adaptive Server generates.
- *Return code testing routines* detect error conditions during execution.

Getting started

Before attempting to run the precompiler, make sure Client-Library version 11.1 or later is installed, as the precompiler uses it as the runtime library. Also, make sure **SQL Server** version 11.1 or later is installed. If one or more products is missing, contact your **System Administrator**.

Invoke the precompiler by issuing the appropriate command at the operating system prompt. See the *Open Client/Server Programmer's Supplement* for your platform for details.

The precompiler command can include several flags that let you determine options for the precompiler, including the input file, login user name and password, and precompiler modes. The *Open Client/Server Programmer's Supplement* contains operating system-specific information on precompiling, compiling, and linking your Embedded SQL application.

Using the examples

The examples in this guide use the pubs2 database. To run the examples, specify the pubs2 database with the Transact-SQL use statement.

Embedded SQL is shipped with several online examples. For information on running these examples, see the *Open Client/Server Programmer's Supplement* for your platform.

Backward compatibility

The System 11 precompiler is compatible with precompilers that are SQL-89-compliant. However, you may have applications created with earlier Embedded SQL releases that are not ANSI compliant. This precompiler uses most of the same Embedded SQL statements used in previous precompiler versions, but it processes them differently.

Follow this procedure to migrate applications created for earlier precompiler releases:

- 1 Remove the following SQL statements and keywords from the application, as System 11 does not support them:

- `release connection_name`
- `recompile`
- `noparse`
- `noproc`
- `pcoptions`
- `cancel`

`release` causes a precompiler error; the precompiler ignores the other keywords. The `cancel` statement causes a runtime error.

- 2 Use the System 11 precompiler to precompile the application again.

Creating and running an Embedded SQL program

Follow these steps to create and run an Embedded SQL application program:

- 1 Write the application program and include the Embedded SQL statements and variable declarations.
- 2 Save the application in a file with a `.cp` extension.

- 3 Precompile the application. If there are no severe errors, the precompiler generates a file containing your application program. The file has the same name as the original source file, with a different extension, depending on the requirements of your C compiler. For details, see the *Open Client/Server Programmer's Supplement* for your platform.
- 4 Compile the new source code as you would compile a standard C program.
- 5 Link the compiled code with Client-Library.
- 6 If you specified the precompiler option to generate stored procedures, load them into Adaptive Server by executing the generated script with isql.
- 7 Run the application program as you would any standard C program.

How the precompiler processes your applications

The Embedded SQL precompiler translates Embedded SQL statements into C data declarations and call statements. After precompiling, you can compile the resulting source program as you would any conventional C program.

The precompiler processes an application in two passes. In the first pass, the precompiler *parses* the Embedded SQL statements and variable declarations, checking the syntax and displaying messages for any errors it detects. If the precompiler detects no severe errors, it proceeds with the second pass, wherein it:

- Adds declarations for the precompiler variables, which begin with “_sql”. To prevent confusion, do not begin your variables' names with “_sql”.
- Converts the text of the original Embedded SQL statements to comments.
- Generates stored procedures and calls to stored procedures if you set this option in the precompile command.
- Converts Embedded SQL statements to Client-Library calls. Embedded SQL uses Client-Library as a runtime library.
- Generates up to three files: a **target file**, an optional **listing file**, and an optional **isql script file**.

Note For detailed descriptions of precompiler command line options, see the *Open Client/Server Programmer's Supplement* for your platform.

Multiple Embedded SQL source files

If the Embedded SQL application consists of more than one source file, the following statements apply:

- Connection names are unique and global to the entire application.
- Cursor names are unique for a given connection.
- Prepared statement names are global to the connection.
- Dynamic descriptors are global to the application.

Precompiler compatibility

Embedded SQL version 11.1 and later is completely ANSI SQL-89 compliant. Therefore, it is compatible with other precompilers that conform to ANSI-89 standards.

Note To run programs created for pre-11.1 precompilers, you must precompile them again with System 11 and make changes, if necessary. For details, see “Backward compatibility” on page 5.

Precompiler-generated files

The target file is similar to the original input file, except that all SQL statements are converted to Client-Library runtime calls.

The listing file contains the input file’s source statements, plus any informational, warning, or error messages.

The isql script file contains the precompiler-generated stored procedures. The stored procedures are written in **Transact-SQL**.

This chapter provides general information about Embedded SQL.

| Topic | Page |
|---------------------------------------|------|
| Five tasks of an Embedded SQL program | 9 |
| General rules for Embedded SQL | 11 |

Five tasks of an Embedded SQL program

In addition to containing the host language code, an Embedded SQL program performs five tasks. Each Embedded SQL program must perform all these tasks to precompile, compile, and execute. Subsequent chapters discuss these five tasks.

- 1 Establish SQL communication via SQLCA, SQLCODE, or SQLSTATE.

Set up the SQL communication area (SQLCA, SQLCODE, or SQLSTATE) to provide a communication path between the application program and Adaptive Server. These structures contain error, warning, and information message codes that Adaptive Server and Client-Library generate. See Chapter 3, “Communicating with Adaptive Server”

- 2 Declare variables.

Identify host variables used in Embedded SQL statements to the precompiler. See Chapter 4, “Using Variables”

- 3 Connect to Adaptive Server.

Connect the application to Adaptive Server. See Chapter 5, “Connecting to Adaptive Server”

- 4 Send Transact-SQL statements to Adaptive Server.

Send Transact-SQL statements to Adaptive Server to define and manipulate data. See Chapter 6, “Using Transact-SQL Statements”

5 Handle errors and return codes.

Handle and report errors returned by Client-Library and Adaptive Server via SQLCA, SQLCODE, or SQLSTATE. See Chapter 8, “Handling Errors”

Simplified Embedded SQL program

Following is a simple Embedded SQL program. At this point, you need not understand everything shown in the program. Its purpose is to demonstrate the parts of an Embedded SQL program. The details are explained in subsequent chapters.

```
/* Establishing a communication area - Chapter 3 */

exec sql include sqlca;

main()
{

/* Declaring variables - Chapter 4 */

exec sql begin declare section;
CS_CHAR  user[31], passwd[31];
exec sql end declare section;

/*Initializing error-handling routines - Chapter 8 */

exec sql whenever sqlerror call err_p();

/*Establishing Adaptive Server connections - Chapter 5
*/

printf("\nplease enter your userid ");
gets(user);
printf("\npassword ");
gets(passwd);
exec sql connect :user identified by :passwd;

/* Issuing Transact-SQL statements - Chapter 6 */

exec sql update titles set price = price * 1.10;
exec sql commit work;

/* Closing server connections - Chapter 5 */
```



```
    exec sql disconnect all;
}

/* Error-handling routines - Chapter 8 */

err_p()
{
    /* Print the error code and error message */

    printf("\nError occurred: code %d.\n%s",
        sqlca.sqlcode, sqlca.sqlerrm.sqlerrmc);
}
```

General rules for Embedded SQL

The following rules apply to Embedded SQL statements in C programs:

- Embedded SQL statements begin with these keywords:

```
    exec sql
```
- Embedded SQL statements must end with a semicolon:

```
    exec sql sql_statement;
```
- Place `exec sql` at the beginning of the source line except when a C label precedes it:

```
    [label:] exec sql sql_statement;
```
- Embedded SQL keywords are not case sensitive. `exec sql`, `EXEC SQL`, `Exec Sql`, or any other of case mix is equally valid. This manual shows Embedded SQL keywords in lowercase. For example:

```
exec sql commit work;
```

Statement placement

An application program can have Embedded SQL statements wherever C statements are valid.

Comments

Comments placed within Embedded SQL and C statements must follow one of two conventions.

The Transact-SQL convention is:

```
/* comments */
```

The ANSI convention is:

```
-- comments
```

Comments placed outside SQL statements must conform to C programming conventions.

Identifiers

Identifiers are used as function or variable names within your application.

Quotation marks

Enclose literal character strings in Embedded SQL statements within single or double quotation marks. If a character string begins with a double quotation mark, end it with a double quotation mark. If a character string begins with a single quotation mark, end it with a single quotation mark.

Reserved words

Do not use C, Transact-SQL, or Embedded SQL reserved words except as intended by the languages. See Appendix C, “Embedded SQL Constructs” for a list of valid constructs in Embedded SQL.

You can write Embedded SQL keywords in upper-, lower-, or mixed case. This guide shows Embedded SQL keywords in lowercase.

Variable naming conventions

Embedded SQL variables must conform to C naming conventions. Do not place variable names within quotation marks. Applicable quotation marks are inserted automatically when the variable names are replaced with actual values.

While parsing the application, declarations are added for precompiler variables. These variables begin with “_sql”. So, to avoid confusion, do not begin variable names with “_sql”.

Scoping rules

Embedded SQL and precompiler-generated statements adhere to **host language** scoping rules. The whenever statement and cursor names are exceptions.

Statement batches

As in Transact-SQL, you can batch several SQL statements in a single `exec sql` statement. Batches are useful and more efficient when an application must execute a fixed set of Transact-SQL statements each time it runs.

For example, some applications create temporary tables and indexes when they start up. You could send these statements in a single batch. See the *Adaptive Server Reference Manual* for rules about statement batches.

The following restrictions apply to statement batches:

- Statements in a batch cannot return results to the program. That is, a batch can contain no `select` statements.
- All statements in a batch must be valid Transact-SQL statements. You cannot place Embedded SQL statements such as `declare cursor` and `prepare` in a statement batch.
- The same rules that apply to Transact-SQL batches apply to Embedded SQL batches. For example, you cannot put a `use database` statement in an Embedded SQL batch.

Communicating with Adaptive Server

This chapter explains how to enable an application program to receive status information from Adaptive Server.

| Topic | Page |
|---|------|
| Scoping rules: SQLCA, SQLCODE, and SQLSTATE | 16 |
| Declaring SQLCA | 16 |
| Declaring SQLCODE as a standalone area | 18 |
| Using SQLSTATE | 20 |

To create a communication path and declare system variables to be used in communications from Adaptive Server to the application, you must create one of three entities:

- A SQL Communication Area (SQLCA), which includes SQLCODE
- A standalone SQLCODE long integer
- A SQLSTATE character array

SQLCODE, SQLCA, and SQLSTATE are variables to be used in communication from Adaptive Server to the application.

After Adaptive Server executes each Embedded SQL statement, it stores return codes in SQLCA, SQLCODE, or SQLSTATE. An application program can access the variables to determine whether the executed SQL statement succeeded or failed.

Note The precompiler automatically sets SQLCA, SQLCODE, and SQLSTATE variables, which are critical for runtime access to the database. You need not initialize or modify them.

For details on detecting and handling errors, multiple error messages, and other return codes, see Chapter 8, “Handling Errors”

Scoping rules: SQLCA, SQLCODE, and SQLSTATE

You can declare SQLCA anywhere in the application program where a C variable can be declared. The scope of the structure follows C scoping rules.

If you declare SQLCA, SQLCODE, or SQLSTATE within your file, each variable must be in scope for all executable Embedded SQL statements in the file. The precompiler generates code to set each of these status variables for each Embedded SQL statement. So, if the variables are not in scope, the generated code will not compile.

If you do not declare SQLCA, SQLCODE, or SQLSTATE within the file being passed to the precompiler, you must declare SQLCODE within a referenced file. The precompiler assumes a declaration of SQLCODE, and generates code to this effect.

Declaring SQLCA

Warning! Although SQLSTATE is preferred over SQLCODE and SQLCA, this version only fully supports SQLCODE. A future version will support SQLSTATE.

The syntax for declaring SQLCA is:

```
exec sql include sqlca;
```

You can use the Embedded SQL include statement to include other files in your application the same way you would use the C preprocessor `#include` command. You can also set a precompiler command option to specify an *include* file directory. At precompile time, the precompiler searches the path specified in the C compile command. The precompiler uses the *include* file path to search for this file. It opens and reads the included file as if were part of the main file. If the included file cannot be found, the precompile fails.

Multiple SQLCAs

You may have multiple SQLCAs, but each must follow C scoping rules for host variables. Each SQLCA need not be in a separate scope.

SQLCA variables

When the precompiler encounters the include `sqlca` statement, it inserts the SQLCA structure declaration into the application program. SQLCA is a data structure containing precompiler-determined *system variables*, each of which can be accessed independently. Your application program should never directly alter these variables.

SQLCA variables pass information to your application program about the status of the most recently executed Embedded SQL statement.

The following table describes the SQLCA variables that hold status information, return codes, error codes, and error messages generated by Adaptive Server:

Table 3-1: Adaptive Server SQLCA variables

| Variable | Datatype | Description |
|--|----------|--|
| <i>sqlcaid</i> | char | Text string that contains “sqlca”. |
| <i>sqlcabc</i> | long | Length of SQLCA. |
| <i>sqlcode</i> | long | Contains the return code of the most recently executed SQL statement. See “SQLCODE values” on page 19 for return code definitions. |
| <i>sqlwarn[0]</i> to <i>sqlwarn[7]</i> | char | Warning flags. Each flag indicates whether a warning has been issued: a ‘W’ for warning, or a blank space for no warning. Chapter 8 describes the <i>sqlwarn</i> flags. |
| <i>sqlerrm.sqlerrmc</i> [] | char | Error message. |
| <i>sqlerrm.sqlerrml</i> | long | Error message length. |
| <i>sqlerrp</i> | char | Procedure that detected error/warning. |
| <i>sqlerrd[6]</i> | long | Details of error/warning. [2] is the number of rows affected. |

Accessing SQLCA variables

SQLCA variables are members of a C structure, `sqlca`, that is declared by the include `sqlca` statement. To access SQLCA variables, use the C structure member operator (`.`), as in the following example:

```
if (sqlca.sqlwarn[1] == 'W')
{
    printf("\nData truncated");
    return;
}
```

You can also pass the address of the sqlca structure to a function, then access the SQLCA variables within that function with the -> operator. The following example shows a function that works this way:

```
warning(p)
struct sqlca *p;
{
    if (p->sqlwarn[3] == 'W')
    {
        printf("\nIncorrect number of variables in
fetch.\n");
    }
    return;
}
```

SQLCA variables are useful for determining whether an Embedded SQL statement executed successfully. The other SQLCA variables listed in the previous section provide additional information about errors and return codes to help in debugging as well as the normal processing of your application.

SQLCODE within SQLCA

The application should test sqlcode after each statement executes, because Adaptive Server updates it after each execution. As a rule, use the whenever statement, described in Chapter 8, “Handling Errors” to perform this task.

Declaring SQLCODE as a standalone area

Warning! Although SQLSTATE is preferred over SQLCODE and SQLCA, this version only fully supports SQLCODE. A future version will fully support SQLSTATE.

As an alternative to creating a SQLCA, use `SQLCODE` independently. It contains the return code of the most recently executed SQL statement. The benefit of declaring `SQLCODE` as a standalone area is that it executes code faster. If you have no need to review the other information that `SQLCA` holds and are solely interested in return codes, consider using `SQLCODE`.

Despite `SQLCODE`'s faster execution speed, `SQLSTATE` is preferred over `SQLCODE`. `SQLCODE` is supported for its compatibility with earlier versions of Embedded SQL.

Note In a future version, you will be advised to use `SQLSTATE` instead of `SQLCODE` for receiving status results.

Following is an example of declaring `SQLCODE` as a standalone area:

```
long SQLCODE;

exec sql open cursor pub_id;
while (SQLCODE == 0)
{
    exec sql fetch pub_id into :pub_name;
```

For details on debugging any errors `SQLCODE` indicates, see Chapter 8, “Handling Errors”

Following is a table of `SQLCODE` values:

Table 3-2: `SQLCODE` values

| Value | Description |
|-------|--|
| 0 | Statement executed successfully. |
| -n | Error occurred. See Server or Client-Library error messages.”-n” represents the number associated with the error or exception. |
| +100 | No data exists, no rows left after fetch, or no rows met search condition for update, delete, or insert. |

Using SQLSTATE

Warning! Although SQLSTATE is preferred over SQLCODE and SQLCA features, this version only fully supports SQLCODE. A future version will fully support both SQLCA and SQLSTATE.

SQLSTATE is a status parameter. Its codes indicate the status of the most recently attempted procedure—either the procedure completed successfully or an error occurred during the execution of the procedure.

SQLSTATE is a character-string parameter whose exceptions values are described in the following table:

Table 3-3: SQLSTATE values

| Value | Description |
|-----------------|----------------------------------|
| 00XXX | Successful execution |
| 01XXX | Warning |
| 02XXX | No data exists; no rows affected |
| Any other value | Error |

Obtaining SQLSTATE Codes and Error Messages

SQLSTATE messages can be informational, warnings, severe, or fatal. Adaptive Server and Open Client Client-Library generate the majority of SQLSTATE messages. See the appropriate documentation for a complete list of SQLSTATE codes and error messages.

See Appendix A, “Precompiler Warning and Error Messages” for the table of SQLSTATE messages that the precompiler can generate.

Summary

This chapter explained SQLCA, SQLCODE, and SQLSTATE. After a statement executes, Adaptive Server stores return codes and information in SQLCA variables, in a standalone SQLCODE area, or in SQLSTATE. These return codes indicate the failure or success of the statement that most recently executed.

Using Variables

| Topic | Page |
|-------------------------------|------|
| Declaring variables | 21 |
| Using host variables | 28 |
| Using indicator variables | 31 |
| Using arrays | 34 |
| Scoping rules | 35 |
| Datatypes and Adaptive Server | 35 |

This chapter details the following two types of variables that pass data between your application and Adaptive Server:

- Host variables, which are C variables you use in Embedded SQL statements to hold data that is retrieved from and sent to Adaptive Server
- Indicator variables, which you associate with host variables to indicate null data and data truncation

Declaring variables

As discussed in Chapter 3, the precompiler automatically sets the system variables when you include `SQLCA`, `SQLCODE`, or `SQLSTATE` in the application program. However, you must explicitly declare host and indicator variables in a declare section before using them in Embedded SQL statements.

Warning! The precompiler generates some variables, all of which begin with “_sql”. Do not begin your variables with “_sql”, or you may receive an error message or unreliable data.

The precompiler ignores macros and `#include` statements in a declare section. It processes include statements as if the contents of the included file were copied directly into the file being precompiled. The syntax for a declare section with an include statement is:

```
exec sql begin declare section;
    exec sql include "filename";
    ...
exec sql end declare section;
```

Host variable declarations must conform to the C rules for variable declarations. You need not declare all variables in one declare section, since you can have an unlimited number of declare sections in a program.

When you declare variables, you must also specify the **datatype**. See “Datatypes and Adaptive Server” on page 35 for valid datatypes. Alternatively, use the Client-Library typedefs, such as `CS_CHAR`, which are declared in the *cspublic.h* file, in declare sections.

The following example shows two character strings defined in a declare section.

```
exec sql begin declare section;
    CS_CHAR name[20];
    CS_CHAR type[3];
exec sql end declare section;
```

When declaring a host variable, you can also initialize it but only if it is a scalar variable, such as this one:

```
exec sql begin declare section;
    int total = 0;
exec sql end declare section;
```

You cannot initialize an array in its declaration.

Using datatypes

In Embedded SQL, you can use the C datatypes `char`, `int`, `float`, `double`, and `void`. You can use the keywords `const` and `volatile`, though not with structures. You can use the keywords `unsigned`, `long`, and `short`. You can use storage class specifiers: `auto`, `extern`, `register`, and `static`.

```
exec sql begin declare section;
    register int frequently_used_host_variable;
    extern char
        shared_string_host_variable[STRING_SIZE];
```

```
/*
** The const restriction is not enforced by
** the precompiler; only the compiler makes use
** of it.
*/
const float
input_only_host_variable = 3.1415926;
/*
** Be careful. You can declare unsigned
** integers, but if you select a negative
** number into one, you will get an incorrect
** result and no error message.
*/
unsigned long int unsigned_host_variable;
exec sql end declare section;
```

You can declare pointers in the declare section, but you cannot use a pointer as a host variable in an Embedded SQL statement.

```
exec sql begin declare section;
int number;
/*
** It's convenient to declare this here,
** but we won't be using it as a host variable.
*/
int *next_number;
exec sql end declare section;
```

You can use the following Sybase datatypes:

CS_BINARY, CS_BIT, CS_BOOL, CS_CHAR, CS_DATETIME,
CS_DATETIME4, CS_DECIMAL, CS_FLOAT, CS_REAL, CS_IMAGE,
CS_INT, CS_MONEY, CS_MONEY4, CS_NUMERIC, CS_RETCODE,
CS_SMALLINT, CS_TEXT, CS_TINYINT, CS_VOID.

CS_CHAR is treated differently from char; CS_CHAR is null-terminated but not blank-padded; char is null-terminated and blank-padded to the length of the array.

```
/*
** Your #define for the array size doesn't
** have to be in the declare section,
** though it would be legal if it were.
*/
#define MAX_NAME 40;

exec sql begin declare section;
    CS_MONEY salary;
```

```
        CS_CHAR print_this[MAX_NAME];
        char print_this_also[MAX_NAME];
exec sql end declare section;

exec sql select salary into :salary from salaries
        where employee_ID = '01234';
/*
** The CS_MONEY type is not directly printable.
** Here's an easy way to do a conversion.
*/
exec sql select :salary into :print_this;

/*
** This will not be blank-padded.
*/
printf("Salary for employee 01234 is %s.\n",
        print_this);

/*
** This will be blank-padded.
*/
exec sql select :salary into :print_this_also;
printf("Salary for employee 01234 is %s.\n",
        print_this_also);
```

Using type definitions

You can use a type definition (typedef) within a declare section to declare variables. For example:

```
exec sql begin declare section;
/*
** The typedef and the use of the typedef
** can be in separate declare sections
** if the typedef comes first.
** The typedef can even be in an "exec
** sql include file".
*/
typedef int STORE_ID;
STORE_ID current_ID;
exec sql end declare section;

exec sql select store_ID into :current_ID
        from sales_table where
```

```
store_name = 'Furniture Kingdom';
```

Using `#define`

You can use `#define` values in a `declare` section to dimension arrays and initialize variables. When you use `#define` in a host variable declaration, place it before the host variable declaration that uses it. For example, the following two examples are valid:

```
#define PLEN 26
CS_CHAR name[PLEN];
```

and:

```
exec sql begin declare section;
#define PLEN 26
exec sql end declare section;

...

exec sql begin declare section;
CS_CHAR name[PLEN];
exec sql end declare section;
```

You can use `#define` to declare *symbolic names*. Make the declaration before using it in the application. For example, to define “10” symbolically, use this nomenclature:

```
exec sql begin declare section;
#define count_1 10
CS_CHAR var1[count_1];
exec sql end declare section;
```

Declaring an array

The precompiler supports *complex definitions*, which are structures and arrays. You may nest structures, but you cannot have an **array** of structures.

The precompiler recognizes single-dimensional arrays of all datatypes.

The precompiler also recognizes double-dimensional arrays of characters, as the following example demonstrates:

```
#define maxrows 25
int numsales [maxrows];

exec sql begin declare section;
#define DATELEN 30
```

```
#define DAYS_PER_WEEK 7
CS_CHAR days_of_the_week[DAYS_PER_WEEK][DATELEN+1];
exec sql end declare section;
```

You can declare arrays of any datatype. However, to select into an array element, its datatype must be scalar—integer, character, floating point, or pointer. You can select into elements of any scalar array, even an array of structures.

```
exec sql begin declare section;
    int sales_totals[100];
    struct sales_record {
        int total_sales;
        char store_name[40];
    } sales_records[100];
exec sql end declare section;

/*
** If there are fewer than 100 stores,
** this will get the sales totals for all
** of them. If there are more than
** 100, it will cause an error at runtime.
*/
exec sql select total_sales into :sales_totals
    from sales_table;

/*
** This gets the sales for just one store.
*/
exec sql select total_sales into :sales_totals[0]
    from sales_table where store_ID = 'xyz';

/*
** This gets two pieces of information on a single **
store.
*/
exec sql select total_sales, store_name
    into :sales_records[i]
    from sales_table where store_ID = 'abc';
```


Declaring character arrays

A character array can be of type `CS_CHAR` or `char[]`; however, the rules governing these two datatypes differ. When an array of type `char[]` is used as input, the precompiler checks that the array terminates with a null character. If the array is not null terminated, a precompiler runtime function returns an error. In contrast, an array of type `CS_CHAR` is not checked for null termination. Rather, the length of the input continues up to the null character, if present, or to the declared length of the array—whichever comes first.

When used as output, arrays of type `char[]` are padded with space characters (blank-padded) and null terminated. Arrays of type `CS_CHAR` are not blank padded, only null terminated.

A character array is scalar, because it represents a single string. Thus, you can select into an array of characters and get back just a single string. Also, unlike arrays of other datatypes, an array of characters can be a host input variable.

For more information on arrays, see “Using arrays” on page 34.

Declaring unions and structures

You can declare unions and structures, either directly or by using a type definition (`typedef`). You can use an element of a union as a host variable, but not the union as a whole. In contrast, a host variable can be either an entire structure or just one of the structure’s elements. The following example declares a union and a structure:

```
exec sql begin declare section;
    typedef int PAYMENT_METHOD;
    PAYMENT_METHOD method;
    union salary_or_percentage {
        CS_MONEY salary;
        CS_NUMERIC percentage;
    } amount;
    struct employee_record {
        char first_name[30];
        char last_name[30];
        char employee_ID[30];
    } this_employee;
    char *employee_of_the_month_ID = "01234567";
exec sql end declare section;

exec sql select first_name, last_name, employee_ID
into :this_employee
```

```
        from employee_table
        where employee_ID = :employee_of_the_month_ID;
exec sql select payment_type into :method
        from remuneration_table where employee_ID =
        :this_employee.employee_ID;
switch (method) {
    case SALARIED:
        exec sql select salary into
            :amount.salary
            from remuneration_table
            where employee_ID =
            this_employee.employee_ID;
        break;
    case VOLUNTEER:
        exec sql select 0 into
            :amount.salary
        break;
    case COMMISSION:
        exec sql select commission_percentage into
            :amount.percentage
            from remuneration_table
            where employee_ID =
            this_employee.employee_ID;
        break;
}
```

Using host variables

Host variables let you transfer values between Adaptive Server and the application program.

Declare the host variable within the application program's Embedded SQL declare section. Only then can you use the variable in SQL statements.

When you use the variable within an Embedded SQL statement, prefix the host variable with a colon. When you use the variable elsewhere in the program, do not use a colon. When you use several host variables successively in an Embedded SQL statement, separate them with commas or follow the grammar rules of the SQL statement.

The following example demonstrates how to use a variable. *user* is defined in a declare section as a character variable. Then, it is used as a host variable in a select statement:

```
exec sql begin declare section;
  CS_CHAR  user[32];
exec sql end declare section;

exec sql select user_name() into :user;
printf("You are logged in as %s.\n", user);
```

There are four ways to use host variables. Use them as:

- Input variables for SQL statements and procedures
- Result variables
- Status variables from calls to SQL procedures
- Output variables for SQL statements and procedures

Declare all host variables as described in “Declaring variables” on page 21, regardless of their function. Following are instructions for using host variables.

Host input variables

These variables pass information to Adaptive Server. The application program assigns values to them. They hold data used in executable statements such as stored procedures, select statements with where clauses, insert statements with values clauses, and update statements with set clauses.

The following example uses the variables *id* and *publisher* as input variables:

```
exec sql begin declare section;
  CS_CHAR id[7];
  CS_CHAR publisher[5];
exec sql end declare section;
...
exec sql delete from titles where title_id = :id;
exec sql update titles set pub_id = :publisher
  where title_id = :id;
```

Host result variables

These variables receive the results of select and fetch statements.

The following example uses the variable *id* as a **result variable**:

```
exec sql begin declare section;
  CS_CHAR id[5];
```

```
exec sql end declare section;

exec sql select title_id into :id from titles
       where pub_id = "0736" and type = "business";
```

Host status variables

These variables receive the return status values of stored procedures. Status variables indicate whether the stored procedure completed successfully or the reasons it failed.

Declare status variables as two-byte integers (CS_SMALLINT).

The following example uses the variable *retcode* as a **status variable**:

```
exec sql begin declare section;
CS_SMALLINT  retcode;
exec sql end declare section;

exec sql begin transaction;
exec sql exec :retcode = update_proc;
if (retcode != 0)
{
    exec sql rollback transaction;
}
```

Host output variables

These variables pass data from stored procedures to the application program. For more information on stored procedures, see “Using Stored Procedures” on page 6-11 . Use host output variables when stored procedures return the value of parameters declared as out.

The following example uses the variables *par1* and *par2* as output variables:

```
exec sql exec a_proc :par1 out, :par2 out;
```

Using indicator variables

You can associate indicator variables with host variables to indicate when a database value is null. Use a space and, optionally, the indicator keyword, to separate each indicator variable from the host variable with which it is associated. Each **indicator variable** must immediately follow its host variable.

Without indicator variables, Embedded SQL cannot indicate null values.

Indicator variables and server restrictions

Embedded SQL is a generic interface that can run on a variety of servers, including Adaptive Server.

Because it is generic, Embedded SQL does not enforce or reflect any particular server's restrictions.

When writing an Embedded SQL application, keep the application's ultimate target **server** in mind. If you are unsure about what is legal on a server and what is not, consult your server documentation.

Using host variables with indicator variables

Declare host and indicator variables in a declare section before using them anywhere in an application program containing Embedded SQL statements. Declare indicator variables as two-byte integers (short or CS_SMALLINT) in a declare section before using them.

Prefix indicator variables with a colon when using them in an Embedded SQL statement.

The syntax for associating an indicator variable with a host variable is:

```
:host_variable [[indicator] :indicator_variable]
```

The association between an indicator and host variable lasts only for the duration of a statement— that is, for the duration of one `exec sql` statement, or between `open` and `close cursor` statements. A value is assigned to the indicator variable at the same time a value is assigned to the host variable.

Adaptive Server sets the indicator variable only when you assign a value to the host variable. Therefore, you can declare an indicator variable once and reuse it with different host variables in different statements.

You can use indicator variables with output, result, and input variables. When used with output and result variables, Embedded SQL sets the variable to indicate the null status of the associated host variable. When used with input variables, you set the value of the indicator variable to show the null status of the **input variable** before submitting it to Adaptive Server.

Using indicator variables with host output and result variables

When you associate an indicator variable with an output or result variable, Client-Library automatically sets it to one of the following values:

Table 4-1: Indicator variable values when used with output or result variable

| Value | Meaning |
|-------|--|
| -1 | The corresponding database column in Adaptive Server contains a null value. |
| 0 | A non-null value was assigned to the host variable. |
| >0 | An overflow occurred while data was being converted for the host variable. The host variable contains truncated data. The positive number represents the length, in bytes, of the value before it was truncated. |

The following example demonstrates associating the indicator variable *indic* with the result variable *id*:

```
exec sql begin declare section;
  CS_CHAR          id[6];
  CS_SMALLINT      indic;
  CS_CHAR          pub_name[41];
exec sql end declare section;

exec sql select pub_id into :id indicator :indic
  from titles where title
    like "%Stress%";

if (indic == -1)
{
  printf("\npub_id is null");
}
else
{
  exec sql select pub_name into :pub_name
    from publishers where pub_id = :id;
  printf("\nPublisher: %s", pub_name);
}
```

Using indicator variables with host input variables

When you associate an indicator variable with an input variable, you must explicitly set the indicator variable, using the values in the following table as a guide.

Table 4-2: Indicator variable values used with input variable

| Value | Meaning |
|-------|--|
| -1 | Treat the corresponding input as a null value. |
| 0 | Assign the value of the host variable to the column. |

You must supply host language code to test for a null input value and set the indicator variable to -1. This informs Client-Library of a null value. When you set the indicator variable to -1, null is used regardless of the host variable's actual value.

The following example demonstrates associating an indicator variable with an input variable. The database royalty column is set to a null value because *indic* is set to -1. Changing the value of *indic* changes the value of royalty.

```
exec sql begin declare section;
  CS_SMALLINT      indic;
  CS_INT           royalty;
exec sql end declare section;

indic = -1;
exec sql update titles set royalty = :royalty
      :indic where pub_id = "0736";
```

Host variable conventions

A **host variable** name must conform to C naming conventions.

You can use a host variable in an Embedded SQL statement wherever a Transact-SQL literal can be used in a Transact-SQL statement at the same location..

A host variable must conform to the valid precompiler datatypes. The datatype of a host variable must be compatible with the datatype of the database column values returned. See Table 4-4 on page 37 and Table 4-5 on page 38 for details.

You cannot use host language reserved words and Embedded SQL keywords as variable names.

A host variable cannot represent Embedded SQL keywords or database objects, except as specified in **dynamic SQL**. For more information on using host variables to represent keywords for database objects, see “Chapter 7, “Using Dynamic SQL”

When a host variable represents a character string in a SQL statement, do not place it within quotes.

The following example is invalid because the precompiler inserts quotes around values when necessary. You should not type the quotes.

```
strcpy (p_id, "12345");
exec sql select pub_id into :p_id from publishers
where pub_id like ":p_id";
```

The following example is valid:

```
strcpy (p_id, "12345");
exec sql select pub_id into :p_id from publishers
where pub_id like :p_id;
```

Using arrays

An array is a group of related pieces of data associated with one variable. You can use arrays as output variables for the into clause of select and fetch statements.

For example:

```
exec sql begin declare section;
  CS_CHAR  au_array [100] [30];
exec sql end declare section;

exec sql
  select au_lname
  into :au_array
  from authors;
```

Note You can fetch a single item anywhere into an array. However, you can fetch multiple rows only into the beginning of an array.

For details on using arrays with select and fetch into, see “Selecting Multiple Rows via Arrays” on page 6-3 .

Multiple arrays

When you use multiple arrays within a single SQL statement, they should be the same size. Otherwise, you will receive an error message.

Scoping rules

The precompiler supports the C programming rules for variable scoping. Host variables defined within nested programs can use the external clause plus the variable name. For example:

```
FILE 1:
CS_CHAR  username[31]
main()
{
    sub1();
    printf("%s\n", username);
}

FILE 2:
void sub1()
{
    exec sql begin declare section;
    extern char username[31];
    exec sql end declare section;

    exec sql select USER() into :username;
    return;
}
```

Datatypes and Adaptive Server

Host variable datatypes must be compatible with the datatypes of the corresponding database columns. So, before writing your application program, check the datatypes of the database columns. To ensure that your host variables are compatible with the Adaptive Server datatypes, use the Sybase-supplied type definitions.

The following table shows and briefly describes the equivalent datatypes. For detailed descriptions of each Adaptive Server datatype, see the *Adaptive Server Enterprise Reference Manual*.

Table 4-3: Comparison of C and Adaptive Server compatible datatypes

| Sybase-supplied typedef | Description | C datatype | Adaptive Server datatype |
|----------------------------|----------------------------------|---------------|--------------------------------|
| CS_BINARY | Binary type | unsigned char | binary, varbinary |
| CS_BIT | Bit type | unsigned char | boolean |
| CS_CHAR | Character type | char[n] | char, varchar |
| CS_DATETIME | 8-byte datetime type | None | datetime |
| CS_DATETIME4 | 4-byte datetime type | None | smalldatetime |
| CS_TINYINT | 1-byte integer type | unsigned char | tinyint |
| CS_SMALLINT | 2-byte integer type | short | smallint |
| CS_INT | 4-byte integer type | long | int |
| CS_DECIMAL | Decimal type | None | decimal |
| CS_NUMERIC | Numeric type | None | numeric |
| CS_FLOAT | 8-byte float type | double | float |
| CS_REAL | 4-byte float type | float | real |
| CS_MONEY | 8-byte money type | None | money |
| CS_MONEY4 | 4-byte money type | None | smallmoney |
| CS_TEXT | Text type -y option required | unsigned char | text |
| CS_IMAGE | Image type -y option required | unsigned char | image |

Converting datatypes

The precompiler automatically compares the datatypes of host variables with the datatypes of table columns in Adaptive Server. If the Adaptive Server datatype and the host language datatype are compatible but not identical, the precompiler converts one type to the other. Datatypes are compatible if the precompiler can convert the data from one type to the other. If the datatypes are incompatible, a conversion error occurs at run time and `sqlcode` is set to `<0`.

Be careful when converting a longer datatype into a shorter one, such as a four-byte into two-byte, because there is always a possibility of truncating data. If a truncation occurs, `sqlwarn1` is set to “W”.

Converting datatypes for result variables

The following table indicates which data conversions are valid for result variables. A bullet indicates that conversion is possible, but be aware that certain types of errors can result if you are not careful when choosing host variable datatypes.

Table 4-4: Datatype conversions for result variables

| From: Adaptive Server datatype | To: C datatype | | | | | | |
|---|----------------|-------------|--------|---------|---------|----------|-------------|
| | CS_TINYINT | CS_SMALLINT | CS_INT | CS_REAL | CS_CHAR | CS_MONEY | CS_DATETIME |
| char | • | • | • | • | • | • | • |
| varchar | • | • | • | • | • | • | • |
| bit | • | • | • | • | • | • | |
| binary | • | • | • | • | • | • | |
| tinyint | • | • | • | • | • | • | |
| smallint | • | • | • | • | • | • | |
| int | • | • | • | • | • | • | |
| float | • | • | • | • | • | • | |
| money | • | • | • | • | • | • | |
| datetime | | | | | • | | • |
| decimal | • | • | • | • | • | • | |
| numeric | • | • | • | • | • | • | |

Converting datatypes for input variables

The following table shows the valid data conversions for input variables. A bullet indicates that conversion is possible, an "x" indicates that conversion is required. Errors, including truncation, can result if you choose nonconvertible host variable datatypes.

Table 4-5: Datatype conversions for input variables

| To: Adaptive Server datatype | | | | | | | | | | | |
|---|----------|----------|----------|----------|----------|----------|----------|----------|----------|----------|--|
| From: C datatype | tinyint | bit | smallint | int | float | char | money | datetime | decimal | numeric | |
| unsigned char | • | • | • | • | • | <i>x</i> | • | | • | • | |
| short int | • | • | • | • | • | <i>x</i> | • | | • | • | |
| long int | • | • | • | • | • | <i>x</i> | • | | • | • | |
| double float | • | • | • | • | • | <i>x</i> | • | | • | • | |
| char | <i>x</i> | <i>x</i> | <i>x</i> | <i>x</i> | <i>x</i> | • | <i>x</i> | • | <i>x</i> | <i>x</i> | |
| money | • | • | • | • | • | • | • | | • | • | |
| datetime | | | | | | <i>x</i> | | • | | | |
| <i>x</i> – indicates that an explicit conversion is required. | | | | | | | | | | | |

Connecting to Adaptive Server

This chapter explains how to connect an Embedded SQL program to Adaptive Server and describes how to specify servers, user names, and passwords.

| Topic | Page |
|-----------------------------------|------|
| Connecting to a server | 39 |
| Changing the current connection | 41 |
| Establishing multiple connections | 41 |
| Disconnecting from a server | 44 |

Connecting to a server

A connection enables an Embedded SQL program to access a database and perform SQL operations.

Use the connect statement to establish a connection between an application program and Adaptive Server. If an application uses both C and COBOL languages, the first connect statement must be issued from a COBOL program. See *Open Client Embedded SQL/COBOL Programmer's Manual* for information.

The syntax for the connect statement is:

```
exec sql connect :user [identified by :password]
[at :connection_name] [using :server]
[labelname labelname labelvalue labelvalue...]
```

Each of the following sections describes one of the connect statement's arguments. Only the *user* argument is required for the connect statement. The other arguments are optional.

user

user is a host variable or quoted string that represents an Adaptive Server user name. The user name must be valid for the server specified.

password

password is a host variable or quoted string that represents the password associated with the specified user name. This argument is necessary only if a password is required to access Adaptive Server. If the password argument is null, the user does not need to supply a password.

connection_name

connection_name uniquely identifies the Adaptive Server connection. It can be a quoted literal. You can create an unlimited number of connections in an application program, one of which can be unnamed. *connection_name* has a maximum size of 128 characters.

When you use *connection_name* in a connect statement, all subsequent Embedded SQL statements that specify the same connection automatically use the server indicated in the connect statement. If the connect statement specifies no server, the default server is used. See the *Open Client/Server Programmer's Supplement* for details on how the default server is determined.

Note To change the current server connection, use the set connection statement described in “Changing the current connection” on page 41.

An Embedded SQL statement should only reference a *connection_name* specified in a connect statement. At least one connect is required for each server that the application program uses.

server

server is a host variable or quoted string that represents a server name. *server* must be a character string that uniquely and completely identifies a server.

connect example

The following example connects to the server SYBASE using the password “passes”.

```
exec sql begin declare section;

CS_CHAR user[16];
CS_CHAR passwd[16];
CS_CHAR server[BUFSIZ];

exec sql end declare section;

strcpy(server, "SYBASE");
strcpy(passwd, "passes");
strcpy(user, "my_id");

exec sql connect :user identified by :passwd using
:server;
```

Changing the current connection

Use the set connection statement to change the current connection. The statement’s syntax is as follows:

```
exec sql set connection {connection_name | default}
```

where default is the unnamed connection, if any.

The following example changes the current connection:

```
exec sql connect "ME" at connect1 using "SERVER1";
exec sql connect "ME" at connect2 using "SERVER2";
exec sql set connection connect1;
exec sql select user_id() into :myid;
```

Establishing multiple connections

Some Embedded SQL applications require or benefit from having more than one active Adaptive Server connection. For example:

- An application that requires multiple Adaptive Server login names can have a connection for each login account.
- By connecting to more than one server, an application can simultaneously access data stored on different servers.

A single application can have multiple connections to a single server or multiple connections to different servers. Use the connect statement's *atconnection_name* clause to name additional connections for an application.

If you open a connection and then another new named or unnamed connection, the new connection is the current connection.

Note If you are generating stored procedures with the precompiler for appropriate SQL statements, then for each Embedded SQL file, the precompiler generates a single file for all stored procedures on all servers. You can load this file into the appropriate server(s). Although the server(s) will report warnings and errors about being unable to read the procedures intended for other servers, ignore them. The stored procedures appropriate for each server will load properly on that server. Be sure to load the stored procedures on all applicable servers or your queries will fail.

Naming a connection

The following table shows how a connection is named:

Table 5-1: How a connection is named

| If this clause is used | But without | Then, the connection name is |
|---------------------------|-------------|---|
| at <i>connection_name</i> | | <i>connection_name</i> |
| using <i>server_name</i> | at | <i>server_name</i> |
| None | | Actual name of the "DEFAULT" connection |

Invalid statements with the *at* clause

The following statements are invalid with the *at* clause:

- connect
- begin declare section

- end declare section
- include file
- include sqlca
- set connection
- whenever

Using Adaptive Server connections

Specify a connection name for any Embedded SQL statement that you want to execute on a connection other than the default unnamed connection. If your application program uses only one connection, you can leave the connection unnamed. Then, you do not need to use the `at` clause.

The syntax for using multiple connections is:

```
exec sql [at connection_name] sql_statement;
```

where `sql_statement` is a Transact-SQL statement.

The following example shows how two connections can be established to different servers and used in consecutive statements:

```
...

exec sql begin declare section;
CS_CHAR user[16];
CS_CHAR passwd[16];
CS_CHAR name;
CS_INT  value, test;
CS_CHAR server_1[BUFSIZ];
CS_CHAR server_2[BUFSIZ];
exec sql end declare section;
...
strcpy (server_1, "sybase1");
strcpy (server_2, "sybase2");
strcpy(user, "my_id");
strcpy(passwd, "mypass");

exec sql connect :user identified by :passwd
at connection_2 using :server_2;

exec sql connect :user identified by :passwd using
:server_1;
```

```
/* This statement uses the current "server_1"
connection */
exec sql select royalty into :value from authors
where author = :name;

if (value == test)
{
/* This statement uses connection "connection_2" */
exec sql at connection_2 update authors
set column = :value*2
where author = :name;
}
```

Disconnecting from a server

The connections your application program establishes remain open until you explicitly close them or your program terminates. Use the disconnect statement to close a connection between the application program and Adaptive Server.

The statement's syntax is as follows:

```
exec sql disconnect {connection_name | current | DEFAULT
| all}
```

- current specifies the current connection
- DEFAULT specifies the unnamed default connection
- all specifies all connections currently in use

The disconnect statement:

- 1 Rolls back the current transactions ignoring any established savepoints.
- 2 Closes the connection.
- 3 Drops all temporary objects, such as tables.
- 4 Closes all open cursors.
- 5 Releases locks established for the current transactions.
- 6 Terminates access to the server's databases.

disconnect does not implicitly commit current transactions.

Warning! Before the program exits, make sure you perform `anexec sql disconnect` or `exec sql disconnect` all statement for each open connection. In some configurations, SQL-Server may not be notified when a **client** exits without disconnecting. If this happens, resources held by the application will not be released.

This chapter explains how to use Transact-SQL statements with Embedded SQL and host variables. It also explains how to use *stored procedures*, which are collections of SQL statements stored in Adaptive Server. Since stored procedures are compiled and saved in the **database**, they execute quickly without being recompiled each time you invoke them.

| Topic | Page |
|---|------|
| Transact-SQL statements in Embedded SQL | 47 |
| Selecting rows | 48 |
| Grouping statements | 61 |

Transact-SQL statements in Embedded SQL

exec sql syntax

Embedded SQL statements must begin with the keywords `exec sql`. The syntax for Embedded SQL statements is:

```
exec sql [at  
connection_name] sql_statement;
```

where:

- *connection_name* specifies the connection for the statement. See Chapter 5, “Connecting to Adaptive Server,” for a description of connections. The `at` keyword is valid for Transact-SQL statements and the `disconnect` statement.
- *sql_statement* is one or more Transact-SQL statements.

Invalid statements

Except for the following Transact-SQL statements, all Transact-SQL statements are valid in Embedded SQL:

- print
- readtext
- writetext

Transact-SQL statements that differ in Embedded SQL

While most Transact-SQL statements retain their functionality and syntax when used in Embedded SQL, the select, update, and delete statements (the Data Manipulation Language, or DML, statements) can be slightly different in Embedded SQL:

- The following four items are specific to the into clause of the select statement.
 - The into clause can assign one row of data to scalar host variables. This clause is valid only for select statements that return just one row of data. If you select multiple rows, a negative SQLCODE results, and only the first row is returned.
 - If the variables in an into clause are arrays, you can select multiple rows. If you select more rows than the array holds, an exception of SQLCODE <0 is raised, and the extra rows are lost.
 - select cannot return multiple rows of data in host variables, except through a cursor or by selecting into an array.
- The update and delete statements can use the search condition where current of *cursor_name*.

Selecting rows

There can be a maximum of 256 columns in a select statement. For the complete listing of the select statement's syntax, see the *Adaptive Server Enterprise Reference Manual*.

Selecting one row

When you use the `select` statement without a cursor or array, it can return only one row of data. Embedded SQL requires a cursor or an array to return more than one row of data.

In Embedded SQL, a `select` statement must have an `into` clause. The clause specifies a list of host variables to be assigned values.

Note The current Embedded SQL precompiler version does not support `into` clauses that specify tables.

The syntax of the Embedded SQL `select` statement is:

```
exec sql [at connect_name ]
        select [all | distinct] select_list into
        :host_variable[[indicator]:indicator_variable]
        [, :host_variable
        [[indicator]:indicator_variable]...];
```

For additional information on `select` statement clauses, see the *Adaptive Server Enterprise Reference Manual*.

The following `select` statement example accesses the `authors` table in the `pubs2` database and assigns the value of `au_id` to the host variable `id`:

```
exec sql select au_id into :id from authors
        where au_lname = "Stringer";
```

Selecting multiple rows via arrays

You can return multiple rows with arrays. The two array actions involve selecting and fetching into arrays.

select into arrays

Use the `select into array` method when you know the maximum number of rows that will be returned. If a `select into` statement attempts to return more rows than the array can hold, the statement returns the maximum number of rows that the smallest array can hold.

Following is an example of selecting into an array:

```
exec sql begin declare section;
        CS_CHAR titleid_array [100] [6];
```

```
exec sql end declare section;
...
exec sql select title_id into :titleid_array
      from titles;
```

Indicator arrays

To use indicators with array fetches, declare an array of indicators of the same length as the *host_variable* array, and use the syntax for associating the indicator with the host variable.

Example

```
exec sql begin declare section;
      int item_numbers [100];
      short i_item_numbers [100];
exec sql end declare section;
...
exec sql select it_n from item.info
      into :item_numbers :i_item_numbers;
...
```

fetch into: batch arrays

fetch returns the specified number of rows from the currently active set. Each *fetch* returns the subsequent batch of rows. For example, if the currently active set has 150 rows and you select and *fetch* 60 rows, the first *fetch* returns the first 60 rows. The next *fetch* returns the following 60 rows. The third *fetch* returns the last 30 rows.

Note To find the total number of rows fetched, see the *SQLERRD* variable in the *SQLCA*, as described in “*SQLCA* variables” on page 17.

Cursors and arrays

Use the *fetch into array* method when you do not know the number of rows to be returned into the array. Declare and open a cursor, then use *fetch* to retrieve *groups of rows*. If a *fetch into* attempts to return more rows than the array can hold, the statement returns the maximum number of rows that the smallest array can hold and *SQLCODE* displays a negative value, indicating that an error or exception occurred.

Selecting multiple rows via cursors

You can also use cursors to return multiple rows. A **cursor** is a data selector that passes multiple rows of data to the host program, one row at a time. The cursor indicates the first row, also called the **current row**, of data and passes it to the host program. With the next fetch statement, the cursor advances to the next row, which has now become the current row. This continues until all requested rows are passed to the host program.

Use a cursor when a select statement returns more than one row of data. Client-Library tracks the rows Adaptive Server returns and buffers them for the application. To retrieve data with a cursor, use the fetch statement.

The cursor mechanism is composed of these statements:

- declare
- open
- fetch
- update and delete where current of
- close

Cursor scoping rules

The rules that govern the initial scope of a cursor differ, depending on whether the cursor is static or dynamic. However, after a static cursor is opened or a dynamic cursor is declared, the scoping rules for both types of cursors are the same. The rules are as follows:

Until a static cursor is open, its scope is limited to the file where the cursor was declared. Any statement that opens the static cursor must be in this file. After a static cursor is open, its scope is limited to the connection on which the cursor was opened.

A dynamic cursor does not exist until it is declared. After it is declared, its scope is limited to the connection on which it was declared.

A cursor name can be open on more than one connection at a time.

Statements that fetch, update, delete, or close a cursor can appear in files other than the one where the cursor is declared. Such statements, however, must execute on the connection where the cursor was opened.

Cursor names must be unique within a program. If, at run time, an application attempts to declare two identically named cursors, the application fails. The following error message results:

There is already another cursor with the name 'XXX'.

Declaring cursors

Declare a cursor for each `select` statement that returns multiple rows of data. You must declare the cursor before using it, and you cannot declare it within a declare section.

The syntax for declaring a cursor is:

```
exec sql declare cursor_name cursor
        for select_statement ;
```

where:

- *cursor_name* identifies the cursor. The name must be unique and have a maximum of 30 characters. The name must begin with a letter of the alphabet or with the symbols “#” or “_”.
- *select_statement* is a select statement that can return multiple rows of data. The syntax for select is the same as that shown in the *Adaptive Server Enterprise Reference Manual*, except that you cannot use `into` or `compute` clauses.

Example: Declaring a cursor

The following example demonstrates declaring cursors:

```
exec sql declare c1 cursor for
        select type, price from titles
        where type like :wk-type;
```

In this example, *c1* is declared as a cursor for the rows that will be returned for the type and price columns. The precompiler generates no code for the declare cursor statement. It simply stores the select statement associated with the cursor.

When the cursor opens, the select statement or procedure in the declare cursor statement executes. When the data is fetched, the results are copied to the host variables.

Note Each cursor's open and declare statements must be in the same file. Host variables used within the declare statement must have the same scope as the one in which the open statement is defined. However, once the cursor is open, you can perform fetch and update or delete where current of on the cursor in any file.

The declare cursor statement is a declaration, not an executable statement. Therefore, it may appear anywhere in a file; SQLCODE, SQLSTATE, and SQLCA are not set after this statement.

Opening cursors

To retrieve the contents of selected rows, you must first open the cursor. The open statement executes the select statement associated with the cursor in the declare statement. The open statement's syntax is:

```
exec sql open cursor_name;
```

After you declare a cursor, you can open it wherever you can issue a select statement. When the open statement executes, Embedded SQL substitutes the values of any host variables referenced in the declare cursor statement's where clause.

The number of cursors you may have open depends on the resource demands of the current session. Adaptive Server does not limit the number of open cursors. However, you cannot open a currently open cursor. Doing so results in an error message.

While an application executes, you can open a cursor as many times as necessary, but you must close it before reopening it. You need not retrieve all the rows from a cursor result set before retrieving rows from another cursor result set.

Fetching data

Use a fetch statement to retrieve data through a cursor and assign it to host variables. The syntax for the fetch statement is:

```
exec sql [at connect_name] fetch cursor_name  
into : host_variable  
[[ indicator]: indicator_variable ]
```

```
[, : host_variable  
[[ indicator]: indicator_variable ]...];
```

where there is one *host_variable* for each column in the result rows.

Prefix each host variable with a colon and separate it from the next host variable with a comma. The host variables listed in the fetch statement must correspond to Adaptive Server values that the select statement retrieves. Thus, the number of variables must match the number of returned values, they must be in the same order, and they must have compatible datatypes.

An *indicator_variable* is a 2-byte signed integer declared in a previous declare section. If a value retrieved from Adaptive Server is null, the runtime system sets the corresponding indicator variable to -1. Otherwise, the indicator is set to 0.

The data that the fetch statement retrieves depends on the cursor position. The cursor points to the *current row*. The fetch statement always returns the current row. The first fetch retrieves the first row and copies the values into the host variables indicated. Each fetch advances the cursor to the next result row.

Normally, you should place the fetch statement within a loop so that all values returned by the select statement can be assigned to host variables.

The following loop uses the whenever not found statement:

```
/* Initialize error-handling routines */  
exec sql whenever sqlerror call err_handle();  
exec sql whenever not found goto end_label;  
for (;;)   
{  
    exec sql fetch cursor_name  
        into :host_variable [, host_variable];  
    ...  
}  
end_label;
```

This loop continues until all rows are returned or an error occurs. In either case, sqlcode or sqlstate, which the whenever statement checks after each fetch, indicates the reason for exiting the loop. The error-handling routines ensure that an action is performed when either condition arises, as described in Chapter 8, “Handling Errors.”

Using cursors to update and delete rows

To update or delete the current row of a cursor, specify the where current of *cursor_name* as the search condition in an update or delete statement.

To update rows through a cursor, the result columns to be used in the updates must be updatable. They cannot be the result of SQL expressions such as `max(colname)`. In other words, there must be a valid correspondence between the result column and the database column to be updated.

The following example demonstrates how to use a cursor to update rows:

```
exec sql declare c1 cursor for
    select title_id, royalty, ytd_sales
    from titles
    where royalty < 25;

exec sql open c1;

for (;;)
{
    exec sql fetch c1 into :title, :roy, :sales;
    if (SQLCODE == 100) break;
    if (sales > 10000)
        exec sql update titles
            set royalty = :roy + 2
            where current of c1;
}
exec sql close c1;
```

The Embedded SQL syntax of the update and delete statements is the same as in Transact-SQL, with the addition of the `where current of cursor_name` search condition.

For details on determining table update protocol and locking, see the *Transact-SQL User's Guide*.

Closing cursors

Use the close statement to close an open cursor. The syntax for the close statement is:

```
exec sql [at connection] close cursor_name;
```

To reuse a closed cursor, issue another open statement. When you re-open a cursor, it points to the first row. Do not issue a close statement for a cursor that is not open or an error will result.

Cursor example

The following example shows how to nest two cursors. Cursor c2 depends upon the value fetched into *title-id* from cursor c1.

The program gets the value of *title-id* at open time, not at declare time.

```
exec sql include sqlca;

main()
{
    exec sql begin declare section;
        CS_CHAR title_id[7];
        CS_CHAR title[81];
        CS_INT  totalsales;
        CS_SMALLINT salesind;
        CS_CHAR au_lname[41];
        CS_CHAR au_fname[21];

    exec sql end declare section;

    exec sql whenever sqlerror call error_handler();
    exec sql whenever sqlwarning call error_handler();
    exec sql whenever not found continue;
    exec sql connect "sa" identified by "";
    exec sql declare c1 cursor for
        select title_id, title, total_sales from pubs2..titles;
    exec sql declare c2 cursor for
        select au_lname, au_fname from pubs2..authors
        where au_id in (select au_id from pubs2..titleauthor
            where title_id = :title_id);
    exec sql open c1;
    for (;;)
    {
        exec sql fetch c1 into :title_id, :title,
            :totalsales :salesind;
```

```
if (sqlca.sqlcode ==100)
    break;
printf("\nTitle ID: %s, Total Sales: %d", title_id, totalsales);
printf("\n%s", title);
if (totalsales > 10)
{
    exec sql open c2;
    for (;;)
    {
        exec sql fetch c2 into :au_lname, :au_fname;
        if (sqlca.sqlcode == 100)
            break;
        printf("\n\tauthor: %s, %s", au_lname, au_fname);
    }
    exec sql close c2;
}
}
exec sql close c1;
exec sql disconnect all;
}
error_handler()
{
printf("%d\n%s\n",sqlca.sqlcode,sqlca.sqlerrm.sqlerrmc);
exec sql disconnect all;
exit(0);
}
```

See the online sample programs for more examples using cursors. For details on accessing the online examples, see the *Open Client/Server Programmer's Supplement* for your platform.

Using stored procedures

There are two types of *stored procedures*—user-defined and precompiler-generated. Both types run faster than stand-alone statements because Adaptive Server preoptimizes the queries. You create user-defined stored procedures, and the precompiler generates stored procedures.

User-defined stored procedures

With Embedded SQL version 11.1 you can execute stored procedures with select statements that return data rows. Stored procedures can return results to your program through output parameters and through a return status variable.

Stored procedure parameters can be either input, output, or both input and output. For details on stored procedures, see the *Transact-SQL User's Guide*.

Syntax

Valid stored procedure names consist of upper- and lowercase letters of the alphabet, “\$”, “_”, and “#”.

Do not include the use statement in a stored procedure.

To execute a stored procedure, use the following syntax:

```
exec [[:status_variable =]status_value] procedure_name
[([@parameter_name=]parameter_value [out[put]]),...]
[into :hostvar_1 [:indicator_1]
[, hostvar_n [indicator_n, ...]]]
[with recompile];
```

where:

- *status_variable* can return either an Adaptive Server return status value or a return code, which either indicates that the stored procedure completed successfully or gives the reasons for the failure. Negative status values are reserved for Adaptive Server use. See the *Transact-SQL User's Guide* for a list of return status values for stored procedures.
- *status_value* is the value of the stored procedure return status variable *status_variable*.
- *procedure_name* is the name of the stored procedure to execute.
- *parameter_name* is the name of a variable in the stored procedure. You can pass parameters either by position or by name. If one parameter is named, all of them must be named. For more information on stored procedures, see the *Transact-SQL User's Guide*.

- *parameter_value* is a literal constant whose value is passed to the stored procedure.
- *output* indicates that the stored procedure returns a parameter value. The matching parameter in the stored procedure must also have been created using the *output* keyword.
- *into:hostvar_1* causes row data returned from the stored procedure to be stored in the specified host variables (*hostvar_1* through *hostvar_n*). Each host variable can have an indicator variable.
- *indicator_n* is a two-byte host variable declared in a previous *declare* section. If the value for the associated *hostvar_n* is null, the indicator variable is set to -1 when the row data is retrieved. If truncation occurs, the indicator variable is set to the actual length of the result column. Otherwise, the indicator variable is 0.
- *with recompile* causes Adaptive Server to create a new query plan for this stored procedure each time the procedure executes.

Note In Embedded SQL, the *exec* keyword is required to execute a stored procedure. You cannot substitute *execute* for *exec*.

Stored procedure example

The following example shows a call to a **stored procedure** where *retcode* is a status variable, *a_proc* is the stored procedure, *par1* is an input parameter, and *par2* is an output parameter:

```
exec sql begin declare section;
    CS_INT    par1;
    CS_INT    par2;
    CS_SMALLINT retcode;
exec sql end declare section;
...
exec sql exec :retcode = a_proc :par1, :par2 out;
```

The next example demonstrates the use of a stored procedure that retrieves data rows. The name of the stored procedure is *get_publishers*:

```
exec sql begin declare section;
    CS_CHAR    pub_id(4);
    CS_CHAR    name(45);
    CS_CHAR    city(25);
    CS_CHAR    state(2);
    CS_SMALLINT retcode;
```

```
exec sql end declare section;

. . .
exec sql exec :retcode = get_publishers :pub_id
                        into :name :city :state;
```

See Chapter 10, “Embedded SQL Statements: Reference Pages” for a more detailed example of the `exec` statement.

Conventions

The datatypes of the stored procedure parameters must be compatible with the C host variables. Client-Library only converts certain combinations. See Chapter 4, “Using Variables” for a table of compatible datatypes.

Precompiler-generated stored procedures

You can set an optional command line switch so that the precompiler automatically generates stored procedures that can optimize the execution of Transact-SQL statements in your program.

For the list of precompiler command line option switches, see the *Open Client/Server Programmer's Supplement*.

Follow these steps to activate precompiler-generated stored procedures:

- 1 Set the appropriate command line switch so that the precompiler automatically generates stored procedures for the Transact-SQL statements to be optimized.

The precompiler generates an `isql` file containing statements that generate the stored procedures.

- 2 Use interactive SQL (the `isql` program) to execute the file.

This loads the stored procedures on Adaptive Server. The precompiler also creates the stored procedure calls in its output file.

By default, precompiler-generated stored procedures have the same name as the source program, minus any file extensions. The stored procedures are numbered sequentially and the file name and number are separated by a semicolon (“;”).

For example, the stored procedures for a source program named `test1.pc`, would be named `test1;1` through `test1;n`, where *n* is the number of the source program's last stored procedure.

Optionally, you can set a command line flag that lets you alter the stored procedures' names. By using this flag, you can test a modified application without deleting a stored procedure already in production. After successfully testing the application, you can precompile it without the flag to install the stored procedure.

Note When you issue the declare cursor statement, only the select clause is saved as a stored procedure. If an application has syntax errors, the precompiler generates neither the target file nor stored procedures.

Grouping statements

Statements can be grouped for execution by batch or by transactions.

Grouping statements by batches

A batch is a group of statements you submit as one unit for execution. The precompiler executes all Transact-SQL statements between the `exec sql` and `;` keywords in batch mode.

Although the precompiler saves stored procedures, it does not save batches for re-execution. The batch is effective only for the current execution.

The precompiler supports only batch mode statements that return no result sets.

```
exec sql insert into TABLE1 values (:val1)
      insert into TABLE2 values (:val2)
      insert into TABLE3 values (:val3);
```

The three insert statements are processed as a group, which is more efficient than being processed individually. Use the `get diagnostics` method of error handling with batches. For details, see “Using get diagnostics” on page 89.

These statements are legal within a batch because none of them returns results. For more information on batches, see the *Transact-SQL User's Guide*.

Grouping statements by transactions

A **transaction** is a single unit of work, whether the unit consists of one or 100 statements. The statements in the transaction execute as a group, so either all or none of them execute.

The precompiler supports two transaction modes—default ANSI/ISO and optional Transact-SQL. In the Transact-SQL transaction mode, each statement is implicitly committed unless it is preceded by a begin transaction statement.

The Transact-SQL mode uses relatively few system resources, while the default ANSI/ISO transaction mode can dramatically affect system response time. For details on choosing the appropriate mode for your application, see the *Transact-SQL User's Guide*.

You can use a precompiler option to determine the transaction mode of the connections your application opens. See the *Open Client/Server Programmer's Supplement* for details.

Transact-SQL transaction mode

In this optional Transaction mode, the Embedded SQL syntax is the same as that used in Transact-SQL. The begin transaction statement explicitly initiates transactions.

The syntax of the Embedded SQL transaction statements is:

```
exec sql [at connect_name ]
        begin transaction [ transaction_name ];

exec sql [at connect_name]
        save transaction [ savepoint_name];

exec sql [at connect_name] commit transaction
        [ transaction_name ];

exec sql [at connect_name] rollback transaction
        [ savepoint_name | transaction_name ];
```

Note disconnect rolls back all open transactions. For details on this statement, see Chapter 5, “Connecting to Adaptive Server”

When you issue a begin transaction on a connection, you must also issue a save, commit, or roll back transaction on the same connection. Otherwise, an error is generated.

Default ANSI/ISO transaction mode

ANSI/ISO SQL does not provide a save transaction or begin transaction statement. Instead, transactions begin implicitly when the application program executes one of the following statements:

- delete
- insert
- select
- update
- open
- exec

The transaction ends explicitly when you issue either a commit work or rollback work statement. You must use the ANSI/ISO forms of the commit and rollback statements.

The syntax is:

```
exec sql commit [work] end-exec  
exec sql rollback [work] end-exec
```

Extended transactions

An **extended transaction** is a unit of work that has multiple Embedded SQL statements. In the Transact-SQL **transaction mode**, you surround an extended transaction statement with the begin transaction and commit transaction statements.

In the default ANSI mode, you are constantly within an extended transaction. When you issue a commit work statement, the current extended transaction ends and another begins. For details, see the *Transact-SQL User's Guide* .

Note Unless the database option `allow ddl in tran` is set, do not use the following Transact-SQL statements in an extended, ANSI-mode transaction: `alter database`, `create database`, `create index`, `create table`, `create view`, `disk init`, `grant`, `load database`, `load transaction`, `reconfigure`, `revoke`, `truncate table`, and `update statistics`.

Using Dynamic SQL

This chapter explains dynamic SQL, an advanced methodology that lets users of Embedded SQL applications enter SQL statements while the application is running. While static SQL will suffice for most of your needs, dynamic SQL provides the flexibility to build diverse SQL statements at run time.

| Topic | Page |
|--|------|
| Dynamic SQL overview | 65 |
| Method 1: Using execute immediate | 67 |
| Method 2: Using prepare and execute | 69 |
| Method 3: Using prepare and fetch with a cursor | 72 |
| Method 4: Using prepare and fetch with dynamic descriptors | 75 |

Dynamic SQL is a set of Embedded SQL statements that permit users of online applications to access the database interactively at application run time.

Use dynamic SQL when one or more of the following conditions is not known until run time:

- SQL statement the user will execute
- Column, index, and table references
- Number of host variables, or their datatypes

Dynamic SQL overview

Dynamic SQL is part of ANSI and the ISO SQL2 standard. It is useful for running an interactive application. If the application only accepts a small set of SQL statements, you can embed them within the program. However, if the application accepts many types of SQL statements, you can benefit from constructing SQL statements, then binding and executing them dynamically.

The following situation would benefit from use of dynamic SQL: The application program searches a bookseller's database of books for sale. A potential buyer can apply many criteria, including price, subject matter, type of binding, number of pages, publication date, language, and so on.

A customer might say, "I want a nonfiction book about business that costs between \$10 and \$20." This request is readily expressed as a Transact-SQL statement:

```
select * from titles where
    type = "business"
    and price between $10 and $20
```

It is not possible to anticipate the combinations of criteria that all buyers will apply to their book searches. Therefore, without using dynamic SQL, an Embedded SQL program could not easily generate a list of prospective books with a single query.

With dynamic SQL, the bookseller can enter a **query** with a different where clause search condition for each buyer. The seller can vary requests based on the publication date, book category, and other data, and can vary the columns to be displayed.

For example:

```
select * from titles
    where type = ?
    and price between ? and ?
```

The question marks ("??") are dynamic parameter markers that represent places where the user can enter search values.

Dynamic SQL protocol

Note The precompiler does not generate stored procedures for dynamic SQL statements because the statements are not complete until run time. At run time, Adaptive Server stores them as temporary stored procedures in the tempdb database. The tempdb database must contain the user name "guest", which in turn must have create procedure permission. Otherwise, attempting to execute one of these temporary stored procedures generates the error message "Server user id *user_id* is not a valid user in database *database_name*", where *user_id* is the user's user ID and *database_name* is the name of the user's database.

The dynamic SQL prepare statement sends the actual SQL statement, which can be any Data Definition Language(DDL) or Data Manipulation Language (DML) statements, or any Transact-SQL statement except create procedure.

The dynamic SQL facility performs these actions:

- 1 Translates the input data into a SQL statement.
- 2 Verifies that the SQL statement can execute dynamically.
- 3 Prepares the SQL statement for execution, sending it to Adaptive Server, which compiles and saves it as a temporary stored procedure (for methods 2, 3, and 4).
- 4 Binds all input parameters or descriptor (for methods 2, 3, and 4).
- 5 Executes the statement.

For a varying-list select, it uses a descriptor to reference the data items and rows returned (for method 4).

- 6 Binds the output parameters or descriptor (for method 2, 3, and 4).
- 7 Obtains results.
- 8 Drops the statement (for method 2, 3, and 4) by deactivating the stored procedure in Adaptive Server.
- 9 Handles all error and warning conditions from Adaptive Server and Client-Library.

Method 1: Using *execute immediate*

Use *execute immediate* to send a complete Transact-SQL statement, stored in a host variable or literal string, to Adaptive Server. The statement cannot return any results—you cannot use this method to execute a select statement.

The dynamically entered statement executes as many times as the user invokes it during a session.

With this method:

- 1 The Embedded SQL program passes the text to Adaptive Server.
- 2 Adaptive Server verifies that the statement is not a select statement.
- 3 Adaptive Server compiles and executes the statement.

With execute immediate, you can let the user enter all or part of a Transact-SQL statement.

The syntax for execute immediate is:

```
exec sql [at connection_name] execute immediate
{:host_variable | string};
```

where:

- *host_variable* is a character-string variable defined in a declare section. Before calling execute immediate, the host variable should contain a complete and syntactically correct Transact-SQL statement.
- *string* is a literal Transact-SQL statement string that can be used in place of *host_variable*.

Embedded SQL sends the statement in *host_variable* or string to Adaptive Server without any processing or checking. If the statement attempts to return results or fails, an error occurs. You can test the value of SQLCODE

after executing the statement or use the whenever statement to set up an error handler. See Chapter 8 for information about handling errors in Embedded SQL programs.

Method 1: Examples

The following two examples demonstrate using method 1, execute immediate. The first example prompts the user to enter a statement and then executes it:

```
exec sql begin declare section;
CS_CHAR statement_buffer[linesize];
exec sql end declare section;
...
printf("\nEnter statement\n");
gets(statement_buffer);

exec sql [at connection] execute immediate
:statement_buffer;
```

The next example prompts the user to enter a search condition to specify rows in the titles table to update. Then, it concatenates the search condition to an update statement and sends the complete statement to Adaptive Server.

```
exec sql begin declare section;
CS_CHAR sqlstring[200];
exec sql end declare section;
```

```
char      cond[150];

exec sql whenever sqlerror call err_p();
exec sql whenever sqlwarning call warn_p();

strcpy(sqlstring,
"update titles set price=price*1.10 where ");

printf("Enter search condition:");
scanf("%s", cond);
strcat(sqlstring, cond);

exec sql execute immediate :sqlstring;

exec sql commit work;
```

Method 2: Using *prepare* and *execute*

Use method 2, prepare and execute, when one of the following cases is true:

- You are certain that no data will be retrieved and you want the statement to execute more than once.
- A select statement is to return a single row. With this method, you cannot associate a cursor with the select statement.

This process is also called a single-row select. If a user needs to retrieve multiple rows, use method 3 or 4.

This method uses prepare and execute to substitute data from C variables into a Transact-SQL statement before sending the statement to Adaptive Server. The Transact-SQL statement is stored in a character buffer with dynamic parameter markers to show where to substitute values from C variables.

Because this statement is prepared, Adaptive Server compiles and saves it as a temporary stored procedure. Then, the statement executes repeatedly, as needed, during the session.

The prepare statement associates the buffer with a statement name and prepares the statement for execution. The execute statement substitutes values from a list of C variables or SQL descriptors into the buffer and sends the completed statement to Adaptive Server. You can execute any Transact-SQL statement this way.

prepare

The syntax for the prepare statement is:

```
exec sql [at connection] prepare statement_name from
{:host_variable | string};
```

where:

- *statement_name* is a name up to 30 characters long that identifies the statement. It is a symbolic name or a C character array host variable containing the name of the statements that the precompiler uses to associate an execute statement with a prepare statement.
- *host_variable* is a character array host variable.

Precede the host variable with a colon, as in standard Embedded SQL statements.

- *string* is a literal string that can be used in place of *host_variable*.

host_variable or *string* can contain dynamic parameter markers (“?”), which indicate places in the dynamic query where values will be substituted when the statement executes.

execute

The syntax for the execute statement is:

```
exec sql [at connection] execute statement_name
[into host_var_list | sql descriptor
    descriptor_name | descriptor sqlda_name]
[using host_var_list | sql descriptor
    descriptor_name | descriptor sqlda_name];
```

where:

- *statement_name* is the name assigned in the prepare statement.
- *into* is used for a single-row select.
- *using* specifies the C variables or descriptors substituted for a dynamic parameter marker in *host_variable*. The variables, which you must define in a declare section, are substituted in the order listed. You need only this clause when the statement contains dynamic parameter markers.
- *descriptor_name* represents the area of memory that holds a description of the dynamic SQL statement’s dynamic parameter markers.

- *host_var_list* a list of host variables to substitute into the parameter markers (“?”) in the query.
- *sqlda_name* is the name of the SQLDA.

Method 2: Example

The following example demonstrates using prepare and execute in method 2. This example prompts the user to enter a where clause that determines which rows in the titles table to update and a multiplier to modify the price. According to what the user elects, the appropriate string is concatenated to the update statement stored in host variable “*sqlstring*”.

```
exec sql begin declare section;
    CS_CHAR sqlstring[200];
    CS_FLOAT multiplier;
exec sql end declare section;

char    cond[150];

exec sql whenever sqlerror perform err_p();
exec sql whenever sqlwarning perform warn_p();

printf("Enter search condition:");
scanf("%s", cond);
printf("Enter price multiplier: ");
scanf("%f", &multiplier);

strcpy(sqlstring,
    "update titles set price = price * ? where ");
strcat(sqlstring, cond);

exec sql prepare update_statement from :sqlstring;

exec sql execute update_statement using
    :multiplier;

exec sql commit;
}
```

Method 3: Using *prepare* and *fetch* with a cursor

Method 3 uses the *prepare* statement with cursor statements to return results from a *select* statement. Use this method for fixed-list *select* statements that may return multiple rows. That is, use it when the application has determined in advance the number and type of *select* column list attributes to be returned. You must anticipate and define host variables to accommodate the results.

When you use method 3, include the *declare*, *open*, *fetch*, and *close* cursor statements to execute the statement. This method is required because the statement returns more than one row. There is an association between the prepared statement identifier and the specified cursor name. You can also include *update* and *delete* where current of cursor statements.

As with method 2, *prepare* and *execute*, a Transact-SQL *select* statement is first stored in a character host variable or string. It can contain dynamic parameter markers to show where to substitute values from input variables. The statement is given a name to identify it in the *prepare*, *declare*, and *open* statements.

Method 3 requires five steps:

- 1 *prepare*
- 2 *declare*
- 3 *open*
- 4 *fetch* (and, optionally, *update* and *delete*)
- 5 *close*

These steps are described below.

prepare

The *prepare* statement is the same as that used with method 2. For details, see “*prepare*” on page 70.

declare

The *declare* statement is similar to the standard *declare* statement for cursors. In dynamic SQL, however, you *declare* the cursor for a prepared *statement_name* instead of for a *select* statement, and any input host variables are referenced in the *open* statement instead of in the *declare* statement.

A dynamic declare statement is an executable statement rather than a declaration. As such, it must be positioned in the code where executable statements are legal, and the application should check status codes (SQLCODE, SQLCA, or SQLSTATE) after executing the declaration.

The dynamic SQL syntax for the declare statement is:

```
exec sql [at connection_name] declare cursor_name
        cursor for statement_name;
```

where:

- *at connection_name* specifies the Adaptive Server connection the cursor will use
- *cursor_name* identifies the cursor, used with the open, fetch, and close statements
- *statement_name* is the name specified in the prepare statement, and represents the select statement to be executed

open

The open statement substitutes any input variables in the statement buffer, and sends the result to Adaptive Server for execution. The syntax for the open statement is:

```
exec sql [at connection_name] open cursor_name [using
        {host_var_list | sql descriptor descriptor_name |
        descriptor sqlda_name}];
```

where:

- *cursor_name* is the name given to the cursor in the declare statement
- *host_var_list* consists of the names of the host variables that contain the values for dynamic parameter markers
- *descriptor_name* is the name of the descriptor that contains the value for the dynamic parameter markers
- *sqlda_name* is the name of the SQLDA

fetch and close

After a cursor opens, the result sets are returned to the application. Then, the data is fetched and loaded into the application program host variables.

Optionally, you can update or delete the data. The fetch and close statements are the same as in static Embedded SQL.

The syntax for the fetch statement is:

```
exec sql [at connection_name] fetch cursor_name into
:host_variable [[indicator]:indicator_variable]
[, :host_variable
[[indicator]:indicator_variable]...];
```

where:

- *cursor_name* is the name given to the cursor in the declare statement.
- There is one *C host_variable* for each column in the result rows. The variables must have been defined in a declare section, and their datatypes must be compatible with the results returned by the cursor.

The syntax for the close statement is:

```
exec sql [at connection_name] close cursor_name;
```

where *cursor_name* is the name assigned to the cursor in the declare statement.

Method 3: Example

The following example uses prepare and fetch, and prompts the user for an order by clause in a select statement:

```
exec sql begin declare section;
CS_CHAR      sqlstring[200];
CS_FLOAT     bookprice,condprice;
CS_CHAR      booktitle[200];
exec sql end declare section;

char         orderby[150];

exec sql whenever sqlerror call err_p();
exec sql whenever sqlwarning call warn_p();

strcpy(sqlstring,
"select title,price from titles\
where price>? order by ");
```



```
printf("Enter the order by clause:");
scanf("%s", orderby);
strcat(sqlstring, orderby);

exec sql prepare select_state from :sqlstring;
exec sql declare select_cur cursor for select_state;

condprice = 10; /* the user can be prompted
                ** for this value */

exec sql open select_cur using :condprice;
exec sql whenever not found goto end;

for (;;)
{
    exec sql fetch select_cur
        into :booktitle,:bookprice;
    printf("%20s  %bookprice=%6.2f\n",
        booktitle, bookprice);
}

end:

exec sql close select_cur;
exec sql commit work;
```

Method 4: Using *prepare* and *fetch* with dynamic descriptors

Method 4 permits varying-list select statements. That is, when you write the application, you need not know the formats and number of items the select statement will return. Use method 4 when you cannot define the host variables in advance because you do not know how many variables are needed or of what type they should be.

Method 4: Dynamic descriptors.

A **dynamic descriptor** is a data structure that holds a description of the variables used in a dynamic SQL statement. There are two kinds of dynamic descriptors—SQL descriptors and SQLDA structures. Both are described later in this chapter.

When a cursor opens, it can have an input descriptor associated with it. The input descriptor contains the values to be substituted for the dynamic SQL statement's parameter markers.

Before the cursor is opened, the user fills in the input descriptor with the appropriate information, including the number of parameters, and, for each parameter, its type, length, precision, scale, indicator, and data.

Associated with the fetch statement is an output descriptor, which holds the resultant data. Adaptive Server fills in the data item's attributes, including its type and the actual data being returned. If you are using an SQL descriptor, use the get descriptor statement to copy the data into host variables.

Dynamic SQL method 4 performs the following steps:

- 1 Prepares the statement for execution.
- 2 Associates a cursor with the statement.
- 3 Defines and binds the input parameters or descriptor and:
 - If using an input descriptor, allocates it
 - If using an input parameter, associates it with the statement or cursor
- 4 Opens the cursor with the appropriate input parameter(s) or descriptor(s).
- 5 Allocates the output descriptor if different from the input descriptor and binds the output descriptor to the statement.
- 6 Retrieves the data by using fetch cursor and the output descriptor.
- 7 Copies data from the dynamic descriptor into host program variables. If you are using an SQLDA, this step does not apply; the data is copied in step 6.
- 8 Closes the cursor.
- 9 Deallocates the dynamic descriptor(s).
- 10 Drops the statement (ultimately, the stored procedure).

Dynamic descriptor statements

There are statements that associate the descriptor with a SQL statement and with a cursor associated with the SQL statement. The following list describes dynamic SQL statements for method 4:

| Statement | Description |
|-----------------------|---|
| allocate descriptor | Notifies Client-Library to allocate a SQL descriptor. |
| describe input | Obtains information about the dynamic parameter marker in the prepare statement. |
| set descriptor | Inserts or updates data in the system descriptor. |
| get descriptor | Moves row or parameter information stored in a descriptor into host variables, thereby allowing the application program to use the information. |
| execute | Executes a prepared statement. |
| open cursor | Associates a descriptor with a cursor and opens the cursor. |
| describe output | Obtains information about the select list columns in the prepared dynamic SQL statement. |
| fetch cursor | Retrieves a row of data for a dynamically declared cursor. |
| deallocate descriptor | Deallocates a dynamic descriptor. |

For complete descriptions of these statements, see Chapter 10, “Embedded SQL Statements: Reference Pages”

About SQL descriptors

A SQL descriptor is an area of memory that stores a description of the variables used in a prepared dynamic SQL statement. A SQL descriptor can contain the following information about data attributes (for details, see the descriptions of the set descriptor and get descriptor commands in Chapter 10, “Embedded SQL Statements: Reference Pages”):

- precision – integer.
- scale – integer.
- nullable – 1 (cs_true) if the column can contain nulls; 0 (cs_false) if it cannot. Valid only with get descriptor statement.
- indicator – value for the indicator associated with the dynamic parameter marker. Valid only with get descriptor statement.

- `name` – name of the dynamic parameter marker. Valid only with `get` descriptor statement.
- `data` – value for the dynamic parameter marker specified by the item number. If the value of `indicator` is -1, the value of `data` is undefined.
- `count` – number of dynamic parameter markers described in the descriptor.
- `type` – datatype of the dynamic parameter marker or host variable.
- `returned_length` – actual length of the data in an output column.

Method 4: Example Using SQL descriptors

The following example that uses prepare and fetch with dynamic parameter markers and SQL descriptors.

```
exec sql begin declare section
    int    index_colcnt, coltype;
    int    int_buff;
    char    char_buff[255], void_buff[255];
    char type[255], title[255];
    char colname[255];
    int sales;
    int descnt, occur, cnt;
    int condcnt, diag_cnt, num_msgs;
    char user_id[30], pass_id[30], server_name[30];
    char str1[1024], str2[1024], str3[1024],
        str4[1024];

exec sql end declare section;

...
void dyn_m4()
{
    printf("\n\nDynamic sql Method 4\n");
    printf("Enter in a Select statement to retrieve
        any kind of ");
    printf("information from the pubs database:");
    scanf("%s", &str4);

    printf("\nEnter the largest number of columns to
        be retrieved or the number ");
    printf("of ? in the sql statement:\n");
    scanf("%d", &occur);

    exec sql allocate descriptor dinout with max
```

```
        :occur;
    exec sql prepare s4 from :str4;
    exec sql declare c2 cursor for s4;

    exec sql describe input s4 using sql descriptor
        dinout;

    fill_descriptor();

    exec sql open c2 using sql descriptor dinout;

    while (sqlca.sqlcode == 0)
    {
        exec sql fetch c2 into sql descriptor dinout;
        if(sqlca.sqlcode == 0) {
            print_descriptor();
        }
    }

    exec sql close c2;

    exec sql deallocate descriptor dinout;
    exec sql deallocate prepare s4;
    printf("Dynamic SQL Method 4 completed\n\n");

}

void
print_descriptor()
{
    exec sql get descriptor dinout :descnt = count;
    printf("Column name \t\tColumn data\n");
    printf("----- \t\t-----\n");
    printf("-----\n");

    for (index_colcnt = 1; index_colcnt <= descnt;
        index_colcnt++)
    { /* get each column attribute */
        exec sql get descriptor dinout value
            :index_colcnt :coltype = TYPE;

        switch(coltype)
        {
            ...
            case 4:/* integer type */
                exec sql get descriptor dinout value
```

```
        :index_colcnt
        :colname = NAME, :int_buff = DATA;
    printf("%s \t\t %d\n", colname, int_buff);
    break;
    ...
}
}

void
fill_descriptor()
{
    exec sql get descriptor dinout :descnt = count;
    for (cnt = 1; cnt <= descnt; cnt++)
    {
        printf("Enter in the data type of the %d ?:",
            cnt);
        scanf("%d", &coltype);
        switch(coltype)
        {
            ...
            case 4:/* integer type */
                printf("Enter in the value of the data:");
                scanf("%d\n", &_buff);
                exec sql set descriptor dinout VALUE :cnt
                    TYPE = :coltype,
                    DATA = :int_buff;
                break;

            default:
                printf("non-supported column type.\n");
                break;
        }
    }
}
```

About SQLDAs

SQLDA is a host-language structure that, like an SQL descriptor, describes the variables used in a dynamic SQL prepared statement. Unlike SQL descriptors, SQLDAs are public data structures whose fields you can access. Statements using SQLDAs may execute faster than equivalent statements using SQL descriptors.

The SQLDA structure is not part of the SQL standard. Different implementations of Embedded SQL define the SQLDA structure differently. Embedded SQL version 11.1 and later supports the SQLDA defined by Sybase; it does not support SQLDA datatypes defined by other vendors.

To define the SQLDA datatype in your Embedded SQL program, you use the Embedded SQL command `include sqlda`. To allocate a SQLDA structure in your program, you use the `malloc` function. To deallocate an SQLDA, you use the `free` function. Your program is responsible for deallocating all SQLDA structures that it creates. Embedded SQL does not limit the number of SQLDA structures that can be created by a program.

Table 7-1 describes the fields of the SQLDA structure.

Table 7-1: Fields of the SQLDA structure

| Field | Datatype | Description |
|-------------------------------------|-------------|--|
| <code>sd_sqln</code> | CS_SMALLINT | The size of the <code>sd_column</code> array. |
| <code>sd_sqld</code> | CS_SMALLINT | The number of columns in the query being described, or 0 if the statement being described is not a query. For <code>fetch</code> , <code>open</code> , and <code>execute</code> statements, this field indicates the number of host variables described by occurrences of <code>sd_column</code> , or the number of dynamic parameter markers for the <code>describe</code> input statement. |
| <code>sd_column[].sd_datafmt</code> | CS_DATAFMT | Identifies the Client-Library CS_DATAFMT structure associated with this column. Refer to descriptions of <code>ct_bind</code> , <code>ct_param</code> , and <code>ct_describe</code> in the <i>Open Client Client-Library/C Reference Manual</i> for more information. |
| <code>sd_column[].sd_sqldata</code> | CS_VOID | For <code>fetch</code> , <code>open</code> , and <code>execute</code> statements, stores the address of the statement's host variable. This field is not used for <code>describe</code> or <code>prepare</code> statements. |
| <code>sd_column[].sd_sqlind</code> | CS_SMALLINT | For <code>fetch</code> , <code>open</code> , and <code>execute</code> statements, this field acts as an indicator variable for the column being described. If the column's value is null, this field is set to -1. This field is not used for <code>describe</code> or <code>prepare</code> statements. |
| <code>sd_column[].sd_sqlen</code> | CS_INT | The actual size of the data pointed to by <code>sd_sqldata</code> associated with this column. |

| Field | Datatype | Description |
|----------------------------------|----------|-------------|
| sd_column[<i>i</i>].sd_sqlmore | CS_VOID | Reserved. |

The Embedded SQL header file *sqlda.h* contains a macro, `SQLDADECL`, that lets you declare `SQLDA` structures in your program. The `SQLDADECL` macro is as follows:

```
#ifndef SQLDADECL
#define SQLDADECL(name, size)
    struct {
        CS_INT      sd_sqln;
        CS_INT      sd_sqln;
        struct {
            CS_DATAFMT      sd_datafmt;
            CS_VOID          sd_sqldata;
            CS_SMALLINT      sd_sqlind;
            CS_INT           sd_sqlllen;
            CS_VOID          sd_sqlmore;
        } sd_column[ (SIZE) ]
    } name
#endif /* SQLDADECL */
```

Method 4: Example using SQLDAs

Following is an example that uses prepare and fetch with dynamic parameter markers and SQL descriptors.

```
exec sql include sqlca;
exec sql include sqlda;

...
SQLDA *input_descriptor, *output_descriptor;
CS_SMALLINT small;
CS_CHAR      character[20];

input_descriptor = (SQLDA *)malloc(SYB_SQLDA_SIZE(3));
input_descriptor->sqlda_sqln = 3;
output_descriptor = (SQLDA *)malloc(SYB_SQLDA_SIZE(3));
output_descriptor->sqlda_sqln = 3;
*p_retcode = CS_SUCCEED;
exec sql connect "sa" identified by "";
/* setup */
exec sql drop table example;
exec sql create table example (fruit char(30), number int);
```



```
exec sql insert example values ('tangerine', 1);
exec sql insert example values ('pomegranate', 2);
exec sql insert example values ('banana', 3);
/* Prepare and describe the select statement */
exec sql prepare statement from
    "select fruit from example where number = ?";
exec sql describe input statement using descriptor    input_descriptor;
input_descriptor->sqlda_column[0].sqlda_datafmt.datatype =
CS_SMALLINT_TYPE;
input_descriptor->sqlda_column[0].sqlda_sqldata = &small;
input_descriptor->sqlda_column[0].sqlda_sqllen = sizeof(small);
small = 2;
exec sql describe output statement using descriptor
    output_descriptor;
if (output_descriptor->sqlda_sqld != 1 ||
    output_descriptor->sqlda_column[0].sqlda_datafmt.datatype !=
CS_CHAR_TYPE)
    FAIL;
else
    printf("First describe output \n");
output_descriptor->sqlda_column[0].sqlda_sqldata = character;
output_descriptor->sqlda_column[0].sqlda_datafmt.maxlength = 20;
exec sql execute statement into descriptor output_descriptor
    using descriptor input_descriptor;
printf("Expected pomegranate, got %s\n", character);
exec sql deallocate prepare statement;
/* Prepare and describe second select statement */
exec sql prepare statement from
    "select number from example where fruit = ?";
exec sql declare c cursor for statement;
exec sql describe input statement using descriptor
    input_descriptor;
input_descriptor->sqlda_column->sqlda_sqldata = character;
input_descriptor->sqlda_column->sqlda_datafmt.maxlength =      CS_NULLTERM;
strcpy(character, "banana");
input_descriptor->sqlda_column->sqlda_sqllen = CS_NULLTERM;
exec sql open c using descriptor input_descriptor;
exec sql describe output statement using descriptor
    output_descriptor;
output_descriptor->sqlda_column->sqlda_sqldata = character;
output_descriptor->sqlda_column->sqlda_datafmt.datatype =      CS_CHAR_TYPE;
output_descriptor->sqlda_column->sqlda_datafmt.maxlength = 20;
output_descriptor->sqlda_column->sqlda_sqllen = 20;
output_descriptor->sqlda_column->sqlda_datafmt.format =
    (CS_FMT_NULLTERM | CS_FMT_PADBLANK);
exec sql fetch c into descriptor output_descriptor;
```

```
printf("Expected pomegranate, got %s\n", character);  
exec sql commit work;
```

Summary

This chapter described dynamic SQL, a set of Embedded SQL statements that permit online applications to access the database interactively. This interaction with the database lets a user define and execute SQL statements at run time.

The four dynamic SQL Methods are:

- Method 1: execute immediate
- Method 2: prepare and execute
- Method 3: prepare and fetch
- Method 4: prepare and fetch with dynamic descriptors

The next chapter describes how to detect and correct Embedded SQL errors.

Handling Errors

This chapter discusses how to detect and correct errors that can occur during the execution of Embedded SQL programs. It covers the `whenever` and `get diagnostics` statements, which you can use to process warnings and errors, and the `SQLCA` variables that pertain to warnings and errors.

| Topic | Page |
|--|------|
| Testing for errors | 86 |
| Testing for warning conditions | 86 |
| Trapping errors with <code>whenever</code> | 87 |
| Using <code>get diagnostics</code> | 89 |
| Writing routines to handle warnings and errors | 90 |
| Precompiler-detected errors | 91 |

While an Embedded SQL application is running, some events may occur that interfere with the application's operation. Following are examples:

- Adaptive Server becomes inaccessible
- The user enters an incorrect password
- The user does not have access to a database object
- A database object is deleted
- A column's datatype changes
- A query returns an unexpected null value
- A dynamic SQL statement contains a syntax error

You can anticipate these events by writing warning and error-handling code to recover gracefully when one of these situations occurs.

Testing for errors

Embedded SQL places a return code in the *SQLCODE* variable to indicate the success or failure of each SQL statement sent to Adaptive Server. You can either test the value of *SQLCODE* after each Embedded SQL statement or use the *whenever* statement to instruct the precompiler to write the test code for you. The *whenever* statement is described later in this chapter.

Using SQLCODE

The following table lists the values *SQLCODE* can contain:

Table 8-1: *SQLCODE* return values

| Value | Meaning |
|-------|--------------------------------------|
| 0 | No warnings or errors |
| <0 | Error |
| 100 | No rows returned from last statement |

When *SQLCODE* is 0, no errors or warnings occurred.

When *SQLCODE* has a negative value, an error occurred. The *SQLCA* variables contain useful information for diagnosing the error.

A *SQLCODE* of 100 indicates that there are no result rows, although the statement executed successfully. This condition is useful for driving a loop that fetches rows from a cursor. When *SQLCODE* becomes 100, the loop and all rows that have been fetched end. This technique is illustrated in Chapter 6, “Using Transact-SQL Statements”

Testing for warning conditions

Even when *SQLCODE* indicates that a statement has executed successfully, a warning condition may still have occurred. The 8-character array *sqlca.sqlwarn* indicates such warning conditions. Each *sqlwarn* array element, or flag, stores either the space character or the character “W”.

The following table describes what the space character or “W” means in each flag:

Table 8-2: sqlwarn flags

| Flag | Description |
|------------|---|
| sqlwarn[0] | If blank, no warning condition of any kind occurred, and all other sqlwarn flags are blank. If sqlwarn[0] is set to “W”, one or more warning conditions occurred, and at least one other flag is set to “W”. |
| sqlwarn[1] | If set to “W”, the character string variable that you designated in a fetch statement was too short to store the statement’s result data, so the result data was truncated. You designated no indicator variable to receive the original length of the data that was truncated. |
| sqlwarn[2] | If set to “W”, the input sent to Adaptive Server contained a null value in an illegal context, such as in an expression or as an input value to a table that prohibits null values. |
| sqlwarn[3] | The number of columns in a select statement’s result set exceeds the number of host variables in the statement’s into clause. |
| sqlwarn[4] | Reserved. |
| sqlwarn[5] | SQL Server generated a conversion error while attempting to execute this statement. |
| sqlwarn[6] | Reserved. |
| sqlwarn[7] | Reserved. |

Test for a warning after you determine that a SQL statement executed successfully. Use the *whenever* statement, as described in the next section, to instruct the precompiler to write the test code for you.

Trapping errors with *whenever*

Use the Embedded SQL *whenever* statement to trap errors and warning conditions. It specifies actions to be taken depending on the outcome of each Embedded SQL statement sent to Adaptive Server.

The *whenever* statement is not executable. Instead, it directs the precompiler to generate C code that tests for specified conditions after each executable Embedded SQL statement in the program.

The syntax of the *whenever* statement is:

```
exec sql whenever {sqlwarning | sqlerror | not found}
    {continue | goto label |
```

```
call function_name ([param [, param]...]) | stop};
```

whenever testing conditions

Each whenever statement can test for one of the following three conditions:

- sqlwarning
- sqlerror
- not found

The precompiler generates warning messages if you do not write a whenever statement for each of the three conditions. If you write your own code to check for errors and warnings, suppress the precompiler warnings by writing a whenever...continue clause for each condition. This instructs the precompiler to ignore errors and warnings.

If you precompile with the verbose option, the precompiler generates a *ct_debug()* function call as part of each connect statement. This causes Client-Library to display informational, warning, and error messages to your screen as your application runs. The whenever statement does not disable these messages. For more information on the precompiler options, see the *Open Client/Server Programmer's Supplement*.

After an Embedded SQL statement executes, the values of *sqlcode* and *sqlwarn0* determine if one of the conditions exists. The following table shows the criteria whenever uses to detect the conditions:

Table 8-3: Criteria for the whenever statement

| Condition | Criteria |
|------------|--------------------------------|
| sqlwarning | sqlcode = 0 and sqlwarn[0] = W |
| sqlerror | sqlcode < 0 |
| not found | sqlcode = 100 |

To change the action of a whenever statement, write a new whenever statement for the same condition. whenever applies to all Embedded SQL statements that follow it, up to the next whenever statement for the same condition.

The whenever statement ignores the application program's logic. For example, if you place whenever at the end of a loop, it does not affect the preceding statements in subsequent passes through the loop.

whenever actions

The whenever statement specifies one of the following four actions:

Table 8-4: whenever actions

| Action | Description |
|----------|---|
| continue | Perform no special action when a SQL statement returns the specified condition. Normal processing continues. |
| goto | Perform a branch to an error handling procedure within your application program. You can write goto as either goto or go to, and you must follow it with a valid statement label name. The precompiler does not detect an error if the label name is not defined in the program, but the C compiler does. |
| call | Call another C routine and, optionally, pass variables. |
| stop | Terminate the program when a SQL statement triggers the specified condition. |

Using *get diagnostics*

The get diagnostics statement retrieves error, warning, and informational messages from Client-Library. It is similar to, but more powerful than, the whenever statement because you can expand it to retrieve more details of the detected errors.

If, within a whenever statement, you specify the application to go to or call another routine, specify get diagnostics in the function code, as follows:

```
void
error_handler()
{
    exec sql begin declare section;
        int num_msgs;
        int condcnt;
    exec sql include sqlca;
    exec sql end declare section;
    exec sql get diagnostics :num_msgs = number;
    for (condcnt=1; condcnt <= num_msgs; condcnt++)
    {
        exec sql get diagnostics exception :condcnt
            :sqlca = sqlca_info;
        printf("sqlcode is :%d\n\ message text:
            %s\n", sqlca.sqlcode,
            sqlca.sqlerrm.sqlerrmc);
    }
}
```

```
    }  
}
```

Writing routines to handle warnings and errors

A good strategy for handling errors and warnings in an Embedded SQL application is to write custom procedures to handle them, then install the procedures with the `whenever...call` statement.

The following example shows sample warning and error handling routines. For simplicity, both routines omit certain conditions that should normally be included. `warning_hndl` omits the code for `sqlwarn[1]`. `error_hndl` omits the code that handles Client-Library errors and operating system errors:

```
/* Declare the sqlca. */  
exec sql include sqlca;  
exec sql whenever sqlerror call error_handler();  
exec sql whenever sqlwarning call  
warning_handler();  
exec sql whenever not found continue;  
/*  
** void error_handler()  
**  
** Displays error codes and numbers from the sqlca  
**/  
void error_handler()  
{  
    fprintf(stderr,  
        "\n**sqlcode=(%d)", sqlca.sqlcode);  
  
/*  
** void warning_handler()  
**  
** Displays warning messages.  
**/  
void warning_handler()  
{  
  
    if (sqlca.sqlwarn[1] == 'W')  
    {  
        fprintf(stderr, "\n** Data truncated.\n");  
    }  
}
```



```
if (sqlca.sqlwarn[3] == 'W')
{
    fprintf(stderr, "\n** Insufficient
        host variables to store results.\n");
}
return;
}
```

Precompiler-detected errors

The Embedded SQL precompiler detects Embedded SQL errors at precompile time. The precompiler detects syntax errors such as missing semicolons and undeclared host variables in SQL statements. These are severe errors, so appropriate error messages are generated.

You can also have the precompiler check Transact-SQL syntax errors. Adaptive Server parses Transact-SQL statements at precompile time if the appropriate precompiler command options are set. See the precompiler reference page in the *Open Client/Server Programmer's Supplement* for your platform.

The precompiler does not detect the error in the following example, in which a table is created and data is selected from it. The error is that the host variables' datatypes do not match the columns retrieved. The precompiler does not detect the error because the table does not yet exist when the precompiler parses the statements:

```
exec sql begin declare section;
    CS_INT    var1;
    CS_CHAR   var2[20];
exec sql end declare section;

exec sql create table
    T1 (col1 int, col2 varchar(20));
....

exec sql select * from T1 into :var2, :var1;
```

Note that the error will be detected and reported at run time.

Improving Performance with Persistent Binding

Persistent binding is a feature of Client-Library, the set of routines that executes Embedded SQL statements. Persistent binding improves a program's performance by enabling the Embedded SQL precompiler to create more efficient code.

| Topic | Page |
|--|------|
| About persistent binding | 94 |
| Precompiler options for persistent binding | 97 |
| Overview of rules for persistent binding | 98 |
| Guidelines for using persistent binding | 105 |

Persistent binding is optional: it takes effect if you request it when you precompile your program. Persistent binding benefits only certain types of Embedded SQL programs.

To understand this chapter, you should be familiar with host variables, cursors, dynamic SQL, and precompiler options. Refer to:

- Chapter 4, “Using Variables” for information about host variables.
- Chapter 6, “Using Transact-SQL Statements” for information about cursors.
- Chapter 7, “Using Dynamic SQL” for information about dynamic SQL.
- The *Open Client/Server Programmer's Supplement* for information about precompiler options and about starting the precompiler.

You need not understand Client-Library to use persistent binding in Embedded SQL. However, understanding Client-Library's command structures, `ct_bind` routine, and `ct_fetch` routine can help you understand why persistent binding works as it does in Embedded SQL.

The general function of command structures, `ct_bind`, and `ct_fetch`

are described briefly in this chapter. For complete descriptions, refer to the *Open Client Client-Library/C Programmer's Guide* and the *Open Client Client-Library/C Reference Manual*.

About persistent binding

To pass values to SQL Server and to store values from it, an Embedded SQL program uses host variables—C variables recognized by Embedded SQL. The program associates these variables with values on SQL Server. For example, the following select statement associates the host output variable `last` with a row value retrieved from SQL Server:

```
id = "998-72-3567";
exec sql select au_lname into :last
from authors where au_id = :id;
```

The statement passes its host input variable, `id`, to SQL Server and associates that variable with the server's `au_id` column.

The act of associating a statement's host variables with SQL Server values is called **binding**. The association itself is also called a binding. Host input variables use only input bindings; host output variable use only output bindings.

Binding governs which data a statement retrieves from the server. If a statement binds a host variable to the wrong server data, the statement will retrieve the wrong value for that host variable. However, unnecessary binding can slow a program's performance.

Embedded SQL lets you control how long bindings remain in effect—how long they “persist.” A binding that persists for more than one execution of a statement is called a **persistent binding**. Persistent bindings enable some Embedded SQL statements to execute faster, thereby improving a program's performance.

In Embedded SQL, each binding is made possible by a Client-Library **command structure**—a data structure that, among other things, defines the bindings of an Embedded SQL statement. For each Embedded SQL statement that executes, there is a corresponding command structure. A single command structure, however, can be used by more than one statement. In fact, when bindings persist from one Embedded SQL statement to another, they do so because the statements share a single command structure.

An Embedded SQL program's source code does not explicitly declare or allocate command structures. Rather, command structures are declared and allocated by the program's generated code.

When binding occurs

By default, binding occurs each time an Embedded SQL statement executes, using a host variable. When an Embedded SQL statement executes more than once, as in a loop, binding occurs at each execution. For example, in the following loop, each execution of the `insert` statement associates its host variables with the same SQL Server values. Yet, by default, binding occurs for each execution:

```
for (i = 1; i <= 3; i++)
{
    exec sql insert into titles (title_id, title)
    values (:bk_id, :bk_title);
    /*
     ** Binding occurs here at each execution.
     ** When a statement undergoes binding, all
     ** its host variables get bound.
     */
}
```

For most statements, bindings do not persist from one statement to the next, even if you request persistent binding. For example, the following `insert` statements, though identical and consecutive, share no bindings:

```
exec sql insert into titles (title_id, title)
    values (:bk_id, :bk_title);
/* Binding occurs for the first statement. */

exec sql insert into titles (title_id, title)
    values (:bk_id, :bk_title);
/* Binding occurs for the second statement. */

exec sql insert into titles (title_id, title)
    values (:bk_id, :bk_title);
/* Binding occurs for the third statement. */
```

For Embedded SQL statements that execute more than once—such as the `insert` statement in the preceding `for` loop—you can specify whether binding should occur only at the first execution or at each subsequent execution as well.

To control persistent binding, you use precompiler options to specify the binding behavior of all the statements in a file. Precompiler options do not let you control the binding behavior of individual statements. The precompiler options that control binding are explained later in this chapter.

Programs that can benefit from persistent binding

Not all Embedded SQL programs benefit from persistent binding. To find out whether persistent binding can benefit your program, answer the following questions:

- 1 Does your program contain at least one Embedded SQL statement that executes more than once?
- 2 If so, does that statement repeatedly use the same host variables to exchange values with SQL Server?

If you answered “yes” to both questions, your program can probably benefit from persistent binding. If you answered “no” to either question, persistent binding would not improve your program’s performance—unless you modify your program so that you can answer “yes” to both questions.

To maximize the benefit from persistent binding, your program should execute a single Embedded SQL statement repeatedly instead of executing two or more identical statements. For example, the following insert statement executes repeatedly:

```
for (i = 1; i <= 3; i++)
{
    exec sql insert into titles (title_id, title)
    values (:bk_id, :bk_title);
}
```

Although the insert statement in this example executes three times, its variables are bound only once. Because binding is not repeated, this example should run faster than a series of identical insert statements that execute only once.

Scope of persistent bindings

The scope of persistent bindings—how long they persist—differs depending on the type of statement and on the precompiler options in effect, as described later in this chapter. However, *bindings never persist beyond the lifetime of a connection*. When a program closes a connection, all bindings for statements issued and all command structures allocated over that connection are canceled.

Precompiler options for persistent binding

Two precompiler options control binding, the `-p` option and the `-b` option. These options affect only Embedded SQL statements that can use persistent binding. (Refer to Table 9-1 for a list of statements that cannot use persistent binding.)

Note The names of the `-p` and `-b` options differ on some platforms. To find out the name of these options on your platform, refer to the *Open Client/Server Programmer's Supplement*. Also, the `-b` option of Embedded SQL version 11.1 and later differs slightly from that for previous releases. Refer to your *Open Client/Server Products Release Bulletin* for more information.

About the `-p` option

The `-p` option controls whether each statement has a **persistent command structure**—one that persists for all executions of a particular statement. Only statements with a persistent command structure can have persistent bindings for host variables. Thus, the

`-p` option controls binding of host input variables, whose values are passed to Adaptive Server. (In this chapter, information about “host input variables” also applies to other variables whose values are passed to SQL Server. Exceptions are noted in the text.

About the **-b** option

The **-b** option controls binding of host variables used in statements that retrieve result data from Adaptive Server. When used in conjunction with the **-p** option, it controls binding of host variables in **select** and **exec** statements. When the **-b** option is used by itself, it can only control statements that fetch with a cursor.

Thus, generally, the **-b** option controls binding of host variables—output variables, result variables, status variables, indicator variables, and so on—whose values are passed from Adaptive Server. (Information about “host output variables” also applies to any other variables whose values are output from SQL Server.) More precisely, the **-b** option controls whether binding occurs at each execution of Client-Library’s `ct_fetch` routine. (The `ct_fetch` routine retrieves a single row of data from SQL Server.

Which option to use: **-p**, **-b**, or both

Most programs that can benefit from persistent bindings for input variables can also benefit from persistent bindings for output variables. In general, you should use both **-p** and **-b** options or use neither option.

Scope of the **-p** and **-b** precompiler options

The **-p** and **-b** options affect only the file being precompiled, unless that file declares a cursor. If the file declares a cursor, **-p** and **-b** affect all statements that use the cursor—even if those statements are in different source files of your program. The effect of **-p** and **-b** on files that use cursors is described in detail later in this chapter.

Overview of rules for persistent binding

The rules of persistent binding differ for different types of Embedded SQL statements. Specifically, the rules differ depending on whether a statement:

- Can use persistent binding
- Uses a cursor
- Is a dynamic SQL statement

- Is a fetch statement with the rebind/norebind clause

Statements that can use persistent binding

Most Embedded SQL statements can use persistent binding. A few, however, cannot. Table 9-1 and Table 9-2 list Embedded SQL commands that cannot use persistent binding. All other Embedded SQL commands—including Transact-SQL commands—can use persistent binding for some or all host variables.

Whether a statement's bindings persist and how long they persist depends on the type of statement—particularly, on whether the statement uses a cursor.

Table 9-1: Embedded SQL commands that cannot use persistent binding

| | |
|-----------------------------|--------------------|
| allocate descriptor | begin transaction |
| close | commit |
| connect | deallocate cursor |
| deallocate descriptor | deallocate prepare |
| describe input | describe output |
| disconnect | end transaction |
| execute | execute immediate |
| get descriptor | get diagnostics |
| open using descriptor | prepare |
| prepare transaction | rollback |
| set descriptor | set connection |
| set transaction diagnostics | |

Table 9-2: Types of Embedded SQL commands that cannot use persistent binding

| | |
|--|--|
| Commands that send text or image data to SQL Server with the -y option | Dynamic SQL commands that use a SQL descriptor or SQLDA for input to Adaptive Server |
|--|--|

Persistent binding in statements without a cursor

If an Embedded SQL statement can use persistent binding but does not use a cursor, you control the statement's bindings with the -p and -b options when precompiling the statement. Table 9-3 describes how these options affect a statement that uses no cursor.

Table 9-3: How the `-p` and `-b` options affect statements that use no cursors

| Options used to precompile statement | Effect on statement's bindings |
|---|--------------------------------|
| Neither <code>-p</code> nor <code>-b</code> | No bindings persist. |
| <code>-p</code> only | Only input bindings persist. |
| <code>-b</code> only | No bindings persist. |
| Both <code>-p</code> and <code>-b</code> | All bindings persist. |

If the statement's bindings persist, they do so until your program closes the connection over which the statement executes. The bindings persist throughout all executions of the statement, even if other statements execute in the meantime. If the statement's bindings do not persist, binding occurs each time the statement executes.

Persistent binding in statements with a cursor

Before your program can use a cursor, you must declare it with the `declare cursor` command. A cursor's declaration governs the binding behavior of all statements that use the cursor—in all source files of your program. The reason for this control is that the command structure for a cursor's declaration is shared by all statements that use the cursor.

When a statement uses a cursor, the cursor's declaration—not the statement using the cursor—controls how long the statement's bindings persist. The bindings persist only if you use the `-b` and `-p` options when precompiling the file that declares the cursor. If you use these options, all statements that use the cursor have persistent bindings as specified by the options.

Strictly speaking, a cursor's declaration controls binding behavior only if the cursor is a **dynamic cursor**—a cursor for a dynamic SQL statement. In cursors for all other SQL statements (**static cursors**), the statement that most recently opened the cursor (open cursor) controls the binding behavior, not the statement that declares the cursor.

Note For a static cursor, the generated code for `open cursor` both declares and opens the cursor. For a dynamic cursor, the generated code for `open cursor` only opens the cursor.

Except for this difference, the binding rules for static cursors and dynamic cursors are the same. Unless you use a particular cursor in more than one source file of your program, the binding behavior of static cursors and dynamic cursors is the same.

In statements that use a cursor, bindings never persist after the cursor is deallocated, even if you use persistent binding. Also, deallocated cursors cannot be reopened. Declaring a new cursor with the name of a deallocated cursor does not reopen the deallocated cursor, nor does it retain bindings associated with that cursor. For more information, refer to the description of the `deallocate cursor` command in Chapter 10, “Embedded SQL Statements: Reference Pages”

The following example shows how the `-b` and `-p` options affect a cursor—in this example, *curs1*. The fetch statement in the example contains host variables. The paragraphs following the example describes how the `-b` and `-p` options affect the bindings of these host variables.

```
#include <stdio.h>
int SQLCODE;

void
main()
{
    exec sql begin declare section;
        char title[100], pub_id[8];
    exec sql end declare section;

    exec sql connect "sa";

    exec sql use pubs2;
    /*
    ** The options used to precompile a cursor's declaration
    ** control whether host variables persist in statements,
    ** such as FETCH, that use the cursor.
    */
    exec sql declare curs1 cursor for select title, pub_id from
        titles;
    exec sql open curs1;

    while (SQLCODE == 0)
    {
        /* If the declaration of curs1 was precompiled without
        ** the -b option, rebind the FETCH statement's variables
        ** each time the statement repeats. Otherwise, bind only
        ** the first time, and let the bindings persist for
```

```
** subsequent repetitions.
*/
exec sql fetch curs1 into :title, :pub_id;
    printf("%s, %s\n", title, pub_id);
}
/* If the declaration of curs1 was precompiled without
** the -p option, cancel the bindings of the FETCH
** statement's variables when curs1 is closed.
** Otherwise, let the bindings persist until the
** program deallocates curs1 or, as here, until the
** program ends.
*/
exec sql close curs1;
exec sql disconnect CURRENT;

exit(0);
```

Preventing persistent binding for all cursor host variables

If you omit both the `-b` and `-p` options when precompiling the above example, no bindings persist. Instead, the generated code binds the host variables each time the fetch statement executes, —regardless of whether the variable is input to Adaptive Server or output from it.

Requesting persistent binding for all cursor host variables

If you use both the `-b` and `-p` options when precompiling the preceding example, the generated code binds the host variables of the fetch statement only the first time that the statement executes. Unlike other Embedded SQL statements (as described in “When binding occurs” on page 95), it does not matter whether there are one or more identical fetch statements in a series, or a simple fetch statement executed in a loop. Using both options together causes the bindings to persist even when the program closes the cursor; the host variables do not need to be rebound when the cursor is reopened. The bindings persist until the program deallocates the cursor—typically, with the `deallocate cursor` or `disconnect` statement.

Requesting persistent binding for cursor output variables only

If you use `-b` but omit `-p` when precompiling the preceding example, the generated code binds the host output variables of the fetch statement only once—the first time that the statement executes. (More precisely, the host variables get bound only if they are not bound already.) The bindings persist for all subsequent executions of the statement, until the program closes *cursor1*. They persist because you used the `-b` option. Because you omitted the `-p` option, bindings for host input variables do not persist.

If your program closes *cursor1* and then reopens it, all bindings for host variables related to *cursor1* are canceled. Any host input variables and host output variables are re-bound when the cursor is reopened. They then persist until your program closes the cursor again.

Requesting persistent binding for cursor input variables only

The preceding example showed how the `-b` and `-p` options affect statements that use host variables with a cursor. The example's only host variables were host output variables. The following code is an example that shows how the `-b` and `-p` options affect statements that use host input variables with a cursor—here, a dynamic cursor named *dyn_cursor1*.

The open statement in the following example contains a host input variable, *min_price*. The following sections describe how the `-b` and `-p` options affect the bindings of this host input variable.

```
#include <stdio.h>
long SQLCODE = 0;

void main()
{
    int i = 0;
    exec sql begin declare section;
        CS_CHAR          sql_string[200];
        CS_FLOAT          min_price;
        CS_CHAR          book_title[200];
    exec sql end declare section;
    exec sql connect "sa";
    exec sql use pubs2;
    strcpy(sql_string,
        "select title from titles where price > ?");
    exec sql prepare sel_stmt from :sql_string;

    /* The options used to precompile a cursor's declaration
    ** control whether host variables persist in statements,
```

```
** such as OPEN, that use the cursor.
*/
exec sql declare dyn_curs1 cursor for sel_stmt;
min_price = 10.00;
/* If the declaration of dyn_curs1 was precompiled
** without -p, bind the OPEN statement's input variable
** (min_price) each time the statement repeats. Otherwise,
** bind only the first time, letting the binding persist
** until dyn_curs1 is deallocated.
*/

for (i = 10; i <= 21; ++i)
{
    min_price = min_price + 1.00;
    exec sql open dyn_curs1 using :min_price;
    while (SQLCODE != 100)
    {
        exec sql fetch dyn_curs1 into :book_title;
        if (SQLCODE != 100) printf("%s\n", book_title);
    }
    printf("\n");
    exec sql close dyn_curs1;
}
exec sql deallocate cursor dyn_curs1;
exec sql disconnect CURRENT;
exit(0);
}
```

If you use `-p` but omit `-b` when precompiling the preceding example, the generated code binds *min_price* only once—the first time that the open statement executes. The binding persists because you used the `-p` option, which controls host input variables.

The binding for *min_price* persists throughout all subsequent iterations of the statement, until the program deallocates *dyn_curs1*. The binding persists even if your program closes *dyn_curs1* and then reopens it.

Persistent binding, cursors, and multiple source files

In the preceding example, the declaration of the cursor *dyn_curs1* controls whether associated host variables persist. For this reason, the host variables in the fetch statement would bind as described in the example, even if the fetch statement were precompiled in a separate source file.

Persistent binding and cursor *fetch* statements

The Embedded SQL *fetch* command has an optional *rebind/norebind* clause that controls whether bindings persist in a particular *fetch* statement. This clause is useful if you need to override the precompiler options that you specified for a file. The *rebind/norebind* clause affects only the statement in which it appears. Bindings for other statements—including other *fetch* statements—are not affected.

If a *fetch* statement omits the *rebind/norebind* clause, the statement obeys the same binding rules as do other types of statements that use the cursor in question.

If a *fetch* statement contains the keyword *rebind*, bindings for host variables in the statement do not persist. Instead, they get rebound each time the statement executes—regardless of whether the *-b* option was used to precompile the declaration of the statement's cursor.

If a *fetch* statement contains the keyword *norebind* but is precompiled with the *-b* option, the keyword has no effect.

Guidelines for using persistent binding

Here are guidelines, tips, and reminders to help you use persistent binding correctly:

- A program benefits from persistent binding only if it meets both these criteria:
 - It contains at least one Embedded SQL statement that executes more than once, and
 - That statement uses the same host variables repeatedly to exchange values with SQL Server.
- The *-p* and *-b* options affect only the file being precompiled, unless that file declares a cursor. If the file declares a cursor, *-p* and *-b* affect all statements that use the cursor. In general, you should use both the *-p* and *-b* options or use neither. If your program consists of more than one Embedded SQL source file, you should generally use the same combination of the *-p* and *-b* options to precompile all the files.

Generally, if you use the same cursor in more than one source file of a program, use the same combination of the -p and -b options when precompiling those files. Otherwise, you will need to understand exactly how different combinations of the options can change which data a statement sends or retrieves.

- A program that uses persistent binding should, where practical, execute a single Embedded SQL statement repeatedly instead of executing two or more identical statements once each.
- The rules controlling a statement's bindings differ depending on whether the statement:
 - Can use persistent binding
 - Uses a cursor
 - Is a dynamic SQL statement
 - Is a fetch statement with the rebind/norebind clause
- Bindings never persist beyond the lifetime of a connection. In statements that use a cursor, bindings never persist after the cursor is deallocated.
- A dynamic cursor's declaration controls the binding behavior of all statements that use the cursor. For a static cursor, the statement that most recently opened the cursor exerts this control. A program should open a static cursor only in the source file that declares it.

Notes on the binding of host variables

Subscripted arrays

If you use -p or -b and bind a subscripted array host variable (input or output), the subscript is ignored after the first execution of the statement, because the actual address of the specified array element is bound. For example:

```
exec sql begin declare section;  
int row;  
int int_table[3] = {  
    10,
```



```
        20,
        30,
    };

    char *string_table[3] = {
        "how",
        "are",
        "you",
    };

    exec sql end declare section;

    for (row=0; row < 3; row++)
    {
        EXEC SQL insert into ... values (:row, :int_table[row],
                                         :string_table[row]);

        /*
        ** If this statement is precompiled with -p, only
        ** int_table[0] and string_table[0] will be bound and
        ** inserted each time.
        ** The same thing applies to output variables
        ** At this time, NO warnings are issued to detect this.
        */
    }
```

To solve this, you can choose among the following solutions:

- Do not use persistent binds when subscripted arrays are used, since you *do* want a rebind (*table[0] is not the same as *table[1] at the next iteration).
- If persistent binds must be used, use an intermediate variable that holds the current value. This method allows persistent binding without errors. However, copying the data creates overhead. Using the above example:

```
exec sql begin declare section;

char    bind_str[80];

int bind_int_variable;

exec sql end declare section;

for (row=0; row < 3; row++)
{
```

```
/*
** Must copy the contents- pointer assignment does
** not suffice host var 'row' is not a subscripted
** array, so it can remain the same.
*/

memcpy(bind_str, string_table[row],80);

bind_int_variable = int_table[row];

EXEC SQL insert into ... values (:row,
                                :bind_int_variable,

                                :bind_str);
}
```

Note No register variables can be used with persistent binding.

Scope of host variables

When host variables remain bound from one execution to the next, you must ensure that they remain in scope. Particular care must be taken when automatic variables such as stack variables are used.

When a possibly problematic situation can be detected by the precompiler, a warning is issued. Whether a host variable remains in scope or not will also depend on the overall program logic.

For example:

```
/*
** a function called by main()
*/

CS_VOID insert(insert_row)

exec sql begin declare section;

int insert_row; /* row will go out of scope once exit
                ** function*/

exec sql end declare section;

{

/*
** id is a stack variable which will go out of scope
** once we exit the function insert()
*/
```

```
    /** it is not likely to be at the same address at the
    ** next call to this function, so if it is bound as
    ** an input variable, there will be errors.
    */

    exec sql begin declare section;

    int id;

    exec sql end declare section;

    exec sql insert values(:row,:id);

}

int fetched_row; /* this variable can be safely bound with
                  ** persistence */

main()
{
    exec sql begin declare section;

    /*
    ** This variable will go out of scope when the program
    ** exits main, which is not a problem.
    */

    int row;

    /*
    ** This variable is a pointer, thus it does not
    ** necessarily pose problems, depending on the scope
    ** of the data it is pointing to.
    */

    char *pointer;

    exec sql end declare section;

    for (row = 0; row < 10; row++)
    {
        insert(row);
    }
}
```


Embedded SQL Statements: Reference Pages

This chapter consists of a reference page for each Embedded SQL statement that either does not exist in Transact-SQL, or works differently from how it does in Transact-SQL. Refer to the *Transact-SQL User's Guide* for descriptions of all other Transact-SQL statements that are valid in Embedded SQL.

| Command statements | Page |
|-----------------------------------|------|
| allocate descriptor | 112 |
| begin declare section | 114 |
| begin transaction | 115 |
| close | 116 |
| commit | 118 |
| connect | 119 |
| deallocate cursor | 121 |
| deallocate descriptor | 123 |
| deallocate prepare | 124 |
| declare cursor (dynamic) | 125 |
| declare cursor (static) | 126 |
| declare cursor (stored procedure) | 128 |
| delete (positioned cursor) | 129 |
| delete (searched) | 131 |
| describe input (SQL descriptor) | 133 |
| describe input (SQLDA) | 134 |
| describe output (SQL descriptor) | 135 |
| describe output (SQLDA) | 136 |
| disconnect | 137 |
| exec | 139 |
| exec sql | 142 |
| execute | 144 |
| execute immediate | 146 |
| exit | 147 |
| fetch | 147 |

| Command statements | Page |
|---------------------------|-------------|
| get descriptor | 150 |
| get diagnostics | 152 |
| include "filename" | 153 |
| include sqlca | 155 |
| include sqlda | 156 |
| initialize_application | 157 |
| open (dynamic cursor) | 158 |
| open (static cursor) | 159 |
| prepare | 161 |
| rollback | 163 |
| select | 164 |
| set connection | 165 |
| set descriptor | 166 |
| update | 167 |
| whenever | 169 |

Except for print, readtext and writetext, all Transact-SQL statements can be used in Embedded SQL, though the syntax of some statements differs as described in this chapter.

The reference pages in this chapter are arranged alphabetically. Each statement's reference page:

- Briefly states what the statement does
- Describes the statement's syntax
- Explains the statement's keywords and options
- Comments on the statement's proper use
- Lists related statements, if any
- Demonstrates the statement's use in a brief example

allocate descriptor

Description

Allocates a SQL descriptor.

Syntax

```
exec sql allocate descriptor descriptor_name  
[with max [host_variable | integer_literal]];
```

Parameters

descriptor_name

The name of the SQL descriptor that will contain information about the dynamic parameter markers in a prepared statement.

with max

The maximum number of columns in the SQL descriptor.

host_variable

An integer host variable defined in a declare section.

integer_literal

A numeric value representing the size, in number of occurrences, of the SQL descriptor.

Examples

```
exec sql begin declare section;
  CS_INT      type;
  CS_INT      numcols, colnum;
exec sql end declare section;
...
exec sql allocate descriptor big_desc
  with max 1000;
exec sql prepare dynstmt from "select * from
  huge_table";
exec sql execute dynstmt into sql descriptor
  big_desc;
exec sql get descriptor :numcols = count;
for (colnum = 1; colnum <= numcols; colnum++)
{
  exec sql get descriptor big_desc :type = type;
  ...
}
exec sql deallocate descriptor big_desc;
...
```

Usage

- The allocate descriptor command specifies the number of item descriptor areas that Adaptive Server allocates.
- You can allocate any number of SQL descriptors.
- When a SQL descriptor is allocated, its fields are undefined.
- If you try to allocate a SQL descriptor that is already allocated, an error occurs.
- If you do not specify a value for the with max clause, one item descriptor is assigned.
- When a SQL descriptor is allocated, the value of each of its fields is undefined.

See also `deallocate descriptor`, `get descriptor`, `set descriptor`

begin declare section

| | |
|-------------|---|
| Description | Begins a declare section, which declares host language variables used in an Embedded SQL source file. |
| Syntax | <pre>exec sql begin declare section; host_variable_declaration; ... exec sql end declare section;</pre> |
| Parameters | <p><i>host_variable_declaration</i></p> <p>The declaration of one or more host language variables.</p> |
| Examples | <pre>exec sql begin declare section; CS_CHAR name(80); CS_INT value; exec sql end declare section;</pre> |
| Usage | <ul style="list-style-type: none">• A declare section must end with the Embedded SQL statement <code>end declare section</code>.• A source file can have any number of declare sections.• declare sections can be placed anywhere that variables can be declared. The declare section that declares a variable must precede any statement that references the variable.• Variable declarations in a declare section must conform to the rules of the host language.• Nested structures are valid in a declare section; arrays of structures are not.• A declare section can contain any number of Embedded SQL include statements.• In Embedded SQL/C routines, the Client-Library datatypes defined in <i>cspubli.c</i> can be used in declare sections.• In C routines, you can declare two-dimensional arrays of characters but only one-dimensional arrays of other datatypes. |

- When processing declare sections, the Embedded SQL precompiler ignores C preprocessor macros and #include statements. When processing Embedded SQL include statements within a declare section, the Embedded SQL precompiler treats the contents of the included file as though had been entered directly into the file being precompiled.

See also `exec sql include "filename"`

begin transaction

Description Marks the starting point of an unchained transaction.

Syntax `exec sql [at connection_name]
begin {transaction | tran} [transaction_name];`

Parameters `transaction | tran`
The keywords `transaction` and `tran` are interchangeable.

transaction_name

The name that you are assigning to this transaction. The name must conform to the rules for Transact-SQL identifiers.

Examples

```
/*
** Use explicit transactions to
** synchronize tables on two servers
*/
exec sql begin declare section;
    char        title_id[7];
    int         num_sold;
exec sql end declare section;
    long        sqlcode;
    ...

exec sql whenever sqlerror goto abort_tran;
try_update:

exec sql at connect1 begin transaction;

exec sql at connect2 begin transaction;
exec sql at connect1 select sum(qty)
    into :num_sold
    from salesdetail
    where title_id = :title_id;
exec sql at connect2 update current_sales
    set num_sold = :num_sold
```

```
        where title_id = :title_id;
exec sql at connect2 commit transaction;
exec sql at connect1 commit transaction;
if (sqlcode != 0)
    printf("oops, should have used 2-phase
    commit\n");
return;
abort_tran:
exec sql whenever sqlerror continue:
exec sql at connect2 rollback transaction;
exec sql at connect1 rollback transaction;
goto try_update;
```

Usage

- This reference page describes aspects of the Transact-SQL `begin transaction` statement that differ when used with Embedded SQL. See the *Adaptive Server Enterprise Reference Manual* for more information about `begin transaction` and Transact-SQL transaction management.
- The `begin transaction` statement is valid only in unchained transaction mode. In chained transaction mode, you cannot explicitly mark the starting point of a transaction.
- When nesting transactions, assign a transaction name only to the outermost `begin transaction` statement and its corresponding `commit transaction` or `rollback transaction` statement.
- Unless you set the database option `ddl in tran`, Adaptive Server does not allow the following statements inside an unchained transaction: `create database`, `create table`, `create index`, `create view`, `drop`, `select into table_name`, `grant`, `revoke`, `alter database`, `alter table`, `truncate table`, `update statistics`, `reconfigure`, `load database`, `load transaction`, and `disk init`.
- A transaction includes only statements that execute on the connection that is current when the transaction begins.
- Remote procedures execute independently of any transaction in which they are included.

See also

`commit transaction`, `commit work`, `rollback transaction`, `rollback work`

close

Description

Closes an open cursor.

| | |
|------------|--|
| Syntax | <code>exec sql [at <i>connection_name</i>] close <i>cursor_name</i>;</code> |
| Parameters | <p><i>cursor_name</i></p> <p>The name of the cursor to be closed; that is, the name that you assigned when declaring the cursor.</p> |
| Examples | <pre> long SQLCODE; exec sql begin declare section; CS_CHAR mlname[40]; CS_CHAR mfname[20]; CS_CHAR phone[12]; exec sql end declare section; exec sql declare author_list cursor for select au_lname, au_fname, phone from authors; exec sql open author_list; while (SQLCODE == 0) { exec sql fetch author_list into :mlname, :mfname, :mphone; if (SQLCODE != 100) printf("%s, %s, %s\n", mlname, mfname, mphone); } exec sql close author_list; </pre> |
| Usage | <ul style="list-style-type: none"> • The close statement closes an open cursor. Unfetched rows are canceled. • Reopening a closed cursor executes the associated query again, positioning the cursor pointer before the first row of the result set. • A cursor must be closed before it is reopened. • Attempting to close a cursor that is not open causes a runtime error. • The commit transaction, rollback transaction, commit work, and rollback work statements close a cursor automatically unless you set a precompiler option to disable the feature. • Closing and then reopening a cursor lets your program see any changes in the tables from which the cursor retrieves rows. |
| See also | <code>declare cursor</code> , <code>fetch</code> , <code>open</code> , <code>prepare</code> |

commit

| | |
|-------------|---|
| Description | Ends a transaction, preserving changes made to the database during the transaction. |
| Syntax | <pre>exec sql [at <i>connection_name</i>] commit [transaction tran work] [<i>transaction_name</i>];</pre> |
| Parameters | <p>transaction trans work</p> <p>The keywords transaction, trans, and work are interchangeable in the rollback statement, except that only work is ANSI-compliant.</p> <p><i>transaction_name</i></p> <p>A name assigned to the transaction.</p> |
| Examples | <pre>/* ** Using chained transaction mode, ** synchronize tables on two servers */ exec sql begin declare section; char title_id[7]; int num_sold; exec sql end declare section; long SQLCODE; ... try_update: exec sql whenever sqlerror goto abort_tran; exec sql at connect1 select sum(qty) into :num_sold from salesdetail where title_id = :title_id; exec sql at connect2 update current_sales set num_sold = :num_sold where title_id = :title_id; exec sql at connect2 commit work; exec sql at connect1 commit work; return; abort_tran: printf("oops, should have used 2-phase commit\n"); exec sql whenever sqlerror continue; exec sql at connect2 rollback work; exec sql at connect1 rollback work; goto try_update;</pre> |

| | |
|----------|---|
| Usage | <ul style="list-style-type: none"> • This reference page mainly describes aspects of the Transact-SQL commit statement that differ when used with Embedded SQL. See the <i>Adaptive Server Enterprise Reference Manual</i> for more information about commit and Transact-SQL transaction management. • Transaction names must conform to the Transact-SQL rules for identifiers. Transaction names are a Transact-SQL extension: they cannot be used with the ANSI-compliant keyword work. • When nesting transactions, assign a transaction name only to the outermost begin transaction statement and its corresponding commit transaction or rollback transaction statement. |
| See also | begin transaction, commit work, rollback transaction, rollback work |

connect

| | |
|-------------|---|
| Description | Creates a connection to Adaptive Server. |
| Syntax | <pre>exec sql connect <i>user_name</i> [identified by <i>password</i>] [at <i>connection_name</i>] [using <i>server_name</i>] [labelname <i>label_name</i> labelvalue <i>label_value</i> ...];</pre> |
| Parameters | <p><i>user_name</i> The user name to be used when logging in to Adaptive Server.</p> <p><i>password</i> The password to use to log in to Adaptive Server.</p> <p><i>connection_name</i> A name that you choose to uniquely identify the Adaptive Server connection.</p> <p><i>server_name</i> The server name of the Adaptive Server to which you are connecting.</p> |
| Examples | <pre>exec sql begin declare section; CS_CHAR user[32]; CS_CHAR password[32]; CS_CHAR server[90]; CS_CHAR conname[20]; exec sql end declare section; strcpy(user, "mylogin"); strcpy(password, "mypass"); strcpy(server, "YOURSERVER");</pre> |

```
strcpy(conname, "con_one");  
  
exec sql connect :user identified by :password  
        using :server at :conname;
```

Usage

- In every Embedded SQL program, the connect statement must be executed before any other executable SQL statement except allocate descriptor.
- If a program uses both C and COBOL languages, the first connect statement must be issued from a COBOL program.
- If a program has multiple connections, only one can be unnamed, and will be the default connection.
- If an Embedded SQL statement does not have an *at connection_name* clause to direct it to a specific named connection, the statement is executed on the current connection.
- To specify a null password, omit the identified by clause or use an empty string.
- If the connect statement does not specify a Adaptive Server, the server named by the DSQUERY environment variable or logical name is used. If DSQUERY is not defined, the default server is SYBASE.
- Client-Library looks up the server name in the interfaces file located in the directory specified by the SYBASE environment variable or logical name.
- The Adaptive Server connection ends when the Embedded SQL program exits or issues a disconnect statement.
- Opening a new connection, named or unnamed, results in the new connection becoming the current connection.
- A program that requires multiple Adaptive Server login names can have a connection for each login account.
- By connecting to more than one server, a program can simultaneously access data stored on different servers.
- A single program can have multiple connections to a single server or multiple connections to different servers.
- The following table shows how a connection is named:

Table 10-1: How a connection is named

| If this clause is used | | Then, the Connection Name is |
|---------------------------|-------------|------------------------------|
| at <i>connection_name</i> | But without | <i>connection_name</i> |

| If this clause is used | But without | Then, the Connection Name is |
|--------------------------|-------------|------------------------------|
| using <i>server_name</i> | at | <i>server_name</i> |
| None | | DEFAULT |

See also

at connection_name, exec sql, disconnect, set connection

deallocate cursor

| | |
|-------------|--|
| Description | Deallocates a cursor for a static SQL statement or for a dynamic SQL statement. |
| Syntax | exec sql [at <i>connection_name</i>] deallocate cursor <i>cursor_name</i> ; |
| Parameters | <p><i>cursor_name</i></p> <p>The name of the cursor to be deallocated. The <i>cursor_name</i> must be a character string enclosed in double quotation marks or in no quotation marks—for example "<i>my_cursor</i>" or <i>my_cursor</i>. It cannot be a host variable.</p> |

Examples

```
exec sql include sqlca;

main()
{
    exec sql begin declare section;
        CS_CHAR title[80];
        CS_SMALLINT i_title;
    exec sql end declare section;

    exec sql whenever sqlerror call error_handler();
    exec sql whenever sqlwarning call error_handler();
    exec sql whenever not found continue;

    exec sql connect "sa";
    exec sql use pubs2;
    exec sql declare title_list cursor for select title from titles;
```

```
exec sql open title_list;
for (;;)
{
    exec sql fetch title_list into :title :i_title;
    if (sqlca.sqlcode == 100) break;

    if (i_title == -1) printf("Title is NULL.\n");

    printf("Title: %s\n", title);
}
exec sql close title_list;
exec sql deallocate cursor title_list;
exec sql disconnect all;
exit(0);
}
error_handler()
{
    printf("%d\n%s\n", sqlca.sqlcode, sqlca.sqlerrm.sqlerrmc);
    exec sql deallocate cursor title_list;
    exec sql disconnect all;
    exit(-1);
}
```

Usage

- Deallocating a cursor releases all resources allocated to the cursor. In particular, `deallocate cursor` drops the Client-Library command handle and `CS_COMMAND` structure associated with the cursor.
- A static cursor can be deallocated at any time after it is opened. A dynamic cursor can be deallocated at any time after it is declared.
- If *cursor_name* is open, `deallocate cursor` closes it and then deallocates it.
- You cannot reference a deallocated cursor, nor can you reopen it. If you try, an error occurs.

- You can declare a new cursor having the same name as that of a deallocated cursor. Opening a cursor with the same name as a deallocated cursor is not the same as reopening the deallocated cursor. Other than the name, the new cursor shares nothing with the deallocated cursor.
- Declaring a new cursor with the same name as that of a deallocated cursor can cause the precompiler to generate a warning message.
- The deallocate cursor statement is a Sybase extension; it is not defined in the SQL standard.

Note If you are using persistent binding in your Embedded SQL program, use the deallocate cursor statement carefully. Needlessly deallocating cursors can negate the advantage of persistent binding.

See also

close cursor, declare cursor, open (static cursor)

deallocate descriptor

| | |
|-------------|---|
| Description | Deallocates a SQL descriptor. |
| Syntax | <code>exec sql deallocate descriptor <i>descriptor_name</i>;</code> |
| Parameters | <p><i>descriptor_name</i></p> <p>The name of the SQL descriptor that contains information about the dynamic parameter markers or return values in a prepared statement.</p> |
| Examples | <pre> exec sql begin declare section; CS_INT numcols, colnum; exec sql end declare section; ... exec sql allocate descriptor big_desc with max 1000; exec sql prepare dynstmt from "select * from huge_table"; exec sql execute dynstmt into sql descriptor big_desc; exec sql get descriptor :numcols = count; for (colnum = 1; colnum <= numcols; colnum++) { exec sql get descriptor big_desc ... } </pre> |

```
exec sql deallocate descriptor big_desc;
...
```

| | |
|----------|---|
| Usage | <ul style="list-style-type: none">• If you attempt to deallocate a SQL descriptor that has not been allocated, an error occurs. |
| See also | allocate descriptor |

deallocate prepare

| | |
|-------------|--|
| Description | Deallocates a dynamic SQL statement that was prepared in a prepare statement. |
| Syntax | <pre>exec sql [at <i>connection_name</i>] deallocate prepare <i>statement_name</i>;</pre> |
| Parameters | <p><i>statement_name</i></p> <p>The identifier assigned to the dynamic SQL statement when the statement was prepared.</p> |
| Examples | <pre>exec sql begin declare section; CS_CHAR sqlstmt[100]; exec sql end declare section; strcpy(sqlstmt, "select * from publishers"); exec sql prepare make_work from :sqlstmt; exec sql declare make_work_cursor cursor for make_work; exec sql deallocate prepare make_work;</pre> |
| Usage | <ul style="list-style-type: none">• A statement must be prepared before it is deallocated. Attempting to deallocate a statement that has not been prepared results in an error.• <i>statement_name</i> must uniquely identify a statement buffer and must conform to the SQL identifier rules for naming variables. <i>statement_name</i> can be either a literal or a character array host variable. |

- The deallocate prepare statement closes and deallocates any dynamic cursors declared for *statement_name*.

Warning! If you are using persistent binds in your embedded SQL program, use the deallocate prepare statement carefully. Needlessly deallocating prepared statements can negate the advantage of persistent binds.

See also

declare cursor (dynamic), execute, execute immediate, prepare

declare cursor (dynamic)

| | |
|-------------|---|
| Description | Declares a cursor for processing multiple rows returned by a prepared dynamic select statement. |
| Syntax | <pre>exec sql [at <i>connection_name</i>] declare <i>cursor_name</i> cursor for <i>prepped_statement_name</i>;</pre> |
| Parameters | <p><i>cursor_name</i></p> <p>The cursor's name, used to reference the cursor in open, fetch, and close statements. A cursor's name must be unique on each connection and must have no more than 128 characters.</p> <p><i>prepped_statement_name</i></p> <p>The name (specified in a previous prepare statement) that represents the select statement to be executed.</p> |
| Examples | <pre>exec sql begin declare section; CS_CHAR sqlstmt[100]; exec sql end declare section; strcpy(sqlstmt, "select * from publishers"); exec sql prepare make_work from :sqlstmt; exec sql declare make_work_cursor cursor for make_work; exec sql deallocate prepare make_work;</pre> |
| Usage | <ul style="list-style-type: none"> • The <i>prepped_statement_name</i> must not have a compute clause. • The <i>cursor_name</i> must be declared on the connection where <i>prepped_statement_name</i> was prepared. |

- The dynamic declare cursor statement is an executable statement, whereas the static declare cursor statement is simply a declaration. The dynamic declare statement must be located where the host language allows executable statements and the program should check return codes (SQLCODE, SQLCA, or SQLSTATE).
- The for update and read only clauses for a dynamic cursor are not part of the declare cursor statement but rather should be included in the prepared statement's select query.

See also

close, connect, fetch, open, prepare

declare cursor (static)

Description Declares a cursor for processing multiple rows returned by a select statement.

Syntax

```
exec sql declare cursor_name
cursor for select_statement
[for update [of col_name_1 [, col_name_n]...]]
for read only];
```

Parameters

cursor_name

The cursor's name, used to reference the cursor in open, fetch, and close statements. A cursor's name must be unique on each connection and must have no more than 128 characters.

select_statement

The Transact-SQL select statement to be executed when the cursor is opened. See the description of the select statement in the *Adaptive Server Enterprise Reference Manual* for more information.

for update

Specifies that the cursor's result list can be updated. (To update the result list, you use the update statement.

of *col_name_1*

The name of the first column to be updated.

of *col_name_n*

The name of the *n*th column to be updated.

for read only

Specifies that the cursor's result list cannot be updated.

Examples

```
main( )
{
```

```

exec sql begin declare section;
        CS_CHAR          b_titleid[TIDSIZE+1];
        CS_CHAR          b_title[65];
        CS_CHAR          b_type[TYPESIZE+1];
exec sql end declare section;
        long            SQLCODE;

exec sql connect "sa";

exec sql use pubs2;

exec sql declare titlelist cursor for
        select title_id, substring(title,1,64)
        from titles where type like :b_type;

strcpy(b_type, "business");
exec sql open titlelist;
for (;;)
{
        exec sql fetch titlelist into :b_titleid,
                :b_title;
        if (SQLCODE == 100)
                break;
        printf("    %-8s %s\n", b_titleid, b_title);
}
exec sql close titlelist;

exec sql disconnect all;
}

```

Usage

- The Embedded SQL precompiler generates no code for the declare cursor statement.
- The *select_statement* does not execute until your program opens the cursor by using the open cursor statement.
- The syntax of the *select_statement* is identical to that shown in the *Adaptive Server Enterprise Reference Manual*, except that you cannot use the compute clause in Embedded SQL.
- The *select_statement* can contain host variables. The values of the host variables are substituted when your program opens the cursor.
- If you omit either the for update or read only clause, Adaptive Server determines whether the cursor is updatable.

See also

close, connect, deallocate cursor, declare cursor (stored procedure), declare cursor (dynamic), fetch, open, update

declare cursor (stored procedure)

| | |
|-------------|---|
| Description | Declares a cursor for a stored procedure. |
| Syntax | <pre>exec sql declare <i>cursor_name</i> cursor for execute <i>procedure_name</i> ([[@<i>param_name</i> =]:<i>host_var</i>] [,[@<i>param_name</i> =]:<i>host_var</i>]...)</pre> |
| Parameters | <p><i>cursor_name</i></p> <p>The cursor's name, used to reference the cursor in open, fetch, and close statements. A cursor's name must be unique on each connection and must have no more than 128 characters.</p> <p><i>procedure_name</i></p> <p>The name of the stored procedure to be executed.</p> <p><i>param_name</i></p> <p>The name of a parameter in the stored procedure.</p> <p><i>host_var</i></p> <p>The name of a host variable to be passed as a parameter value.</p> |
| Examples | <pre>main() { exec sql begin declare section; CS_CHAR b_titleid[7]; CS_CHAR b_title[65]; CS_CHAR b_type[13]; exec sql end declare section; long SQLCODE; exec sql connect "sa"; exec sql use pubs2; exec sql create procedure p_titles (@p_type varchar(30)) as select title_id, substring(title,1,64) from titles where type like @p_type; exec sql declare titlelist cursor for execute p_titles (:b_type); strcpy(b_type, "business"); exec sql open titlelist; for (;;) </pre> |

```
        {
            exec sql fetch titlelist into :b_titleid,
                :b_title;
            if (SQLCODE == 100)
                break;
            printf("    %-8s %s\n", b_titleid, b_title);
        }

    exec sql close titlelist;
    exec sql disconnect all;
}
```

| | |
|----------|--|
| Usage | <ul style="list-style-type: none">• <i>procedure_name</i> must consist of only one select statement.• It is not possible to retrieve output parameter values from a stored procedure executed using a cursor.• It is not possible to retrieve the return status value of a stored procedure executed using a cursor. |
| See also | close, deallocate cursor, declare cursor (static), declare cursor (dynamic), fetch, open, update |

delete (positioned cursor)

| | |
|-------------|---|
| Description | Removes, from a table, the row indicated by the current cursor position for an open cursor. |
| Syntax | <pre>exec sql [at <i>connection_name</i>] delete [from] <i>table_name</i> where current of <i>cursor_name</i>;</pre> |
| Parameters | <p><i>table_name</i></p> <p>The name of the table from which the row will be deleted.</p> <p>where current of <i>cursor_name</i></p> <p>Causes Adaptive Server to delete the row of the table indicated by the current cursor position for the cursor <i>cursor_name</i>.</p> |

Examples

```
exec sql include sqlca;

main()
{
    char answer[1];
```

```
exec sql begin declare section;
    CS_CHAR disc_type[40];
    CS_CHAR store_id[5];
    CS_SMALLINT ind_store_id;
exec sql end declare section;
exec sql connect "sa";
exec sql use pubs2;
exec sql declare purge_cursor cursor for
    select discounttype, stor_id
    from discounts;
exec sql open purge_cursor;
exec sql whenever not found goto alldone;
while (1)
    {
    exec sql fetch purge_cursor into :disc_type, :store_id
        :ind_store_id;
    if (ind_store_id != -1)
        {
        printf("%s, %s\n", disc_type, store_id);
        printf("Delete Discount Record? (y/n) >");
        gets(answer);
        if (strncmp(answer, "y", 1) == 0)
            {
            exec sql delete from discounts where
                current of purge_cursor;
            }
        }
    }
/*
** No changes will be committed to the database because
** this program does not contain an "exec sql commit work;"
** statement. The changes will be rolled back when the
```



```

**  user disconnects.
*/

alldone:

    exec sql close purge_cursor;

    exec sql disconnect all;

}

```

Usage

- This reference page mainly describes aspects of the Transact-SQL delete statement that differ when used with Embedded SQL. See the *Adaptive Server Enterprise Reference Manual* for more information about the delete statement.
- This form of the delete statement must execute on the connection where the cursor *cursor_name* was opened. If the delete statement includes the *at connection_name* clause, the clause must match the *at connection_name* clause of the open cursor statement that opened *cursor_name*.
- The delete statement fails if the cursor was declared for read only, or if the select statement included an order by clause.

See also

close, declare cursor, fetch, open, update

delete (searched)

Description

Removes rows specified by search conditions.

Syntax

```

exec sql [at connection_name] delete table_name_1
[from table_name_n
[, table_name_n]...]
[where search_conditions];

```

Parameters

table_name_1

The name of the table from which this delete statement deletes rows.

from *table_name_n*

The name of a table to be joined with *table_name_1* to determine which rows of *table_name_1* will be deleted. The delete statement does *not* delete rows from *table_name_n*.

where *search_conditions*

Specifies which rows will be deleted. If you omit the where clause, the delete statement deletes all rows of *table_name_1*.

Examples

```
/*
```

```

** Function to FAKE a cascade delete of an author **
**by name -- this function assumes that pubs2 is
** the current database.
** Returns 1 for success, 0 for failure
**/
    int          drop_author(fname, lname)
    char         *fname;
    char         *lname;
{
exec sql begin declare section;
    CS_CHAR      f_name[41], l_name[41];
    CS_CHAR      titleid[10], auid[10];
exec sql end declare section;
    long         SQLCODE;
strcpy(f_name, fname);
strcpy(l_name, lname);
exec sql whenever sqlerror goto roll_back;
exec sql select au_id from authors into :auid
            where au_fname = :f_name
            and au_lname = :l_name;

exec sql delete from au_pix where au_id = :auid;
exec sql delete from blurbs where au_id = :auid;
exec sql declare curl cursor for
            select title_id from titleauthor
            where au_id = :auid;
exec sql open curl;
while (SQLCODE == 0)
{
    exec sql fetch curl into :titleid;

    if(SQLCODE == 100) break;

    exec sql delete from salesdetail
            where title_id = :titleid;
    exec sql delete from rowsched
            where title_id = :titleid;
    exec sql delete from titles
            where title_id = :titleid;
    exec sql delete from titleauthor
            where current of curl;
}
exec sql close curl;
exec sql delete from authors
            where au_id = :auid;
exec sql commit work;
return 1;

```

```
roll_back:
    exec sql rollback work;
    return 0;
}
```

- Usage
- This reference page describes mainly aspects of the Transact-SQL delete statement that differ when used with Embedded SQL. See the *Adaptive Server Enterprise Reference Manual* for more information about the delete statement.
 - If you need to remove rows specified by the current position of a cursor pointer, use the delete (positioned cursor) statement.
- See also
- close, declare cursor, fetch, open, update

describe input (SQL descriptor)

Description

Obtains information about dynamic parameter markers in a prepared dynamic SQL statement and stores that information in a SQL descriptor.

For a list of possible SQL descriptor datatype codes, see “SQL descriptor datatype codes” on page 172.

Syntax

```
exec sql describe input statement_name
using sql descriptor descriptor_name;
```

Parameters

statement_name

The name of the prepared statement about which you want information. *statement_name* must identify a prepared statement.

sql descriptor

Identifies *descriptor_name* as a SQL descriptor.

descriptor_name

The name of the SQL descriptor that is to store information about the dynamic parameter markers in the prepared statement.

Examples

```
exec sql begin declare section;
char          query[maxstmt];
int           nin, nout, i;
exec sql end declare section;
int          j;

...

exec sql allocate descriptor din with max 256;
exec sql allocate descriptor dout with max 256;
```

```
exec sql whenever sqlerror stop;
exec sql prepare dynstmt from :query;

exec sql describe input dynstmt
      using sql descriptor din;

exec sql get descriptor din :nin = count;
      for (i = 0; i < nin; i++)
```

Usage

- Information about the statement is written into the descriptor provided in the using clause. Use the get descriptor statement after executing the describe input statement to extract information from the descriptor into host variables.
- The descriptor must be allocated before the describe input statement can be executed.

See also

allocate descriptor, deallocate descriptor, describe output, get descriptor, prepare, set descriptor

describe input (SQLDA)

Description

Obtains information about dynamic parameter markers in a prepared dynamic SQL statement and stores that information in a SQLDA structure.

Syntax

```
exec sql describe input statement_name
      using descriptor descriptor_name;
```

Parameters

statement_name

The name of the prepared statement about which you want information. *statement_name* must identify a prepared statement.

descriptor

Identifies *descriptor_name* as an SQLDA structure.

descriptor_name

The name of the SQLDA structure that is to store information about the dynamic parameter markers in the prepared statement.

Examples

```
...

exec sql prepare s4 from :str4;
exec sql declare c2 cursor for s4;
exec sql describe input s4 using descriptor dinout;
printf("Number of input parameters is %hd\n",
```

```
dinout.sd.sqld);
```

- Usage
- Information about the statement is written into the descriptor specified in the using clause. After the get descriptor statement is executed, you can read the information out of the SQLDA structure.
- See also
- allocate descriptor, deallocate descriptor, describe output, get descriptor, prepare, set descriptor

describe output (SQL descriptor)

- Description
- Obtains row format information about the result set of a prepared dynamic SQL statement.
- For a list of possible SQL descriptor datatype codes, see Table 10-5 on page 172.
- Syntax
- ```
exec sql describe [output] statement_name
using sql descriptor descriptor_name;
```
- Parameters
- output
- An optional keyword that has no effect on the describe output statement but provides conformance to the SQL standard.
- statement\_name*
- The name (specified in a prepare statement) that represents the select statement to be executed.
- sql descriptor
- Identifies *descriptor\_name* as a SQL descriptor.
- descriptor\_name*
- The name of a SQL descriptor that is to store the information returned by the describe output statement.
- Examples
- ```
...
exec sql open curs2 using sql descriptor descr_out;

exec sql describe output prep_stmt4
using sql descriptor descr_out;

while (sqlca.sqlcode != 100 && sqlca.sqlcode >= 0)
{
    exec sql fetch curs2 into sql descriptor
        descr_out;
    print_descriptor();
}
```

```
exec sql close curs2;
exec sql deallocate descriptor descr_out;
exec sql deallocate prepare prep_stmt4;
printf("dynamic sql method 4 completed\n\n");
}
```

...

- Usage
- The information obtained is the type, name, length (or precision and scale, if a number), nullable status, and number of items in the result set.
 - The information is about the result columns from the select column list.
 - Execute this statement before the prepared statement executes. If you perform a describe output statement after you execute and before you perform a get descriptor, the results will be discarded.

See also allocate descriptor, describe input, execute, get descriptor, prepare

describe output (SQLDA)

Description Obtains row format information about the result set of a prepared dynamic SQL statement and stores that information in a SQLDA structure.

Syntax `exec sql describe [output] statement_name
 using descriptor sqlda_name;`

Parameters **output**
 An optional keyword that has no effect on the describe output statement but provides conformance to the SQL standard.

statement_name
 The name (specified in a prepare statement) that represents the select statement to be executed.

descriptor
 Identifies *sqlda_name* as a SQLDA structure.

sqlda_name
 The name of a SQLDA structure that is to store the information returned by the describe output statement.

Examples

...

```
exec sql open curs2 using descriptor input_descriptor;
```

```
exec sql describe output statement using descriptor
    output_descriptor;

output_descriptor->sqllda_column->sqllda_sqldata = character;

output_descriptor->sqllda_column->sqllda_datafmt.datatype =      CS_CHAR_TYPE;

output_descriptor->sqllda_column->sqllda_datafmt.maxlength = 20;

output_descriptor->sqllda_column->sqllda_sqllen = 20;

output_descriptor->sqllda_column->sqllda_datafmt.format =
    (CS_FMT_NULLTERM | CS_FMT_PADBLANK);

exec sql fetch curs2 into descriptor output_descriptor;
```

- Usage
- The information obtained is the data held in the SQLDA fields, such as the type, name, length (or precision and scale, if a number), nullable status, and number of items in the result set.
 - The information is about the result columns from the select column list.
- See also
- describe input, execute, prepare

disconnect

Description Closes one or more connections to a Adaptive Server.

Syntax `exec sql disconnect`
`{connection_name | current | DEFAULT | all};`

Parameters *connection_name*

The name of a connection to be closed.

current

Specifies that the current connection is to be closed.

DEFAULT

Specifies that the default connection is to be closed. This keyword must be in uppercase letters if you specify the default *connection_name* using a character string variable, for example:

```
exec sql disconnect :hv;
```

all

Specifies that all active connections be closed.

Examples

```
#include <stdio.h>
```

```
exec sql include sqlca;

main()
{
    exec sql begin declare section;
    CS_CHAR servname[31], username[31],
    password[31], conname[129];
    exec sql end declare section;

    exec sql whenever sqlerror call error_handler();
    exec sql whenever sqlwarning call error_handler();
    exec sql whenever not found continue;

    printf ("Username: ");
    gets  (username);
    printf ("Password: ");
    gets  (password);
    printf ("SQL Server name: ");
    gets  (servname);
    printf ("Connection name: ");
    gets  (conname);

    /*
    ** Make a named connection.
    */
        exec sql connect :username identified by :password
            at :conname using :servname;

    /*
    ** Make an unnamed (default) connection.
    */
        exec sql connect :username identified by :password
            using :servname;

    /*
    ** The second (default) connection is the current connection.
    */
        exec sql disconnect current;

    /*
    ** We now have neither a default connection nor a current one.
    */
        exec sql disconnect :conname;

    /*
    ** Now there are no open connections.
    */
        exec sql exit;
}
```



```
error_handler()
{
    printf("%d\n%s\n", sqlca.sqlcode, sqlca.sqlerrm.sqlerrmc);
    exit(0);
}
```

- Usage
- By itself, the disconnect keyword is not a valid statement. Instead, it must be followed by *connection_name*, current, DEFAULT, or all.
 - Closing a connection releases all memory and resources associated with that connection.
 - disconnect does not commit current transactions; it rolls them back. If an unchained transaction is active on the connection, disconnect rolls it back, ignoring any savepoints.
 - Closing a connection closes open cursors, drops temporary Adaptive Server objects, releases any locks the connection has in the Adaptive Server, and closes the network connection to the Adaptive Server.
- See also
- commit work, commit transaction, connect, rollback transaction, rollback work

exec

Description

Runs a system procedure or a user-defined stored procedure.

Syntax

```
exec sql [at connection_name]
exec [[:status_var =]status_value] procedure_name
[[[[@parameter_name =]param_value [out[put]]],...]]
[into :hostvar_1 [:indicator_1]
[, hostvar_n [indicator_n,...]]]
[with recompile];
```

Note Do not confuse the exec statement with the Embedded SQL execute statement; they are not related. The Embedded SQL exec statement is, however, the equivalent of the Transact-SQL execute statement.

Parameters

status_var

A host variable to receive the return status of the stored procedure.

status_value

The value of the stored procedure return status variable *status_var*.

procedure_name

The name of the stored procedure to be executed.

parameter_name

The name(s) of the stored procedure's parameter(s).

param_value

A host variable or literal value.

output

Indicates that the stored procedure returns a parameter value. The matching parameter in the stored procedure must also have been created using the output keyword.

into :hostvar_1

Causes row data returned from the stored procedure to be stored in the specified host variables (*hostvar_1* through *hostvar_n*). Each host variable can have an indicator variable.

with recompile

Causes Adaptive Server to create a new query plan for this stored procedure each time the procedure executes.

Examples**Example 1**

```
exec sql begin declare section;
    char          titleid[10];
    int           total_discounts;
    short         retstat;
exec sql end declare section exec;
exec sql create procedure get_sum_discounts
    (@titleid tid, @discount int output) as
begin
    select @discount = sum( qty * discount)
    from salesdetail
    where title_id = @titleid

end;
printf("title id: ");
gets(titleid);

exec sql exec
    :retstat = get_sum_discount :titleid,
    :total_discounts out;

printf("total discounts for title_id %s were
    %s\n", titleid, total_discounts);
exec sql begin declare section;
    CS_INT          status;
```

```

        CS_CHAR          city(30);
        CS_INT           result;
    exec sql end declare section;
    LONG                SQLCODE;

    input "City", city ;
    exec sql exec countcity :city, :result out;
    if (SQLCODE = 0)
        print city + " occurs " + result + "
            times." ;

```

Example 2

```

EXEC SQL BEGIN DECLARE SECTION;
/* storage for login name and password */
CS_CHAR          username[30], password[30];
CS_CHAR          pub_id[4][5], pub_name[4][40], stmt[100] ;
CS_CHAR          city[4][15], state[4][3];
CS_INT           ret_status;
EXEC SQL END DECLARE SECTION ;

...

EXEC SQL set chained off;

strcpy(stmt,"create proc get_publishers as select * from publishers
return ");

EXEC SQL EXECUTE IMMEDIATE :stmt;

EXEC SQL EXEC :ret_status = get_publishers INTO
                                :pub_id,
                                :pub_name,
                                :city,
                                :state;

printf("Pub Id      Publisher Name          City          State \n");
printf("\n-----  ----- \n");

for ( i = 0 ; i < sqlca.sqlerrd[2] ; i++ )
{
    printf("%-8s", pub_id[i])    ;
    printf("%-25s", pub_name[i]) ;
    printf("%-12s", city[i])    ;
    printf("%-6s\n", state[i])  ;
}

printf("\n(%d rows affected, return status = %d)\n", sqlca.sqlerrd[2],
ret_status);

...

```

}

Usage

- Only one select statement can return rows to the client application.
- If the stored procedure contains select statements that can return row data, you must use one of two methods to store the data. You can either use the into clause of the exec statement or declare a cursor for the procedure. If you use the into clause, the stored procedure must not return more than one row of data, unless the host variables that you specify are arrays.
- The value *param_value* can be a host variable or literal value. If you use the output keyword, *param_value* must be a host variable.
- You can specify the output keyword for *parameter_name* only if that keyword was also used for the corresponding parameter of the create procedure statement that created *procedure_name*.
- The Embedded SQL exec statement works much like the Transact-SQL execute statement.

See also

declare cursor (stored procedure), select

exec sql

Description

Marks the beginning of a SQL statement embedded in a host language program.

Syntax

exec sql [at *connection_name*] *sql_statement*;

Parameters

at

Causes the SQL statement *sql_statement* to execute at the SQL Server connection *connection_name*.

connection_name

The connection name that identifies the SQL Server connection where *sql_statement* is to execute. The *connection_name* must be defined as a previous connect statement.

sql_statement

A Transact-SQL statement or other Embedded SQL statement.

Examples

```
exec sql
begin declare section;
    char    site1(20);
    int     sales1;
exec sql end declare section;
```

```

exec sql connect "user1" identified by "password1"
    using "server1";
exec sql connect "user2" identified by "password2"
    using "server2"
/* Remember that a connection that has not been
   explicitly named has the name of its server */
exec sql at server1 select count(*) from sales
    into :sales1;

sitel = sitename("server1");
exec sql at server2 insert into numsales
    values(:sitel, :sales1);

```

Usage

- SQL statements embedded in a host language must begin with `exec sql`. The keywords `exec sql` can appear anywhere that a host language statement can begin.
- The statement *sql_statement* can occupy one or more program lines; however, it must conform to host language rules for line breaks and continuation lines.
- The `at` clause affects only the statement *sql_statement*. The clause does not affect subsequent SQL statements, and does not reset the current connection.
- The `at` clause is not valid when *sql_statement* is one of the following SQL statements:

Table 10-2: Statements that cannot use the `at` clause of `exec sql`

| | | |
|-----------------------|-----------------------------|---------------------|
| allocate descriptor | begin declare section | connect |
| deallocate descriptor | declare cursor (dynamic) | end declare section |
| exit | get diagnostics | include file |
| include sqlca | set connection | set diagnostics |
| whenever | | |

- *connection_name* must be defined in a previous connect statement.
- Each Embedded SQL statement must end with a terminator. In C, the terminator is the semicolon (;).

See also

begin declare section, connect, disconnect, set connection

execute

| | |
|-------------|---|
| Description | <p>Executes a dynamic SQL statement from a prepared statement.</p> <p>For details on the execute immediate statement, see “execute immediate” on page 146.</p> |
| Syntax | <pre>exec sql [at <i>connection_name</i>] execute <i>statement_name</i> [into {<i>host_var_list</i> descriptor <i>descriptor_name</i> sql descriptor <i>descriptor_name</i>}] [using {<i>host_var_list</i> descriptor <i>descriptor_name</i> sql descriptor <i>descriptor_name</i>}];</pre> <hr/> <p>Note Do not confuse the Embedded SQL execute statement with the Embedded SQL exec statement or the Transact-SQL execute statement.</p> <hr/> |
| Parameters | <p><i>statement_name</i></p> <p>A unique identifier for the statement, defined in a previous prepare statement.</p> <p><i>descriptor_name</i></p> <p>Specifies the area of memory, or the SQLDA structure, that describes the statement’s dynamic parameter markers or select column list.</p> <p>into</p> <p>An into clause is required when the statement executes a select statement, which must be a single-row select. The target of the into clause can be a SQL descriptor, a SQLDA structure, or a list of one or more Embedded SQL host variables.</p> <p>Each host variable in the <i>host_var_list</i> must first be defined in a declare section. An <i>indicator variable</i> can be associated with a host variable to show when a null data value is retrieved.</p> <p>descriptor</p> <p>Identifies <i>descriptor_name</i> as a SQLDA structure.</p> <p>sql descriptor</p> <p>Identifies <i>descriptor_name</i> as a SQL descriptor.</p> |

using

The host variables that are substituted for dynamic parameter markers in *host_var_list*. The host variables, which you must define in a *declare* section, are substituted in the order listed. Use this clause only when *statement_name* contains dynamic parameter markers. The dynamic descriptor can also contain the values for the dynamic parameter markers.

Examples

```
exec sql begin declare section;
    CS_CHAR          dymo_buf(128);
    CS_CHAR          title_id(6);
    CS_INT           qty;
    CS_CHAR          order_no(20);
exec sql end declare section;

dymo_buf = "INSERT salesdetail
(ord_num, title_id, qty) VALUES (:?, :?, :?)"

exec sql prepare ins_com from :dymo_buf;

print "Recording Book Sales";
input "Order number?", order_no;
input "Title ID?", title_id;
input "Quantity sold?", qty;

exec sql execute ins_com
    using :order_no, :title_id, :qty;

exec sql disconnect;
```

Usage

- *execute* is the second step in method 2 of dynamic SQL. The first step is the *prepare* statement.
- *prepare* and *execute* are valid with any SQL statement except a multirow *select* statement. For multirow *select* statements, use either dynamic cursor.
- The statement in *statement_name* can contain dynamic parameter markers ("?"). They mark the positions where host variable values are to be substituted before the statement executes.
- The *execute* keyword distinguishes this statement from *exec*. See *exec* on page 139.

See also

declare section, *get descriptor*, *prepare*, *set descriptor*

execute immediate

| | |
|-------------|---|
| Description | Executes a dynamic SQL statement stored in a character-string host variable or quoted string. |
| Syntax | <code>exec sql [at <i>connection_name</i>] execute immediate {:<i>host_variable</i> "string"};</code> |
| Parameters | <p><i>host_variable</i></p> <p>A character-string host variable defined in a declare section. Before calling <code>execute immediate</code>, the host variable should contain a complete and syntactically correct Transact-SQL statement.</p> <p><i>string</i></p> <p>A quoted literal Transact-SQL statement string that can be used in place of <i>host_variable</i>.</p> |
| Examples | <pre>exec sql begin declare section; CS_CHAR host_var(128); exec sql end declare section; printf("Enter a non-select SQL statement: "); gets(host_var); exec sql execute immediate :host_var;</pre> |
| Usage | <ul style="list-style-type: none">• Using the <code>execute immediate</code> statement is dynamic SQL Method 1. See Chapter 7, "Using Dynamic SQL" for information about the four dynamic SQL methods.• Except for messages, the statement in <i>host_variable</i> cannot return results to the your program. Thus, the statement cannot be, for example, a <code>select</code> statement.• The Embedded SQL precompiler does not check the syntax of the statement stored in <i>host_variable</i> before sending it to Adaptive Server. If the statement's syntax is incorrect, Adaptive Server returns an error code and message to your program.• Use <code>prepare</code> and <code>execute</code> (dynamic SQL method 2) to substitute values from host variables into a dynamic SQL statement.• Use <code>prepare</code>, <code>open</code>, and <code>fetch</code> (dynamic SQL method 3) to execute <code>select</code> statements with dynamic SQL statements that return results. |
| See also | <code>execute</code> , <code>prepare</code> |

exit

| | |
|-------------|--|
| Description | Closes Client-Library and deallocates all Embedded SQL resources allocated to your program. |
| Syntax | <code>exec sql exit;</code> |
| Examples | <pre>exec sql include sqlca; main() { /* The body of the main function goes here, ** including various Embedded SQL statements. */ ... /* The exit statement must be the last ** embedded SQL statement in the program. */ exec sql exit; } /* end of main */</pre> |
| Usage | <ul style="list-style-type: none">• The <code>exit</code> statement closes all connections that your program opened. Also, <code>exit</code> deallocates all Embedded SQL resources and Client-Library resources allocated to your program.• Although the <code>exit</code> statement is valid on all platforms, it is required only on some. For more information, see the <i>Open Client/Server Programmer's Supplement</i>.• You cannot use Client-Library functions after using the <code>exit</code> statement, unless you initialize Client-Library again. See the <i>Open Client Client-Library/C Programmer's Guide</i> for information about initializing Client-Library.• The <code>exit</code> statement is a Sybase extension; it is not defined in the SQL standard. |
| See also | <code>disconnect</code> |

fetch

| | |
|-------------|---|
| Description | Copies data values from the current cursor row into host variables or a dynamic descriptor. |
| Syntax | <code>exec sql [at <i>connection_name</i>] fetch [rebind norebind] <i>cursor_name</i> into {:<i>host_variable</i> [[<i>indicator</i>]:<i>indicator_variable</i>]</code> |

```
[, :host_variable
[[indicator]:indicator_variable]]... |
descriptor descriptor_name |
sql descriptor descriptor_name);
```

Parameters

rebind | *norebind*

Specifies whether host variables require rebinding for this `fetch` statement. The `rebind` clause overrides precompiler options that control rebinding.

cursor_name

The name of the cursor. The name is defined in a preceding `declare cursor` statement.

host_variable

A host language variable defined in a `declare` section.

indicator_variable

A 2-byte host variable declared in a previous `declare` section. If the value for the associated variable is null, `fetch` sets the indicator variable to -1. If truncation occurs, `fetch` sets the indicator variable to the actual length of the result column. Otherwise, it sets the indicator variable to 0.

descriptor

Identifies *descriptor_name* as a SQLDA structure.

sql descriptor

Identifies *descriptor_name* as a SQL descriptor.

descriptor_name

The name of the dynamic descriptor that is to hold a result set.

Examples

```
exec sql begin declare section;
      CS_CHAR          title_id[6];
      CS_CHAR          title[80];
      CS_CHAR          type[12];
      CS_SMALLINT      i_title;
      CS_SMALLINT      i_type;
exec sql end declare section;
exec sql declare title_list cursor for
      select type, title_id, title from titles
      order by type;

exec sql open title_list
while (sqlca.sqlcode != 100) {
exec sql fetch title_list into
      :type :i_type, :title_id, :title :i_title;

      if (i_type != -1) {
```

```
        printf("Type: %s\n", type);
    }
    else {
        printf("Type: undecided\n");
    }

    printf("Title id: %s\n", title_id);

    if (i_title <> -1) {
        print "Title: ", title;
    }
    else {
        print "Title: undecided";
    }
}

exec sql close title_list;
```

Usage

- The fetch statement can be used both with static cursors and with cursors in dynamic SQL.
- The open statement must execute before the fetch statement executes.
- The first fetch on an open cursor returns the first row or group of rows from the cursor's result table. Each subsequent fetch returns the next row or group of rows.
- You can fetch multiple rows into an array.
- The "current row" is the row most recently fetched. To update or delete it, use the where current of *cursor_name* clause with the update or delete statement. These statements are not valid until after a row has been fetched.
- After all rows have been fetched from the cursor, calling fetch sets SQLCODE to 100. If the select furnishes no results on execution, SQLCODE is set to 100 on the first fetch.
- There must be one, and only one, *host_variable* for each column of the result set.
- When neither the rebind nor the norebind is specified, the binding behavior is determined by the precompiler option -b. See "Guidelines for using persistent binding" on page 105 for information on persistent binds and the *Open Client/Server Programmer's Supplement* for your platform for details on precompiler options.

- An *indicator_variable* must be provided for a *host_variable* that can receive a null value. A runtime error occurs when a null value is fetched for a host variable that has no indicator variable.
- When possible, Client-Library converts the datatype of a result column to the datatype of the corresponding host variable. If Client-Library cannot convert a datatype, it issues an error message. If conversion is not possible, an error occurs.

See also

allocate descriptor, close, declare, delete (positioned cursor), open, prepare, update

get descriptor

Description Retrieves attribute information about dynamic parameter markers and select column list attributes and data from a SQL descriptor.

For a list of SQL descriptor datatype codes, see Table 10-5 on page 172.

Syntax

```
exec sql get descriptor descriptor_name
{:host_variable = count |
value item_number :host_variable = item_name
[, :host_variable = item_name]...};
```

Parameters

descriptor_name
The name of the SQL descriptor that contains information about the dynamic parameter markers or return columns in a prepared statement.

host_variable
A variable defined in a declare section.

count
The number of dynamic parameters retrieved.

item_number
A number specifying the *n*th dynamic parameter marker or select column for which get descriptor retrieves information.

item_name
The name of an attribute to be retrieved. See Table 10-3 on page 10-57 .

Table 10-3: Valid item_name values

| Value | Description |
|------------------------|---|
| <i>data</i> | Value for the dynamic parameter marker or target associated with the specified SQL descriptor. If indicator is negative, this field is undefined. |
| <i>indicator</i> | Value for the indicator parameter associated with the dynamic parameter marker or target. |
| <i>length</i> | The length, in characters, of the dynamic parameter marker or target for the specified SQL descriptor. |
| <i>name</i> | The name of the specified SQL descriptor containing information about the dynamic parameter markers. |
| <i>nullable</i> | Equals 0 if the dynamic parameter marker can accept a null value; otherwise, equals 1. |
| <i>precision</i> | An integer specifying the total number of digits of precision for the CS_NUMERIC variable. |
| <i>returned_length</i> | The length of character types of the values from the select column list. |
| <i>scale</i> | An integer specifying the total number of digits after the decimal point for the CS_NUMERIC variable. |
| <i>type</i> | The datatype of this column (item number) in the row. For values, see “SQL descriptor datatype codes” on page 172. |

Examples

```

exec sql begin declare section;
    int      numcols, colnum, type, intbuf;
    char     charbuf[100];
exec sql end declare section;
...
exec sql allocate descriptor big_desc
    with max 1000;
exec sql prepare dynstmt from "select * from \
    huge_table";
exec sql execute dynstmt into sql descriptor
    big_desc;

exec sql get descriptor big_desc :numcols = count;
for (colnum = 1; colnum <= numcols; colnum++)
{
    exec sql get descriptor big_desc
        value :colnum :type = type;
}

```

```
if (type == 4)
{
    exec sql get descriptor big_desc
        value :colnum :intbuf = data;

    /* Display intbuf. */
    ...
}
else if (type == 1)
{
    big_desc
        value :colnum :charbuf = data;

    /* Display charbuf. */
    ...
}
}
exec sql deallocate descriptor big_desc;
...
```

Usage

- The get descriptor statement returns information about the number or attributes of dynamic parameters specified or the select list columns in a prepared statement.
- This statement should be executed after a describe input, describe output, execute, or fetch (dynamic) statement has been issued.
- It is not possible to retrieve *data*, indicator, or *returned_length* until the data associated with the descriptor is retrieved from the server by an execute statement or fetch statement.

See also

describe input, describe output, fetch, set descriptor

get diagnostics

Description

Retrieves error, warning, and informational messages from Client-Library.

Syntax

```
get diagnostics
{ :hv = statement_info [, :hv = statement_info]...|
  exception :condition_number
  :hv = condition_info [, :hv = condition_info]...}
```

| | |
|------------|---|
| Parameters | <p><i>statement_info</i></p> <p>The keyword <code>number</code> is currently the only supported <i>statement_info</i> type. It returns the total number of exceptions in the diagnostics queue.</p> <p><i>condition_info</i></p> <p>Any one of the keywords <i>sqlca_info</i>, <i>sqlcode_number</i>, and <i>returned_sqlstate</i>.</p> |
| Examples | <pre>exec sql begin declare section; CS_INT num_msgs; CS_INT condcnt=1; exec sql include sqlca; exec sql end declare section; exec sql exec sp_password "bass", "foo"; exec sql get diagnostics :num_msgs = number; printf("Number of messages is %d.\n", num_msgs); /* Loop through and print the messages. */ while (condcnt <= num_msgs) { exec sql get diagnostics exception :condcnt :sqlca = sqlca_info; printf("SQLCODE = %d \n", sqlca.sqlcode); printf("%s \n", sqlca.sqlerrm.sqlerrmc); condcnt = condcnt + 1; }</pre> |
| Usage | <ul style="list-style-type: none">• Many Embedded SQL statements are capable of causing multiple warnings or errors. Typically, only the first error is reported via <code>SQLCODE</code>, <code>SQLCA</code>, or <code>SQLSTATE</code>. Use <code>get diagnostics</code> to process all the errors.• You can use <code>get diagnostics</code>, which is the target of the call, <code>perform</code>, or <code>go to</code> clause of a <code>whenever</code> statement, in the code.• You can use <code>get diagnostics</code> after a statement for which you want to retrieve informational messages. |
| See also | <code>whenever</code> |

include "filename"

| | |
|-------------|---|
| Description | Includes an external file in an Embedded SQL source file. |
|-------------|---|

Syntax `exec sql include "filename";`

Parameters *filename*

The name of the file to be included in the Embedded SQL source file containing this statement.

Examples

```
common.h:

/* This file contains definitions and
** declarations used in the file getinfo.c.
*/

#include <stdio.h>
#include "../common.h"
void    err_handler();
void    warning_handler();
exec sql include sqlca;
{
    exec sql begin declare section;
        CS_CHAR username[33], password[33], date[33];
    exec sql end declare section;

    exec sql whenever sqlerror call err_handler();
    exec sql whenever sqlwarning call warning_handler();
    exec sql whenever not found continue;

/*
** Copy the user name and password defined in common.h to
** the variables decalred for them in the declare section.
*/
strcpy (username, USER);
strcpy(password, PASSWORD);

printf("Today's date: %s\n", date);
...
}
void    err_handler()
{
...
}
void    warning_handler()
{
...
}
/* common.h */
#define USER "sa"
```



```
#define PASSWORD ""
=====

exec sql begin declare section;
      char      global_username[100];
      char      global_password[100];
exec sql end declare section;
```

getinfo.c

```
#include <common.h>
printf("uid?\n");
gets(global_username);
printf("password?\n");
gets(global_password);
```

do_connect.c

```
exec sql include "common.h";

exec sql connect :global_username
      identified by :global_password;
```

Usage

- The Embedded SQL precompiler processes the included file as though it were part of the Embedded SQL source file, recognizing all declare sections and SQL statements. The Embedded SQL precompiler writes the resulting host language source code into the generated file.
- Use the include path precompiler command line option to specify the directories to be searched for any included files. Refer to the *Open Client/Server Programmer's Supplement* for more information on precompiler command line options.
- Included files can be nested up to a maximum depth of 32 files.
- The include *"filename"* statement can be used anywhere.

See also

declare section

include sqlca

Description Defines the SQL Communications Area (SQLCA) in an Embedded SQL program.

Syntax `exec sql include sqlca;`

Examples

```
exec sql include SQLCA;

...
exec sql update t1 set c1      =      123 where c2      >
47;
if (sqlca.sqlcode      ==      0)
{
    printf("%d rows updated/n", sqlca.sqlerrd[2]);
}
else if (sqlca.sqlcode      ==      100)
{
    printf("No rows matched the query\n");
} else {
    printf("An error occured\n%s\n",
        sqlca.sqlerrm.sqlerrmc);
}
```

Usage

- The include sqlca statement can be used anywhere that host language declarations are allowed.

See also

begin declare section

include sqllda

Description

Defines the SQLDA structure in an Embedded SQL program.

Syntax

```
exec sql include sqllda;
```

Examples

```
exec sql include sqllda;

...

SQLDA *input_descriptor, *output_descriptor;
CS_SMALLINT small;
CS_CHAR      character[20];

input_descriptor = (SQLDA *)malloc(SYB_SQLDA_SIZE(3));
input_descriptor->sqllda_sqln = 3;
output_descriptor = (SQLDA *)malloc(SYB_SQLDA_SIZE(3));
output_descriptor->sqllda_sqln = 3;
```

Usage

- The include sqllda statement can be used anywhere that host language declarations are allowed.

initialize_application

| | |
|-------------|--|
| Description | Generates a call to set the application name on the global CS_CONTEXT handle. If precompiled with the -x option, it will also set the cs_config(CS_SET, CS_EXTERNAL_CONFIG, CS_TRUE) property. |
| Syntax | <pre>exec sql initialize_application [application_name "=" application_name];</pre> |
| Examples | <pre>exec sql include sqlca; main() { exec sql initialize_application application_name = :appname; /* ** The body of the main function goes here, ** including various Embedded SQL statements. */ ... /* The init statement must be the first ** embedded SQL statement in the program. */ exec sql exit; } /* end of main */</pre> |
| Usage | <ul style="list-style-type: none"> • <i>application_name</i> is either a literal string or a character variable containing the name of the application. • If <i>initialize_application</i> is the <i>first</i> Embedded SQL statement executed by an application, -x causes ct_init to use external configuration options to initialize the Client-Library part of the CS_CONTEXT structure. • If <i>initialize_application</i> is not the first Embedded SQL statement, ct_init does <i>not</i> pick up external configuration options. • Regardless of whether or not <i>initialize_application</i> is the first Embedded SQL statement, -x causes exec sql connect statements to use external configuration data. If -e is also specified, Sybase uses the server name as a key to the configuration data. If -e is not specified, then the application name (or DEFAULT) is used as the key to the configuration data. • If you specify -x and the application name, the following applies: <ul style="list-style-type: none"> • ct_init uses the application name to determine which section of the external configuration file to use for initialization. • The application name is passed to Adaptive Server as part of the connect statement. The application name is entered in the <i>sysprocesses.program_name</i> table. |

- If `-e` is specified without `-x`, then `ct_init` will use external configuration data when initializing, but every connection will use the server name as a key to the external configuration data. See the *Open Client/Server Programmer's Supplement* for information on command-line options.

See also

`exit`

open (dynamic cursor)

Description Opens a previously declared dynamic cursor.

Syntax `exec sql [at connection_name] open cursor_name
[row_count = size] [using {host_var_list |
descriptor descriptor_name |
sql descriptor descriptor_name}]`;

Parameters

cursor_name

Names a cursor that has been declared using the `declare cursor` statement.

size

The number of rows moved in a network roundtrip, not the number fetched into the host variable. The *size* argument can be either a literal or a declared host variable.

host_var_list

Names the host variables that contain the values for dynamic parameter markers.

descriptor

Identifies *descriptor_name* as a SQLDA structure.

sql descriptor

Identifies *descriptor_name* as a SQL descriptor.

descriptor_name

Names the dynamic descriptor that contains information about the dynamic parameter markers in a prepared statement.

Examples

```
exec sql begin declare section;  
    CS_CHAR      dyna_buf[128];  
    CS_CHAR      title_id[6];  
    CS_CHAR      lastname[40];  
    CS_CHAR      firstname[20];  
    CS_CHAR      phone[12];  
exec sql end declare section;
```

```
dyna_buf = "SELECT a.au_lname, a.au_fname, a.phone"
          + "FROM authors a, titleauthor t "
          + "WHERE a.au_id = t.au_id "
          + "AND t.title_id = ?";

exec sql prepare dyna_comm from :dyna_buf;

exec sql declare who_wrote cursor for dyna_comm;

printf("List authors for what title? ");
gets(title_id);

exec sql open who_wrote using :title_id;

while (TRUE){
    exec sql fetch who_wrote into
        :lastname, :firstname, :phone;
    if (sqlcode == 100) break;
    printf("Last name is %s\n", lastname,
        "First name is %s\n", firstname,
        "Phone number is %s\n", phone);
}

exec sql close who_wrote;
```

Usage

- open executes the statement specified in the corresponding declare cursor statement. You can then use the fetch statement to retrieve the results of the prepared statement.
- You can have any number of open cursors.
- The using clause substitutes host-variable or dynamic-descriptor contents for the dynamic parameter markers ("?",) in the select statement.

See also

close, declare, fetch, prepare

open (static cursor)

Description

Opens a previously declared static cursor. This statement can be used to open any static cursor, including one for a stored procedure.

Syntax

```
exec sql [at connection_name] open cursor_name
[row_count = size];
```

Parameters

cursor_name

The name of the cursor to be opened.

row_count

The number of rows moved in a network roundtrip, not the number fetched into the host variable.

size

The number of rows that are moved at the same time from Adaptive Server to the client. The client buffers the rows until they are fetched by the application. This parameter allows you to tune network efficiency.

Examples

```
exec sql begin declare section;
    char          b_titleid[tidsize+1];
    char          b_title[65];
    char          b_type[typesize+1];
exec sql end declare section;
    long          sqlcode;
    char          response[10];

    ...
exec sql declare titlelist cursor for
    select title_id, substring(title,1,64)
    from titles where type like :b_type;
    strcpy(b_type, "business");

exec sql open titlelist;

for (;;)
    exec sql fetch titlelist into :b_titleid,
        :b_title;
        if (sqlcode == 100)
            break;
        printf("    %-8s %s\n", b_titleid, b_title);
        printf("update/delete? ");
        gets(response);
        if (!strncasecmp(response,"u",1))
        {
            printf("enter the new titleid\n>");
            gets(b_titleid);
            exec sql update titles
                set title_id = :b_titleid
                where current of titlelist;
        }
        else if (!strncasecmp(response,"d",1))
        {
            exec sql delete from titles
                where current of titlelist;
        }
    }
exec sql close titlelist;
```

Usage

- open executes the select statement given by the declare cursor statement and prepares results for the fetch statement.
- You can have an unlimited number of open cursors.
- A static cursor must be opened only in the file where the cursor is declared. The cursor can be closed in any file.
- The values of host variables embedded in the declare cursor statement are taken at open time.
- When specifying *cursor_name*, you can use the name of a deallocated static cursor. If you do, the precompiler declares and opens a new cursor having the same name as that of the deallocated cursor. Thus, the precompiler does not reopen the deallocated cursor but instead creates a new one. The results sets for the two cursors can differ.

prepare

Description

Declares a name for a dynamic SQL statement buffer.

Syntax

```
exec sql [at connection_name] prepare statement_name from {:host_variable
| "string";
```

Parameters

statement_name

An identifier used to reference the statement. *statement_name* must uniquely identify the statement buffer and must conform to the SQL identifier rules for naming variables. The *statement_name* can also be a *host_variable* string containing a valid SQL identifier. *statement_name* can be up to 30 characters.

host_variable

A character-string host variable that contains an executable SQL statement. Place dynamic parameter markers (“?”) anywhere in the select statement where a host variable value will be substituted.

string

A literal string that can be used in place of *host_variable*.

Examples

```
exec sql begin declare section;
      CS_CHAR      dyn_buffer[128];
      CS_CHAR      state[2];
exec sql end declare section;
```

```
-- The select into table_name statement returns no
```

```
-- results to the program, so it does not
-- need a cursor.

dyn_buffer = "select * into #work from authors"
            + "where state = ?";

printf("State? ");
gets(state);

exec sql prepare make_work from :dyn_buffer;
exec sql execute make_work using :state;
```

Usage

- In the current implementation, Sybase creates a temporary stored procedure for a dynamic SQL statement stored in a character string literal or host variable.
- `prepare` sends the contents of *host_variable* to the Adaptive Server to convert into a temporary stored procedure. This temporary stored procedure remains in tempdb on Adaptive Server until the statement is deallocated or the connection is disconnected.
- The scope of *statement_name* is global to your program but local to the connection *connection_name*. The statement persists until the program either deallocates it or closes the connection.
- `prepare` is valid with Dynamic SQL methods 2, 3, and 4.
- With method 2, (`prepare` and `execute`), an `execute` statement substitutes values from host variables, if any, into the prepared statement and sends the completed statement to Adaptive Server. If there are no host variables to substitute and no results, you can use `execute immediate`, instead.
- With method 3, `prepare` and `fetch`, a `declare cursor` statement associates the saved `select` statement with a cursor. An `open` statement substitutes values from host variables, if any, into the `select` statement and sends the result to Adaptive Server for execution.
- With methods 2, 3, and 4, `prepare` and `fetch` with parameter descriptors, the dynamic parameter descriptors, represented by question marks ("?"), indicate where host variables will be substituted.
- A prepared statement must be executed on the same connection on which it was prepared. If the prepared statement is used to declare a cursor, all operations on that cursor use the same connection as the prepared statement.
- The statement in *host_variable* can contain dynamic parameter markers that indicate where to substitute values of host variables into the statement.

See also `declare cursor`, `execute`, `execute immediate`, `deallocate prepare`

rollback

| | |
|-------------|---|
| Description | Rolls a transaction back to a savepoint inside the transaction or to the beginning of the transaction. |
| Syntax | <pre>exec sql [at <i>connection_name</i>] rollback [transaction tran work] [<i>transaction_name</i> <i>savepoint_name</i>];</pre> |
| Parameters | <p><code>transaction</code> <code>trans</code> <code>work</code></p> <p>The keywords <code>transaction</code>, <code>trans</code>, and <code>work</code> are interchangeable in the <code>rollback</code> statement, but only <code>work</code> is ANSI-compliant.</p> <p><i>transaction_name</i></p> <p>The name of the transaction being rolled back.</p> <p><i>savepoint_name</i></p> <p>The name assigned to the savepoint in a <code>save transaction</code> statement. If you omit <i>savepoint_name</i>, SQL Server rolls back the entire transaction.</p> |
| Examples | <pre>abort_tran: exec sql whenever sqlerror continue: exec sql at connect2 rollback transaction; exec sql at connect1 rollback transaction; goto try_update;</pre> |
| Usage | <ul style="list-style-type: none"> • This reference page mainly describes aspects of the Transact-SQL <code>rollback</code> statement that differ when used with Embedded SQL. See the <i>Adaptive Server Enterprise Reference Manual</i> for more information about the <code>rollback</code> statement, savepoints, and Transact-SQL transaction management. • Transaction names and savepoint names must conform to the Transact-SQL rules for identifiers. • Transaction names and savepoints are Transact-SQL extensions; they are not ANSI-compliant. Do not use a transaction name or savepoint name with the ANSI-compliant keyword <code>work</code>. |
| See also | <code>begin transaction</code> , <code>commit</code> |

select

| | |
|-------------|--|
| Description | Retrieves rows from database objects. |
| Syntax | <pre>exec sql [at <i>connect_name</i>] select <i>select_list</i> into <i>destination</i> from <i>table_name</i>...;</pre> |
| Parameters | <p><i>select_list</i></p> <p>Same as <i>select_list</i> in the Transact-SQL select statement, except that the <i>select_list</i> cannot perform variable assignments in Embedded SQL.</p> <p><i>destination</i></p> <p>A table or a series of one or more Embedded SQL host variables. Each host variable must first be defined in a previous declare section. <i>Indicator variables</i> can be associated with the host variables.</p> |
| Examples | <pre>/* This example retrieves columns from a ** single row of the authors table and ** stores them in host variables. Because the ** example's select statement cannot return more ** than one row, no cursor is needed. */ exec sql begin declare section; character last[40]; character first[20]; character phone[12]; character id[11]; exec sql end declare section; printf("Enter author id: "); gets(id); exec sql select au_lname, au_fname, phone into :last, :first, :phone from authors where au_id = :id; if (sqlcode != 100) { print "Information for Author ", id, ":"; print last, first, phone; } else { print "Could not locate author ", id; };</pre> |

| | |
|----------|---|
| Usage | <ul style="list-style-type: none">• This reference page mainly describes aspects of the Transact-SQL select statement that differ when the statement is used in Embedded SQL. See the <i>Adaptive Server Enterprise Reference Manual</i> for more information about the select statement.• The compute clause of the Transact-SQL select statement cannot be used in Embedded SQL programs.• Host variables in a select statement are input variables only, except in the statement's into clause. Host variables in the into clause are output variables.• Previously declared input host variables can be used anywhere in a select statement that a literal value or Transact-SQL variable is allowed. Indicator variables can be associated with input host variables to specify null values.• If a select statement returns more than one row, each host variable in the statement's into clause must be an array with enough space for all the rows. Otherwise, you must use a cursor to bring the rows back one at a time. |
| See also | declare cursor |

set connection

| | |
|-------------|--|
| Description | Causes the specified existing connection to become the current connection. |
| Syntax | <code>set connection {<i>connection_name</i> DEFAULT};</code> |
| Parameters | <p><i>connection_name</i></p> <p>The name of an existing connection that you want to become the current connection.</p> <p>default</p> <p>Specifies that the unnamed default connection is to become the current connection.</p> |
| Examples | <pre>exec sql connect "ME" at connect1 using "SERVER1"; exec sql connect "ME" at connect2 using "SERVER2"; /* The next statement executes on connect2. */ exec sql select userid() into :myid; exec sql set connection connect1; /* The next statement executes on connect1. */</pre> |

Usage

```
exec sql select count(*)from t1;
```

- The set connection statement specifies the current connection for all subsequent SQL statements, except those preceded by the `exec sql` clause at.
- A set connection statement remains in effect until you choose a different current connection by using the set connection statement again.

See also

at connection_name, connect

set descriptor

Description

Inserts or updates data in a SQL descriptor.

For a list of possible SQL descriptor datatypes, see Table 10-5 on page 172.

Syntax

```
exec sql set descriptor descriptor_name  
{count = host_variable} |  
{value item_number {item_name =  
:host_variable}{,...];
```

Parameters

descriptor_name

The name of the SQL descriptor that contains information about the dynamic parameter markers in a prepared statement.

count

The number of dynamic parameter specifications to be described.

host_variable

A host variable defined in a declare section.

item_number

Represents the *n*th occurrence of either a dynamic parameter marker or a select column.

item_name

Represents the attribute information of either a dynamic parameter marker or a select list column. Table 10-4 lists the values for *item_name*.

Table 10-4: Values for item_name

| Value | Description |
|-------------|---|
| <i>data</i> | Value for the dynamic parameter marker or target associated with the specified SQL descriptor. If indicator is negative, this field is undefined. |

| Value | Description |
|------------------|--|
| <i>length</i> | The length, in characters, of the dynamic parameter marker of target for the specified SQL descriptor. |
| <i>precision</i> | An integer specifying the total number of digits of precision for the CS_NUMERIC variable. |
| <i>scale</i> | An integer specifying the total number of digits after the decimal point for the CS_NUMERIC variable. |
| <i>type</i> | The datatype of this column (item number) in the row. For values, see “SQL descriptor datatype codes” on page 172. |

Examples

```

exec sql prepare get_royalty
from "select royalty from roysched
where title_id = ? and lorange <= ? and
hirange > ?";

exec sql allocate descriptor roy_desc with max 3;

exec sql set descriptor roy_desc
value 1 data = :tid;
exec sql set descriptor roy_desc
value 2 data = :sales;
exec sql set descriptor roy_desc
value 3 data = :sales;

exec sql execute get_royalty into :royalty
using sql descriptor roy_desc;

```

Usage

- An Embedded SQL program passes attribute and value information to Client-Library, which holds the data in the specified SQL descriptor until the program issues it a request to execute a statement.

See also

allocate descriptor, describe input, describe output, execute, fetch, get descriptor, open(dynamic cursor)

update**Description**

Modifies data in rows of a table.

| | |
|------------|--|
| Syntax | <pre>exec sql [at <i>connection_name</i>] update <i>table_name</i> set [<i>table_name</i>] <i>column_name1</i> = {<i>expression1</i> NULL (<i>select_statement</i>)} [, <i>column_name2</i> = {<i>expression2</i> NULL (<i>select_statement</i>)}]... [<i>from table_name</i> [, <i>table_name</i>]... [where {<i>search_conditions</i> current of <i>cursor_name</i>}]];</pre> |
| Parameters | <p><i>table_name</i></p> <p>The name of a table or view, specified in any format that is valid for the update statement in Transact-SQL.</p> |
| Examples | <pre>exec sql begin declare section; CS_CHAR store_name[40]; CS_CHAR disc_type[40]; CS_INT lowqty; CS_INT highqty; CS_FLOAT discount; exec sql end declare section; CS_CHAR answer[1]); exec sql declare update_cursor cursor for select s.stor_name, d.discounttype, d.lowqty, d.highqty, d.discount from stores s, discounts d where d.stor_id = s.stor_id; exec sql open update_cursor; exec sql whenever not found goto alldone; while (TRUE) { exec sql fetch update_cursor into :store_name, :disc_type, :lowqty, :highqty, discount; print store_name, disc_type, lowqty, highqty, discount; printf("New discount? "); gets(discount); exec sql update discounts set discount = :discount where current of update_cursor; }</pre> |

```
alldone:
exec sql close update_cursor;
exec sql disconnect all;
```

Usage

- This reference page mainly describes aspects of the Transact-SQL update statement that differ when the statement is used in Embedded SQL. See the *Adaptive Server Enterprise Reference Manual* for more information about the update statement.
- Host variables can appear anywhere in an expression or in any where clause.
- You can use the where clause to update selected rows in a table. Omit the where clause to update all rows in the table. Use where current of *cursor_name* to update the current row of an open cursor.
- When where current of *cursor_name* is specified, the statement must be executed on the connection specified in the open cursor statement. If the at *connection_name* clause is used, it must match the open cursor statement.

See also

close, delete cursor, fetch, open, prepare

whenever

Description

Specifies an action to occur whenever an executable SQL statement causes a specified condition.

Syntax

```
exec sql whenever {sqlerror | not found | sqlwarning}
{continue | go to label | goto label |
stop | call routine_name [args]};
```

Parameters

sqlerror

Specifies an action to take when an error is detected, such as a syntax error returned to the Embedded SQL program from SQL Server.

not found

Specifies an action to take when a fetch or select into statement retrieves no data or when a searched update or delete statement affects no rows.

sqlwarning

Specifies an action to take when a warning is received; for example, when a character string is truncated.

continue

Take no action when the condition occurs.

go to | goto

Transfer control to the program statement at the specified *label*.

label

A host language statement label, such as a C label.

stop

Terminate the Embedded SQL program when the condition occurs.

call

Transfer control to a callable routine in the program, such as a user-defined function or subroutine.

routine_name

A host language routine that can be called. The routine must be able to be called from the source file that contains the *whenever* statement. You may need to declare the routine as external to compile the Embedded SQL program.

args

One or more arguments to be passed to the callable routine, using the parameter-passing conventions of the host language. The arguments can be any list of host variables, literals, or expressions that the host language allows. A space character should separate each argument from the next.

Examples

```
exec sql whenever sqlerror call err_handler();
exec sql whenever sqlwarning call warn_handler();

long SQLCODE;
exec sql begin declare section;
    CS_CHAR      lastname[40];
    CS_CHAR      firstname[20];
    CS_CHAR      phone[12];
exec sql end declare section;

exec sql declare au_list cursor for
    select au_lname, au_fname, phone
    from authors
    order by au_lname;

exec sql open au_list;

exec sql whenever not found go to list_done;
```



```
while (TRUE){
    exec sql fetch au_list
        into :lastname, :firstname, :phone;
    printf("Lastname is: %s\n", lastname,
"Firstname is: %s\n", firstname,
        "Phone number is: %s\n", phone;
}
list_done:
exec sql close au_list;
exec sql disconnect current;
```

Usage

- The whenever statement causes the Embedded SQL precompiler to generate code following each executable SQL statement. The generated code includes the test for the condition and the host language statement or statements that carry out the specified action.
- The Embedded SQL precompiler generates code for the SQL statements that follow the whenever statement in the source file, including SQL statements in subroutines that are defined in the same source file.
- Use whenever ...continue to cancel a previous whenever statement. The continue action causes the Embedded SQL precompiler to ignore the condition. To prevent infinite loops, use whenever ...continue in an error handler before executing any Embedded SQL statements.
- When you use whenever ...go to *label*, *label* must represent a valid location to resume execution. In C, for example, *label* must be declared in any routine that has executable SQL statements within the scope of the whenever statement. C does not allow a goto statement to jump to a label declared in another function.
- If you have a whenever statement in your program but you have not declared SQLCA or SQLSTATE status variables, the Embedded SQL precompiler assumes that you are using the SQLCODE variable. Be sure that SQLCODE is declared. Otherwise, the generated code will not compile.

SQL descriptor codes

The following table pertains to the SQL descriptor used for dynamic SQL statements. Sybase's use of dynamic SQL values conforms to the ANSI/ISO 185-92 SQL-92 standards. For more information, see the appropriate ANSI/ISO documentation.

Table 10-5: SQL descriptor datatype codes

| ANSI SQL datatype | Code |
|--------------------------|-------------|
| bit | 14 |
| character | 1 |
| character varying | 12 |
| date, time | 9 |
| decimal | 3 |
| double precision | 8 |
| float | 6 |
| integer | 4 |
| numeric | 2 |
| real | 7 |
| smallint | 5 |

| Sybase-defined datatype | Client-Library code |
|--------------------------------|----------------------------|
| smalldatetime | -9 |
| money | -10 |
| smallmoney | -11 |
| text | -3 |
| image | -4 |
| tinyint | -8 |
| binary | -5 |
| varbinary | -6 |
| long binary | -7 |
| longchar | -2 |

Table 10-6: SQL descriptor identifier values

| Value | Description |
|------------------------|--|
| <i>type</i> | The datatype of this column (item number) in the row. For values, see “SQL descriptor datatype codes” on page 172. |
| <i>length</i> | The length, in characters, of the dynamic parameter marker of target for the specified SQL descriptor. |
| <i>returned_length</i> | The length of char types of the values from the select column list. |
| <i>precision</i> | An integer specifying the total number of digits of precision for the CS_NUMERIC variable. |

| Value | Description |
|------------------|---|
| <i>scale</i> | An integer specifying the total number of digits after the decimal point for the CS_NUMERIC variable. |
| <i>nullable</i> | Equals 0 if the dynamic parameter marker can accept a null value; otherwise, equals 1. |
| <i>indicator</i> | Value for the indicator parameter associated with the dynamic parameter marker or target. |
| <i>data</i> | Value for the dynamic parameter marker or target associated with the specified SQL descriptor. If indicator is negative, this field is undefined. |
| <i>name</i> | The name of the specified SQL descriptor containing information about the dynamic parameter markers. |

Open Client/Server Configuration File

Open Client/Server applications can easily be configured using the Open Client/Server configuration file. By default, the file is named *ocs.cfg* and is located in the *\$\$SYBASE/config* directory. This chapter describes how the configuration file can be used with Embedded SQL.

| Topic | Page |
|--|------|
| Purpose of the open Client/Server configuration file | 175 |
| Accessing the configuration functionality | 175 |
| Default settings | 176 |
| Syntax for the Open Client/Server configuration file | 177 |
| Sample programs | 179 |

Purpose of the open Client/Server configuration file

The Open Client/Server configuration file provides a single location where all Open Client/Server application connections can be configured. Using the configuration file simplifies the tasks of establishing configuration standards and managing configuration changes.

Accessing the configuration functionality

This feature is available through two new command-line options and the `initialize_application` statement.

- -x – this option allows for external configuration. The application needs to initialize an application with a name. The Open Client/Server configuration file will have a section with this application name. Under this section, place all properties that need to be set for this application. The -x option is useful only when used with `initialize_application`. If initializing is not done, and the -x option is used, the default section of the configuration file will be accessed.
- -e – this option allows us to configure by SERVER NAME. No call to `initialize_application` is required. The server name will be used as a key to look up in the configuration file for properties to be set the section defined by the server name. This will allow users to associate connection names with specific connection properties.

Note If `INITIALIZE_APPLICATION` is not the first Embedded SQL statement to be executed, external configuration properties will not be set. If it is the first Embedded SQL statement to be executed, then the external configuration options will be used for initialization.

Default settings

The following is the Open Client/Server configuration file with default settings. You can customize the file as needed.

[DEFAULT]

;This is the default section loaded by applications that use the external configuration feature, but which do not specify their own application name. Initially this section is empty. Defaults from all properties will be the same as earlier releases of Open Client libraries.

[ANSI_ESQL]

;This section defines configuration which an ANSI conforming Embedded SQL application should use to get ANSI-defined behavior from SQL Servers and Open Client libraries. This set of configuration properties matches the set which earlier releases of Embedded SQL (version 10.0.x) automatically set for applications during execution of a CONNECT statement.

```
CS_CAP_RESPONSE=CS_RES_NOSTRIPBLANKS
CS_EXTRA_INF=CS_TRUE
```

```
CS_ANSI_BINDS=CS_TRUE
CS_OPT_ANSINULL=CS_TRUE
CS_OPT_ANSIPERM=CS_TRUE
CS_OPT_STR_RTRUNC=CS_TRUE
CS_OPT_ARITHABORT=CS_FALSE
CS_OPT_TRUNCIGNORE=CS_TRUE
CS_OPT_ISOLATION=CS_OPT_LEVEL3
CS_OPT_CHAINXACTS=CS_TRUE
CS_OPT_CURCLOSEONXACT=CS_TRUE
CS_OPT_QUOTED_IDENT=CS_TRUE
;End of default sections
```

Syntax for the Open Client/Server configuration file

The syntax for the Open Client/Server configuration file will match the existing syntax for Sybase localization and configuration files supported by CS-Library with minor variations.

Syntax

- `;` – Signifies a comment line.
- `[section_name]` – Section names are wrapped in square brackets. The Open Client/Server configuration file comes with sections named `DEFAULT` and `ANSI_ESQL`. The application name will be used as the section name for an application that has been compiled with the `-x` option. For an application that has been compiled with the `-e` option, the server name will be used for the section name. Any name can be used as a section name for those sections that contain settings that will be used in multiple sections. The following example shows a section arbitrarily named, `GENERIC`, and how that section is included in other sections:

```
[GENERIC]
    CS_OPT_ANSINULL=CS_TRUE

[APP_PAYROLL]
    include=GENERIC
    CS_CAP_RESPONSE=CS_RES_NOSTRIPBLANKS

[APP_HR]
    include=GENERIC
    CS_OPT_QUOTED_IDENT=CS_TRUE
```

- `entry_name=entry_value`
 - Entry values can be anything: integers, strings and so on. If an entry value line ends with `"\<newline>` the entry value continues to the next line.
 - White spaces are trimmed from the beginning and end of entry values.
 - If white spaces are required at the beginning or end of an entry value, wrap them in double quotes.
 - An entry that begins with a double quote must end with a double quote. Two double quote characters in a row within a quoted string represent a single double quote in the value string. If a newline is encountered within double quotes, it is considered to be literally part of the value.
 - Entry names and section names can consist of alphabetic characters (both upper and lower case), the digits 0-9, and punctuation characters. The first letter MUST be alphabetic.
 - Entry and section names are case sensitive.
 - `Include=earlier_section`

If a section contains the entry `include`, then the entire contents of that previously defined section are considered to be replicated within this section. In other words, the properties defined in the previous section are inherited by this section.

Note that the included section must have been defined prior to it being included in another section. This allows the configuration file parsing to happen in a single pass and eliminates the need to detect recursive included directives.

If an included section in turn includes another section, the order of entry values is defined by a “depthfirst” search of the included sections.

Sections cannot include a reference to themselves. In other words, recursion is not possible because you must include a previously defined section—you cannot include the section being defined.

All direct entry values defined in a given section supersede any values which may have been included from another section. In the following example, `CS_OPT_ANSINULL` will be set to false in the `APP.PAYROLL` application. Note that the position of the `include` statement does not affect this rule.


```
[GENERIC]
    CS_OPT_ANSINULL=CS_TRUE

[APP_PAYROLL]
    CS_OPT_ANSINULL=CS_FALSE
    include=GENERIC
```

Sample programs

Consider the following scenario: An Embedded SQL program defines a cursor to retrieve rows from the titles table in the pubs2 database. The WHERE clause uses non-ANSI standard NULL checking. To clarify, IS NULL and IS NOT NULL are ANSI standards which is the default used by Embedded SQL programs, whereas an Embedded SQL program wishing to use = NULL or != NULL will need to turn OFF ANSINULL behavior and use Transact-SQL syntax instead. If you wanted to make comparisons with NULLs in Transact-SQL syntax in Embedded SQL prior to version 11.1, you would need to make the call:

```
EXEC SQL set ansinull off;
```

In the following example, no change is made to the Embedded SQL code, but the desired behavior is attained by setting appropriate properties in the Open Client/Server configuration file.

There are two versions of the same program listed below. One is to be used with the -e option and the other with the -x option.

Embedded SQL/C sample makefile on Windows NT

The *libcobct.lib* and *mfmts32.lib* libraries do not need to be included in the Embedded SQL/C sample makefile.

You must change the CC_INCLUDE variable in the makefile to:

```
CC_INCLUDES= -I$(SYBASE)\include
```

Note On Windows NT, the command to compile all the example programs is `nmake`, not `make`.

Embedded SQL/C sample programs

Before you build Embedded SQL/C sample programs on UNIX platforms, you must:

- Set execute permission on the *sybopts.sh* file for the file's owner:

`chmod u+x sybopts.sh`
- If you have not already done so, include the current directory in the search path:

```
setenv PATH .:$PATH
```

Embedded SQL program version for use with the -x option

```
/* Program name: ocs_test.cp
**
** Description : This program declares a cursor which retrieve rows
** from the 'titles' table based on condition checking for NULLS
** in the NON-ANSI style.
** The program will be compiled using the -x option which will
** use an external configuration file (ocs.cfg) based on the
** name of the application. The name of the application is
** defined at the time of INITIALIZING the application. Note that
** this is a new 11.x feature too.
*/

#include <stdio.h>

/* Declare the SQLCA */
EXEC SQL INCLUDE sqlca;

EXEC SQL BEGIN DECLARE SECTION;
    /* storage for login name and password */
    CS_CHARusername[30], password[30];
    CS_CHARtitle_id[7], price[30];
EXEC SQL END DECLARE SECTION;

/*
** Forward declarations of the error and message handlers and
** other subroutines called from main().
*/
void    error_handler();
void    warning_handler();
```

```
int main()
{
    int i=0 ;

    EXEC SQL WHENEVER SQLERROR CALL error_handler();
    EXEC SQL WHENEVER SQLWARNING CALL warning_handler();
    EXEC SQL WHENEVER NOT FOUND CONTINUE ;

    /*
    ** Copy the user name and password defined in sybsqllex.h to
    ** the variables declared for them in the declare section.
    */

    strcpy(username, "sa");
    strcpy(password, "");

    EXEC SQL INITIALIZE_APPLICATION APPLICATION_NAME = "TEST1";

    EXEC SQL CONNECT :username IDENTIFIED BY :password ;
    EXEC SQL USE pubs2 ;

    EXEC SQL DECLARE title_list CURSOR FOR
    SELECT title_id, price FROM titles
        WHERE price != NULL;

    EXEC SQL OPEN title_list ;
    for ( ;; )
    {
        EXEC SQL FETCH title_list INTO
            :title_id, :price;
        if ( sqlca.sqlcode == 100 )
        {
            printf("End of fetch! \n");
            break;
        }
        printf("Title ID : %s\n", title_id );
        printf("Price      : %s\n", price) ;
        printf("Please press RETURN to continue .. ");
        getchar();
        printf("\n\n");
    }
    EXEC SQL CLOSE title_list;
    exit(0);
}
```

```
void error_handler()
{
    . . .}

void warning_handler()
{
    . . .}
```

Note Precompiler option to set in the makefile: `cpre -x`.

The following is a sample configuration file for the preceding program:

```
[DEFAULT]
;

[TEST1]
;This is name of the application set by INITIALIZE_APPLICATION. ;Therefore this
is the section that will be referred to a runtime.

CS_OPT_ANSINULL=CS_FALSE

;The above option will enable comparisons of nulls in the NON-ANSI
;style.
```

Same Embedded SQL program with the -e option

```
/* Program name: ocs_test.cp
**
** Description : This program declares a cursor which retireves rows
** from the 'titles' table based on condition checking for NULLS
** in the NON-ANSI style.
** The program will be compiled using the -e option which will
** use the server name that the application connects to, as the
** corresponding section to look up in the configuration file.
*/

#include <stdio.h>

/* Declare the SQLCA */
EXEC SQL INCLUDE sqlca;

EXEC SQL BEGIN DECLARE SECTION;
    /* storage for login name and password */
```

```
        CS_CHARusername[30], password[30];
        CS_CHARtitle_id[7], price[30];
EXEC SQL END DECLARE SECTION;

/*
** Forward declarations of the error and message handlers and
** other subroutines called from main().
*/
void    error_handler();
void    warning_handler();

int main()
{
    int i=0 ;

    EXEC SQL WHENEVER SQLERROR CALL error_handler();
    EXEC SQL WHENEVER SQLWARNING CALL warning_handler();
    EXEC SQL WHENEVER NOT FOUND CONTINUE ;

    /*
    ** Copy the user name and password defined in sybsqlx.h to
    ** the variables declared for them in the declare section.
    */

    strcpy(username, "sa");
    strcpy(password, "");

    EXEC SQL CONNECT :username IDENTIFIED BY :password ;
    EXEC SQL USE pubs2 ;

    EXEC SQL DECLARE title_list CURSOR FOR
        SELECT title_id, price FROM titles
            WHERE price != NULL;

    EXEC SQL OPEN title_list ;
    for ( ;; )
    {
        EXEC SQL FETCH title_list INTO
            :title_id, :price;
        if ( sqlca.sqlcode == 100 )
        {
            printf("End of fetch! \n");
            break;
        }
        printf("Title ID : %s\n", title_id );
    }
}
```

```
        printf("Price      : %s\n", price) ;
        printf("Please press RETURN to continue .. ");
        getchar();
        printf("\n\n");
    }
    EXEC SQL CLOSE title_list;
    exit(0);

}

void error_handler()
{
    . . .}
```

Note Precompiler option to set in the makefile: `cpre -e`.

The following is a sample configuration file for the preceding program:

```
[DEFAULT]
;

[SYBASE]
;This is name of the server that the application connect to. Therefore
;this is the section that will be referred to a runtime.
;
CS_OPT_ANSINULL=CS_FALSE
;The above option will enable comparisons of nulls in the NON-ANSI
;style.
```

The above configuration files have been vastly simplified. A typical Open Client/Server configuration file would be in the following format:

```
[DEFAULT]
;
[ANSI_ESQL]
CS_CAP_RESPONSE=CS_RES_NOSTRIPBLANKS
CS_EXTRA_INF=CS_TRUE
CS_ANSI_BINDS=CS_TRUE
CS_OPT_ANSINULL=CS_TRUE
CS_OPT_ANSIPERM=CS_TRUE
CS_OPT_STR_RTRUNC=CS_TRUE
CS_OPT_ARITHABORT=CS_FALSE
CS_OPT_TRUNCIGNORE=CS_TRUE
CS_OPT_ISOLATION=CS_OPT_LEVEL3
CS_OPT_CHAINXACTS=CS_TRUE
CS_OPT_CURCLOSEONXACT=CS_TRUE
```

```
CS_OPT_QUOTED_IDENT=CS_TRUE
;
;The following is a sample section showing how to alter standard
;configuration:
;
[RELEVANT_SECTION_NAME]
;
;Use most of the ANSI properties defined above,
;
include=ANSI_ESQL

;but override some default properties

CS_OPT_ANSINULL=CS_TRUE      ; enable non-ansi style null comparisons
CS_OPT_CHAINXACTS=CS_FALSE ; run in autocommit mode
```

Summary

The Open Client/Server configuration file serves as a single location where environment settings can be managed for multiple Embedded SQL applications. The default name of this file is *ocs.cfg*, and is located in the *\$SYBASE/config* directory. The use of the configuration file is regulated by the use of the *-x* and *-e* precompiler options. The syntax used for modifying the Open Client/Server configuration file matches the existing syntax for Sybase localization and configuration files supported by CS-Library with minor variations.

Precompiler Warning and Error Messages

The Embedded SQL precompiler generates the informational, warning, and error messages in different tables.

| Topic | Page |
|------------------------------------|------|
| Command line option messages | 188 |
| First pass parser messages | 190 |
| Second pass parser messages | 193 |
| Code generation messages | 194 |
| FIPS flag messages | 195 |
| Internal error messages | 195 |
| Sybase and Client-Library messages | 196 |
| Runtime messages | 197 |

Each table contains four fields.

- “Message_ID” lists the identification code of the message you may receive.
- “Message Text” lists the online text associated with the message you may receive.
- “Severity” lists the seriousness of the message you may receive. A message can be:
 - Information – No error or warning was detected, and the precompiler succeeded. The message is purely informational.
 - A warning – A noncritical error was detected, but the program precompiled.
 - Severe – An error occurred, and no code was generated. The precompilation failed.
 - Fatal – A severe error occurred from which the precompiler cannot recover. No further attempt will be made to process your files. Precompiler exits.

- The fourth field, “Fix,” suggests a means of correcting the situation that caused the error or warning.

Table A-1: Command line option messages

| Message ID | Message text | Severity | Fix |
|------------------------|---|-------------|---|
| M_COMPAT_INFO | Compatibility mode specified. | Information | No fix required. |
| M_DUPOPT | Duplicate command line option specified. | Severe | Do not duplicate the options specified on the command line. Remove the offending duplicate option. |
| M_EXCFG_OVERRIDE | The switch value will have no effect because the external switch value has been specified. | Warning | When you use an external configuration file, you may override configuration options set on the command line. Choose one means of setting options. |
| M_INVALID_COMPAT | Unrecognized compatibility mode specified. | Information | No fix required. |
| M_INVALID_FILE_FMT | Invalid character in file value at line value. | Severe | Check that characters in the input file are valid. Also check that you have correctly set the character set you want to use. |
| M_INVALID_FIPLEVEL | Invalid FIPS level specified. | Severe | Legal values are SQL92E and SQL89. |
| M_INVALID_SYNTAX | Invalid syntax checking level specified. | Severe | Legal values are NONE, SYNTAX, SEMANTIC. |
| M_INVLD_HLANG | Host Language specified is invalid. | Severe | Valid options are ANSI_C, KR_C. |
| M_INVLD_OCLIB_VER | The Open Client Client-Library version is invalid. | Severe | The correct version string is "CS_VERSION_110". |
| M_INVOPT | Option is invalid. | Severe | Invalid option specified. Substitute the correct value. |
| M_LABEL_SYNTAX | Security label is improperly specified; the proper format is 'labelname=labelvalue'. | Severe | Use the allowed syntax. |
| M_MSGINIT_FAIL | Error initializing localized error messages. | Warning | Verify that the Sybase installation is complete and that there is a valid entry for the LANG variable in the <i>locales.dat</i> file. |
| M_MULTI_IN_USE_DEF_OUT | When precompiling multiple input files, you cannot specify output (Listing, SQL, or Language) file names. | Severe | Remove all -G, -L, and -O flags from the command line or precompile the files one at a time. |

| Message ID | Message text | Severity | Fix |
|------------------------|---|----------|--|
| M_NO_INPUT_FILE | Error: No input file is specified to be precompiled. | Severe | Specify an input file for precompilation. Note This error may occur if you precede the input file name with a flag (such as -G, for generate stored procedures) which takes an optional argument. To fix, put another flag in front of the input file name. For example, replace <code>cpre -G file.pc</code> with <code>cpre -G -Ccompilername</code> . |
| M_OPEN_INCLUDE | Unable to open the specified include file <i>file</i> . | Severe | The specified file is either not in the path or is missing the required read permission. Specify the path with the -I flag and verify the read permission. |
| M_OPEN_INPUT | Unable to open the specified input file <i>file</i> . | Severe | Check the validity of the path and filename specified. If the filename extension is not provided, the precompiler searches for the default extension. |
| M_OPEN_ISQL | Unable to open the specified ISQL file <i>file</i> . | Severe | Check the validity of the isql filename (the file in which the stored procedures are written). Verify that you have the write permission in the directory where the file is being created. |
| M_OPEN_LIST | Unable to open the specified listing file <i>file</i> . | Severe | Check the validity of the listing filename. Verify that you have write permission in the directory where the file is being created. |
| M_OPEN_TARGET | Unable to open the specified target file <i>file</i> . | Severe | Check the validity of the output filename. Verify that you have write permission in the directory where the file is being created. |
| M_OPT_MUST_BE_PROVIDED | Option <i>value</i> must be provided. | Severe | Provide a value for option. |

| Message ID | Message text | Severity | Fix |
|-------------------|--|----------|--|
| M_OPT_REINIT | Warning: <i>value</i> switch initialized multiple times. | Warning | The specified switch has been initialized multiple times. The second and subsequent values are ignored. |
| M_PATH_OFL | Error: Max allowed paths for "INCLUDE" files is 64 (OVERFLOWED). | Severe | The maximum allowed paths on the command line have been exceeded. Reduce the number of directories from which the "INCLUDE" files are fetched. |
| M_STATIC_HV_CNAME | Static cursor names cannot be host-variables; <i>line</i> . | Severe | Replace the host variable with a SQL identifier. |
| M_UNBALANCED_DQ | Unbalanced quotes in delimited identifier. | Severe | Balance the quote. |

Table A-2: First pass parser messages

| Message ID | Message text | Severity | Fix |
|----------------------|---|----------|--|
| M_64BIT_INT | Warning: 64 bit integer host variables are not supported. Line <i>value</i> . | Warning | Use some other host variable type (float, numeric, or 32-bit integer) and, if necessary, copy the value between the host variable and the 64-bit program variable. |
| M_BLOCK_ERROR | Non-matching block terminator in <i>value</i> at line: <i>value</i> . | Severe | Correct your program syntax. |
| M_CONST_FETCH | Error: Attempted fetch into CONST storage class variable <i>value</i> . | Severe | You cannot fetch into a constant type. To fetch the value, remove the constant qualifier in its declaration. |
| M_DUP_HV | Duplicate host variable in <i>file</i> at line <i>line</i> . | Severe | Another host variable with the same name is already declared in the same block. Verify that each variable within a given block has a unique name. |
| M_DUP_STRUNION | Duplicate structure/union in <i>file</i> at line <i>line</i> . | Severe | Another structure with the same name is already being declared in the same block. Verify that each variable within a given block has a unique name. |
| M_IDENT_OR_STRINGVAR | Error: item must be a SQL-identifier or a string-type variable. | Severe | Verify that the connection, cursor, or statement name is of type string or SQL identifier. |

| Message ID | Message text | Severity | Fix |
|----------------------|--|----------|--|
| M_ILL_LITERAL_USAGE | Error: Use of literal parameters to an RPC with an OUTPUT qualifier is not legal. | Severe | Do not use a literal as an OUTPUT parameter to a stored procedure. |
| M_ILL_PARAM_MODE | Error: Mixing calling modes in an rpc call in <i>file</i> at line <i>line</i> . | Severe | Call the stored procedure with arguments passed by name or by position. Mixing these modes in the same call is illegal. |
| M_INDICVAR | Error: item must be an indicator-type variable. | Severe | Use a short integer. |
| M_INTVAR | Error: item must be an integer-type variable. | Severe | Use an integer. |
| M_MISMATCHED_QUOTES | Error: mismatched quotes on hex literal <i>value</i> . | Severe | Make quotes match. |
| M_MULTIDIM_ARRAY | Error: at line <i>line</i> . Multiple-dimensioned array variables are not supported. | Severe | Multiple-dimensioned arrays are not supported. Break up a $m \times n$ array into m arrays of n elements each. |
| M_MULTI_RESULTS | Error: Embedded Query at line <i>line</i> returns multiple result sets. | Severe | Break the query into multiple queries, each returning one result set. Alternatively, rewrite the queries to fill a temporary table with all the values, then select from the temporary table, thus giving a single result set. |
| M_NODCL_NONANSI | Warning: Neither SQLCODE nor SQLCA declared in non-ANSI mode. | Warning | In non-ANSI mode, declare either SQLCA, SQLCODE, or both. Verify that the scope is applicable for all Embedded SQL statements within the program. |
| M_NOLITERAL | Error: item may not be an unquoted name. | Severe | Use a quoted name or host variable. |
| M_NOSQUOTE | Error: item may not be a single quoted string. Use double quotes. | Severe | Use double quotes. |
| M_NOT_AT_ABLE | An “at” clause is used with a statement type which does not allow it. This occurred at line <i>value</i> . | Severe | Remove the at clause from the specified statement. |
| M_NUMBER_OR_INDICVAR | Error: item must be an integer or an indicator-type variable. | Severe | Use a literal integer or a short integer or CS_SMALLINT. |

| Message ID | Message text | Severity | Fix |
|------------------------|---|-------------|--|
| M_NUMBER_OR_INTVAR | Error: item must be an integer constant or an integer type variable. | Severe | Unused. May be used to raise an error if some field in the dynamic SQL statements (such as, MAX, Value <i>n</i> ,) are not an integer type or an integer constant. |
| M_PARAM_RESULTS | Error: Embedded Query at line <i>line</i> returns unexpected parameter result sets. | Severe | Arises only during optional server syntax checking. Determine why the query is returning parameters and rewrite it. |
| M_PASS1_ERR | File <i>file</i> : Syntax errors in Pass 1: Pass 2 not done. | Information | Errors in Pass 1 resulted in an aborted precompilation. Correct Pass 1 errors, then proceed. |
| M_PTR_IN_DEC_SEC | Warning: Pointers are not yet supported in Declare section. | Warning | |
| M_QSTRING_OR_STRINGVAR | Error: item must be a quoted string or a type string variable. | Severe | Verify that server name, user name, and password are either double-quoted strings or of type string. |
| M_SCALAR_CHAR | Error: non-array character variable <i>value</i> is being used illegally as a host variable at line <i>line</i> . | Severe | Use a character array. |
| M_SQLCA_IGNR | Warning: Both SQLCODE and SQLCA declared: SQLCA ignored. | Warning | Remove one of the two declarations. |
| M_SQLCA_WARN | Warning: An INCLUDE SQLCA seen while in ANSI mode: SQLCA ignored. | Warning | |
| M_SQLCODE_UNDECL | Warning: SQLCODE not declared while in ANSI mode. | Warning | Declare SQLCODE. |
| M_STATE_CODE | Warning: both SQLSTATE and SQLCODE declared: SQLCODE ignored. | Warning | Remove one of the two declarations. |
| M_STATE_SQLCA | Warning: both SQLSTATE and SQLCA declared: SQLCA ignored. | Warning | Remove one of the two declarations. |
| M_STATUS_RESULTS | Error: Embedded Query at line <i>line</i> returns unexpected status result sets. | Severe | Arises only during optional server syntax checking. Determine why the query is returning status results and rewrite it. |

| Message ID | Message text | Severity | Fix |
|-------------------|--|----------|--|
| M_STICKY_AUTOVAR | Warning: automatic variable <i>value</i> used with sticky binds at line <i>line</i> . This may cause incorrect results or errors at runtime. | Warning | Be certain that your program logic will not allow errors in this case. Alternatively, use a static or global variable. |
| M_STICKY_REGVAR | Error: register variable <i>value</i> cannot be used with sticky binds at line <i>line</i> . | Severe | Remove the register qualifier. |
| M_STRUCT_NOTFOUND | Structure/union definition not found in scope in <i>file</i> at <i>line</i> . | Severe | Verify that the definition of the structure or union is within the scope of the specified line. |
| M_SYNTAX_PARSE | Syntax error in file <i>file</i> at <i>line</i> . | Severe | Check the indicated line number for a syntax error in the Embedded SQL grammar. |
| M_UNBALANCED_DQ | Unbalanced quotes in delimited identifier. | Severe | Balance the quotes. |
| M_UNDEF_ELM | Error <i>value</i> : illegal structure/union element. | Severe | The specified element of the structure is not included in the structure definition. Correct the definition. |
| M_UNDEF_HV | Host variable <i>value</i> undefined. | Severe | Define the host variable in the proper place. |
| M_UNDEF_IV | Indicator variable <i>value</i> undefined. | Severe | Define the indicator variable in the proper place. |
| M_UNDEF_STR | Error structure <i>value</i> undefined. | Severe | Undefined structure on the specified line. Define the structure in the proper scope. |
| M_UNSUP | The <i>value</i> , feature is not supported in this version. | Fatal | This feature is not supported. |

Table A-3: Second pass parser messages

| Message ID | Message text | Severity | Fix |
|---------------------|---|----------|---|
| M_CURSOR_RD | The cursor <i>value</i> is redefined at line <i>line</i> in <i>file</i> . | Warning | A cursor with same name has already been declared. Use a different name. |
| M_HOSTVAR_MULTIBIND | Warning: host variable used as a bind variable <i>value</i> more than once per statement. | Warning | Do not use a host variable multiple times in a single fetch statement. You cannot fetch multiple results into one location. Client-Library causes the last value fetched to be put in the variable. |

| Message ID | Message text | Severity | Fix |
|------------------|---|----------|---|
| M_INVTYPE_IV | Indicator variable is an incorrect type. | Severe | The indicator variable should be of type CS_SMALLINT or of type INDICATOR. |
| M_PARSE_INTERNAL | Internal parser error at line <i>line</i> . Please contact a Sybase representative. | Fatal | Immediately report this internal consistency parser error to Sybase Technical Support. |
| M_SQLCANF | 'INCLUDE SQLCA' statement not found. | Warning | Add statement. |
| M_WHEN_ERROR | Unable to find the SQL statement 'WHENEVER SQLERROR'. | Warning | Add 'WHENEVER SQLERROR' statement or use command line option to suppress warning and 'INTO' messages (see the <i>Open Client/Server Programmer's Supplement</i>). |
| M_WHEN_NF | Unable to find the SQL statement 'WHENEVER NOT FOUND'. | Warning | Enter a 'WHENEVER NOT FOUND' statement or use command line option to suppress warning and 'INTO' messages (see the <i>Open Client/Server Programmer's Supplement</i>). |
| M_WHEN_WARN | Unable to find the SQL statement 'WHENEVER WARNING'. | Warning | Enter a 'WHENEVER WARNING' statement or use command line option to suppress warning and 'INTO' messages (see the <i>Open Client/Server Programmer's Supplement</i>). |

Table A-4: Code generation messages

| Message ID | Message text | Severity | Fix |
|-------------------|--|----------|---|
| M_INCLUDE_PATHLEN | An included or copied file path was too long. Leaving the path off the generated file name: <i>value</i> . | Warning | Use links or move the file to a shorter path. |
| M_WRITE_ISQL | Unable to write to the isql file. Return code: <i>value</i> . | Fatal | Verify your permission to create and write to the isql file and in the directory. Also, verify that the file system is not full. |
| M_WRITE_TARGET | Unable to write to the target file. Return code: <i>value</i> . | Fatal | Verify your permission to create and write to a file in the directory where the precompiler is generating the target file. Also, verify that the file system is not full. |

Table A-5: FIPS flag messages

| Message ID | Message text | Severity | ANSI extension |
|-----------------|---|-------------|---|
| M_FIPS_ARRAY | FIPS-flagger Warning: ANSI extension ARRAY type at <i>line</i> . | Information | Arrays. As for all FIPS messages, do not use this feature if you need to be ANSI compliant. |
| M_FIPS_DATAINIT | FIPS-flagger Warning: ANSI extension Data Initialization at <i>line</i> . | Information | Data initialization. |
| M_FIPS_HASHDEF | FIPS-flagger Warning: ANSI extension "#DEFINE" <i>line</i> . | Information | Using #define in a declare section. |
| M_FIPS_LABEL | FIPS-flagger Warning: ANSI extension ':' with label in a "WHENEVER" clause. | Information | Allowing ":" with a label in a "WHENEVER" clause. |
| M_FIPS_POINTER | FIPS-flagger Warning: ANSI extension POINTER type at <i>line</i> . | Information | The type POINTER. |
| M_FIPS_SQLDA | FIPS-flagger Warning: ANSI extension sqlda. (line <i>line</i>). | Information | The SQLDA structure. |
| M_FIPS_STMT | FIPS-flagger Warning: ANSI extension statement (line <i>line</i>) | Information | The statement at this line is an extension. |
| M_FIPS_SYBTYPE | FIPS-flagger Warning: ANSI extension Sybase SQL-Type <i>line</i> . | Information | Sybase-specific datatypes. |
| M_FIPS_TYPE | FIPS-flagger Warning: ANSI extension data type at <i>line</i> . | Information | The specified syntax is not ANSI compliant. |
| M_FIPS_TYPEDEF | FIPS-flagger Warning: ANSI extension TYPEDEF <i>line</i> . | Information | TYPEDEF. |
| M_FIPS_VOID | FIPS-flagger Warning: ANSI extension VOID type <i>line</i> . | Information | The type VOID. |

Table A-6: Internal error messages

| Message ID | Message text | Severity | Fix |
|-------------------|--|----------|---|
| M_ALC_MEMORY | Unable to allocate a block of memory. | Fatal | Check system resources. |
| M_FILE_STACK_OVFL | File stack overflow: Max allowed nesting is <i>value</i> . | Fatal | The file stack overflowed while trying to process the nested INCLUDE statement. Do not exceed the nested depth maximum of 32. |

| Message ID | Message text | Severity | Fix |
|------------------|---|----------|--|
| M_INTERNAL_ERROR | Fatal Internal Error at file <i>file</i> line <i>line</i> : Argument inconsistency error. Please contact Sybase representative. | Fatal | This is an internal error. Contact your Sybase representative. |

Table A-7: Sybase and Client-Library messages

| Message ID | Message text | Severity | Fix |
|------------|--|-------------|--|
| M_COLMCNT | The bind count of the <i>bind variable count</i> and the column count of result set are incompatible. | Warning | The number of returned columns is different from the number of results columns returned with the bind variable types and number. |
| M_COLVARLM | The host variable <i>name</i> length <i>value</i> is less than the column length of <i>value</i> . | Warning | The host variable may not be able to hold the fetched column. Check the column length and adjust the length of the host variable accordingly. |
| M_COLVARPS | The host variable <i>name</i> precision and scale: <i>value</i> are different from the column's precision <i>value</i> and scale: <i>value</i> | Warning | The precision and scale of the host variable is different from that of the column being fetched or inserted into. Make the scale and precision compatible. |
| M_COLVARTM | Open Client unable to convert type <i>value</i> to type <i>value</i> for host variable name. | Warning | Illegal type. Use <code>cs_convert</code> , as Open Client will not convert by default. |
| M_CTMSG | Client Library message: <i>value</i> . | | |
| M_OCAPI | Error during execution of the Open Client API <i>value</i> . Error: <i>value</i> . | Warning | Depending on the context in which this warning occurs, you may be required to take corrective action before proceeding. |
| M_OPERSYS | Operating system error: <i>value</i> occurred during execution of the Open Client API. | Warning | An operating system error occurred. Speak with your system administrator. |
| M_PRECLINE | Warning(s) during check of query on line <i>value</i> . | Information | Examine the query for problems. |
| M_SYBSERV | Sybase server error. Server: <i>value</i> . Message: name. | Warning | Check the syntax of the statement sent to the server that caused this error. Verify that all resources are available in the Server to process the SQL statement. |

Table A-8: Runtime messages

| SQLCODE Value, SQLSTATE Code | Message Text | Severity | Fix |
|---|---|-----------------|--|
| -25001 ZZ000 | Unrecoverable error occurred. | Fatal | Immediately report this error to Sybase Technical Support. |
| -25002 ZA000 | Internal error occurred. | Fatal | Immediately report this error to Sybase Technical Support. |
| -25003 ZD000 | Unexpected CS_COMPUTE_RESULT received. | Severe | Embedded SQL cannot retrieve compute results. Rewrite the query so it does not return them. |
| -25004 ZE000 | Unexpected CS_CURSOR_RESULT received. | Severe | Verify that the value returned by the CS_LIBRARY routine is valid. Consult your CS-Library documentations for details. |
| -25005 ZF000 | Unexpected CS_PARAM_RESULT received. | Severe | Verify that the value returned by the CS_LIBRARY routine is valid. Consult your CS-Library documentation for details. |
| -25006 ZG000 | Unexpected CS_ROW_RESULT received. | Severe | Verify that the value returned by the CS_LIBRARY routine is valid. Consult your CS-Library documentation for details. |
| -25007 ZB000 | No message(s) returned for SQLCA, SQLCODE, or SQLSTATE. | Information | Informational message. No action is required. |
| -25008 ZC000 | Connection has not been defined yet. | Severe | Enter a valid connect statement. |
| -25009 ZH000 | Unexpected CS_STATUS_RESULT received. | Severe | Verify that the value returned by the CS_LIBRARY routine is valid. Consult your CS-Library documentation for details. |
| -25010 ZI000 | Unexpected CS_DESCRIBE_RESULT received. | Severe | Verify that the value returned by the CS_LIBRARY routine is valid. Consult your CS-Library documentation for details. |
| -25011 22005 | Data exception—error in assignment of item descriptor type. | Severe | Enter a valid descriptor type. |
| -25012 ZJ000 | Memory allocation failure. | Severe | There is an insufficient amount of memory to allocate to this operation. |

| SQLCODE Value, SQLSTATE Code | Message Text | Severity | Fix |
|---------------------------------|---|----------|---|
| -25013 ZK000 | SQL-Server must be version 10 or greater. | Severe | Verify that your installation has an installed, functioning copy of SQL Server 10.0 or higher. If you do not have SQL Server 10.0 or higher, have your installation's designated person contact Sybase Technical Support. |
| -25014 22024 | Data exception — unterminated C string. | Severe | Be sure to null-terminate all C strings. |
| -25015 ZL000 | Error retrieving thread identification. | Severe | An internal error probably occurred – call Technical Support. |
| -25016 ZM000 | Error initializing Client Library. | Severe | Check your \$SYBASE directory setup. |
| -25017 ZN000 | Error taking a mutex. | Severe | Unused. |
| -25018 08002 | Connection name in use. | Severe | Check your program logic – are you re-opening an open connection? Or use a new name for the second connection. Note You cannot have two 'DEFAULT' connections. |

Type Definitions and Limits

The following type definitions are valid in Embedded SQL:

Table B-1: Valid typedefs

| Typedef | Description |
|--------------|---|
| CS_BINARY | Binary type |
| CS_BIT | Bit type |
| CS_CHAR | Character type |
| CS_DATETIME | Datetime type |
| CS_FLT8 | 8-byte float type |
| SQLINDICATOR | Used for indicator variables (2-byte integer) |
| CS_INT | 4-byte integer |
| CS_MONEY | Money type |
| CS_SMALLINT | 2-byte integer |
| CS_TINYINT | 1-byte integer |
| CS_SMALLINT | Unsigned 2-byte integer |
| CS_TEXT | Text type |
| CS_IMAGE | Image type |

Implementation limits

The nesting depth for `exec sql include filename` limit is 32.

Embedded SQL Constructs

The following constructs are valid in Embedded SQL statements:

Table C-1: Embedded SQL constructs

| | |
|--|--|
| begin declare section | dump database |
| begin tran | dump tran |
| begin work | end declare section |
| checkpoint | <i>exec procedure_name</i> |
| close <i>cursor_name</i> | <i>execute name</i> |
| commit tran | execute immediate |
| commit work | fetch <i>cursor_name</i> |
| connect | grant |
| create database | include sqlca or include <i>filename</i> |
| create default | insert |
| create table | open <i>cursor_name</i> |
| create index | prepare <i>statement_name</i> |
| create unique index | revoke |
| create clustered index | rollback tran |
| create nonclustered index | rollback work |
| create unique clustered index | select |
| create unique nonclustered index | set |
| create proc | truncate |
| create rule | update |
| create trigger | use |
| create view | <i>whenever condition action</i> |
| declare cursor | |
| delete | |
| disconnect | |
| drop table default index proc rule trigger view | |

Glossary

Adaptive Server Enterprise

A server in Sybase's client/server architecture. Adaptive Server manages multiple databases and multiple users, keeps track of the actual location of data on disks, maintains mapping of logical data description to physical data storage, and maintains data and procedure caches in memory.

Note Prior to version 11.5, Adaptive Server Enterprise was known as SQL Server.

array

A structure composed of multiple identical variables that can be individually addressed.

array binding

The process of binding a result column to an array variable. At fetch time, multiple rows' worth of the column are copied into the variable.

batch

A group of commands or statements.

A Client-Library command batch is one or more Client-Library commands terminated by an application's call to `ct_send`. For example, an application can batch together commands to declare, set rows for, and open a cursor.

A Transact-SQL statement batch is one or more Transact-SQL statements submitted to Adaptive Server by means of a single Client-Library command or Embedded SQL statement.

browse mode

Browse mode is a method that DB-Library and Client-Library applications can use to browse through database rows, updating their values one row at a time. Cursors provide similar functionality and are generally more portable and flexible.

bulk copy

A utility for copying data in and out of databases. Also called `bcp`.

callback event

In Open Client and Open Server, a callback event is an occurrence that triggers a callback routine.

callback routine

A routine that Open Client or Open Server calls in response to a triggering event, known as a callback event.

| | |
|---------------------------------|---|
| capabilities | A client/server connection's capabilities determine the types of client requests and server responses permitted for that connection. |
| character set | A set of specific (usually standardized) characters with an encoding scheme that uniquely defines each character. ASCII and ISO 8859-1 (Latin 1) are two common character sets. |
| character set conversion | Changing the encoding scheme of a set of characters on the way into or out of a server. Conversion is used when a server and a client communicating with it use different character sets. For example, if Adaptive Server uses ISO 8859-1 and a client uses Code Page 850, character set conversion must be turned on so that both server and client interpret the data passing back and forth in the same way. |
| client | In client/server systems, the client is the part of the system that sends requests to servers and processes the results of those requests. |
| Client-Library | Part of Open Client, Client-Library is a collection of routines for use in writing client applications. Client-Library accommodates cursors and other advanced features in the Sybase product line. |
| code set | See <i>character set</i> . |
| collating sequence | See <i>sort order</i> . |
| command | In Client-Library, a command is a server request initiated by an application's call to <code>ct_command</code> , <code>ct_dynamic</code> , or <code>ct_cursor</code> and terminated by the application's call to <code>ct_send</code> . |
| command structure | A command structure (<code>CS_COMMAND</code>) is a hidden Client-Library structure that Client-Library applications use to send commands and process results. |
| connection structure | A connection structure (<code>CS_CONNECTION</code>) is a hidden Client-Library structure that defines a client/server connection within a context. |
| context structure | A context structure (<code>CS_CONTEXT</code>) is a CS-Library hidden structure that defines an application "context," or operating environment, within a Client-Library or Open Server application. The CS-Library routines <code>cs_ctx_alloc</code> and <code>cs_ctx_drop</code> allocate and drop a context structure, respectively. |
| conversion | See <i>character set conversion</i> . |
| CS-Library | Included with both the Open Client and Open Server products, CS-Library is a collection of utility routines that are useful to both Client-Library and Server-Library applications. |

| | |
|-------------------------|---|
| current row | With respect to cursors, the current row is the row to which a cursor points. A fetch against a cursor retrieves the current row. |
| cursor | <p>A cursor is a symbolic name that is associated with a SQL statement.</p> <p>In Embedded SQL, a cursor is a data selector that passes multiple rows of data to the host program, one row at a time.</p> |
| database | A set of related data tables and other database objects that are organized to serve a specific purpose. |
| datatype | A defining attribute that describes the values and operations that are legal for a variable. |
| DB-Library | Part of Open Client, DB-Library is a collection of routines for use in writing client applications. |
| deadlock | A situation that arises when two users, each having a lock on one piece of data, attempt to acquire a lock on the other's piece of data. Adaptive Server detects deadlocks and resolves them by killing one user's process. |
| default | Describes the value, option, or behavior that Open Client/Server products use when none is explicitly specified. |
| default database | The database that a user gets by default when he or she logs in to a database server. |
| default language | <ol style="list-style-type: none">1. The language that Open Client/Server products use when an application does no explicit localization. The default language is determined by the "default" entry in the locales file.2. The language that Adaptive Server uses for messages and prompts when a user has not explicitly chosen a language. |
| Dynamic SQL | Dynamic SQL allows an Embedded SQL or Client-Library application to execute SQL statements containing variables whose values are determined at runtime. |
| error message | A message that an Open Client/Server product issues when it detects an error condition. |
| event | An occurrence that prompts an Open Server application to take certain actions. Client commands and certain commands within Open Server application code can trigger events. When an event occurs, Open Server calls either the appropriate event-handling routine in the application code or the appropriate default event handler. |

| | |
|-----------------------------|---|
| event handler | In Open Server, a routine that processes an event. An Open Server application can use the default handlers Open Server provides or can install custom event handlers. |
| exposed structure | An exposed structure is a structure whose internals are exposed to Open Client/Server programmers. Open Client/Server programmers can declare, manipulate, and de-allocate exposed structures directly. The CS_DATAFMT structure is an example of an exposed structure. |
| extended transaction | In Embedded SQL, an extended transaction is a transaction composed of multiple Embedded SQL statements. |
| FIPS | FIPS is an acronym for Federal Information Processing Standards. If FIPS flagging is enabled, Adaptive Server or the Embedded SQL precompiler issue warnings when a non-standard extension to a SQL statement is encountered. |
| gateway | A gateway is an application that acts as an intermediary for clients and servers that cannot communicate directly. Acting as both client and server, a gateway application passes requests from a client to a server and returns results from the server to the client. |
| hidden structure | A hidden structure is a structure whose internals are hidden from Open Client/Server programmers. Open Client/Server programmers must use Open Client/Server routines to allocate, manipulate, and de-allocate hidden structures. The CS_CONTEXT structure is an example of a hidden structure. |
| host language | The programming language in which an application is written. |
| host program | In Embedded SQL, the host program is the application program that contains the Embedded SQL code. |
| host variable | In Embedded SQL, a variable that enables data transfer between Adaptive Server and the application program. See also <i>indicator variable</i> , <i>input variable</i> , <i>output variable</i> , <i>result variable</i> , and <i>status variable</i> . |
| indicator variable | <p>A variable whose value indicates special conditions about another variable's value or about fetched data.</p> <p>When used with an Embedded SQL host variable, an indicator variable indicates when a database value is null.</p> |
| input variable | A variable that is used to pass information to a routine, a stored procedure, or Adaptive Server. |

| | |
|-------------------------|--|
| interfaces file | A file that maps server names to transport addresses. When a client application calls <code>ct_connect</code> or <code>dbopen</code> to connect to a server, Client-Library or DB-Library searches the interfaces file for the server's address. Note that not all platforms use the interfaces file. On these platforms, an alternate mechanism directs clients to server addresses. |
| isql script file | In Embedded SQL, an isql script file is one of the three files the precompiler can generate. An isql script file contains precompiler-generated stored procedures, which are written in Transact-SQL. |
| key | A subset of row data that uniquely identifies a row. Key data uniquely describes the <i>current row</i> in an open cursor. |
| keyword | A word or phrase that is reserved for exclusive use in Transact-SQL or Embedded SQL. Also called a <i>reserved word</i> . |
| listing file | In Embedded SQL, a listing file is one of the three files the precompiler can generate. A listing file contains the input file's source statements and informational, warning, and error messages. |
| locale name | A character string that represents a language/character set pair. Locale names are listed in the <i>locales file</i> . Sybase predefines some locale names, but a system administrator can define additional locale names and add them to the locales file. |
| locale structure | A locale structure (<code>CS_LOCALE</code>) is a CS-Library hidden structure that defines custom localization values for a Client-Library or Open Server application. An application can use a <code>CS_LOCALE</code> to define the language, character set, datepart ordering, and sort order it will use. The CS-Library routines <code>cs_loc_alloc</code> and <code>cs_loc_drop</code> allocate and drop a locale structure. |
| locales file | A file that maps locale names to language/character set pairs. Open Client/Server products search the locales file when loading localization information. |
| localization | Localization is the process of setting up an application to run in a particular national language environment. An application that is localized typically generates messages in a local language and character set and uses local datetime formats. |
| login name | The name a user uses to log in to a server. An Adaptive Server login name is valid if Adaptive Server has an entry for that user in the system table <code>syslogins</code> . |
| message number | A number that uniquely identifies an error message. |

| | |
|---------------------------------|---|
| message queue | In Open Server, a linked list of message pointers through which threads communicate. Threads can write messages into and read messages from the queue. |
| multi-byte character set | A character set that includes characters encoded using more than 1 byte. EUC JIS and Shift-JIS are examples of multibyte character sets. |
| mutex | A mutual exclusion semaphore. This is a logical object that an Open Server application uses to ensure exclusive access to a shared object. |
| null | Having no explicitly assigned value. NULL is not equivalent to zero, or to blank. A value of NULL is not considered to be greater than, less than, or equivalent to any other value, including another value of NULL. |
| Open Server | A Sybase product that provides tools and interfaces for creating custom servers. |
| Open Server application | A custom server constructed with Open Server. |
| output variable | In Embedded SQL, a variable that passes data from a stored procedure to an application program. |
| parameter | <ol style="list-style-type: none">1. A variable that is used to pass data to and retrieve data from a routine.2. An argument to a stored procedure. |
| passthrough mode | When in passthrough mode, a gateway relays Tabular Data Stream (TDS) packets between a client and a remote data source without unpacking the packets' contents. |
| property | A property is a named value stored in a structure. Context, connection, thread, and command structures have properties. A structure's properties determine how it behaves. |
| query | <ol style="list-style-type: none">1. A data retrieval request; usually a <code>select</code> statement.2. Any SQL statement that manipulates data. |
| registered procedure | In Open Server, a collection of C statements stored under a name. Open Server-supplied registered procedures are called <i>system registered procedures</i> . |
| remote procedure call | <ol style="list-style-type: none">1. One of two ways in which a client application can execute an Adaptive Server stored procedure. (The other is with a <code>Transact-SQL execute</code> statement.) A Client-Library application initiates a remote procedure call command by calling <code>ct_command</code> . A DB-Library application initiates a remote procedure call command by calling <code>dbRPCinit</code> . |

| | |
|-------------------------|---|
| | <ol style="list-style-type: none">2. A type of request a client can make of an Open Server application. In response, Open Server either executes the corresponding registered procedure or calls the Open Server application's RPC event handler.3. A stored procedure executed on a different server from the server to which the user is connected. |
| result variable | In Embedded SQL, a variable which receives the results of a select or fetch statement. |
| server | In client/server systems, the server is the part of the system that processes client requests and returns results to clients. |
| Server-Library | A collection of routines for use in writing Open Server applications. |
| sort order | Used to determine the order in which character data is sorted. Also called collating sequence. |
| SQLCA | <ol style="list-style-type: none">1. In an Embedded SQL application, SQLCA is a structure that provides a communication path between Adaptive Server and the application program. After executing each SQL statement, Adaptive Server stores return codes in SQLCA.2. In a Client-Library application, SQLCA is a structure that the application can use to retrieve Client-Library and server error and informational messages. |
| SQLCODE | <ol style="list-style-type: none">1. In an Embedded SQL application, SQLCODE is a structure that provides a communication path between Adaptive Server and the application program. After executing each SQL statement, Adaptive Server stores return codes in SQLCODE. A SQLCODE can exist independently or as a variable within a SQLCA structure.2. In a Client-Library application, SQLCODE is a structure that the application can use to retrieve Client-Library and server error and informational message codes. |
| SQL Server | see Adaptive Server Enterprise. |
| statement | In Transact-SQL or Embedded SQL, an instruction that begins with a keyword. The keyword names the basic operation or command to be performed. |
| status variable | In Embedded SQL, a variable that receives the return status value of a stored procedure, thereby indicating the procedure's success or failure. |
| stored procedure | In Adaptive Server, a collection of SQL statements and optional control-of-flow statements stored under a name. Adaptive Server-supplied stored procedures are called <i>system procedures</i> . |

| | |
|-------------------------------------|---|
| System Administrator | The user in charge of Adaptive Server system administration, including creating user accounts, assigning permissions, and creating new databases. On Adaptive Server, the System Administrator's login name is "sa". |
| system descriptor | In Embedded SQL, a system descriptor is an area of memory that holds a description of variables used in Dynamic SQL statements. |
| system procedures | Stored procedures that Adaptive Server supplies for use in system administration. These procedures are provided as shortcuts for retrieving information from system tables, or as mechanisms for accomplishing database administration and other tasks that involve updating system tables. |
| system registered procedures | Internal registered procedures that Open Server supplies for registered procedure notification and status monitoring. |
| target file | In Embedded SQL, a target file is one of three files the precompiler can generate. A target file is similar to the original input file, except that all SQL statements are converted to Client-Library function calls. |
| TDS | (Tabular Data Stream) An application-level protocol that Sybase clients and servers use to communicate. It describes commands and results. |
| thread | A path of execution through Open Server application and library code and the path's associated stack space, state information, and event handlers. |
| Transact-SQL | Transact-SQL is an enhanced version of the database language SQL. Applications can use Transact-SQL to communicate with Sybase Adaptive Server. |
| transaction | One or more server commands that are treated as a single unit for the purposes of backup and recovery. Commands within a transaction are committed as a group; that is, either all of them are committed or all of them are rolled back. |
| transaction mode | Transaction mode refers to the manner in which Adaptive Server manages transactions. Adaptive Server supports two transaction modes: Transact-SQL mode (also called "unchained transactions") and ANSI mode (also called "chained transactions"). |
| user name | See <i>login name</i> . |

Index

Symbols

58
#define 25
\$ 58
? (question mark)
 dynamic parameter marker 66
_sql 13

A

allocate descriptor 112
allow ddl in tran 116
ANSI
 and dynamic SQL 65
array
 double-dimensional 25, 114
arrays 49
 multiple 34
 persistent binding 106
 select into 49
 using 34
arrays, batch 50
at 40
at connect_name
 named connection 120
at connection_name 43
at connection_name clause
 in exec sql statement 142
automatic variables 108

B

-b precompiler option 98
batch arrays
 fetch into 50
batches
 and get diagnostics 61

 restrictions 13
 statements 13
begin transaction 62, 63
binding 58, 65
 and loops 95
 persistent 93, 109
 variables 94

C

call 89
case sensitivity
 Embedded SQL 11
character array
 declaring 25
Client-Library
 run-time 2
close 116
close and cursors 55
close cursor 55
colons
 and indicator variables 31
 and variables 28
command line options, precompiler 6
command structure
 persistent 97
comments
 in Embedded SQL 11
commit 45
commit transaction 63, 118
commit work 63
compatibility 7, 48
 backward 5
complex definition 25
compute clause, disallowed 165
configuration file 175
connect 39
 and multiple connections 41
connection

- naming 42
- connection_name 40
- connections
 - closing 120, 137
 - default 120
 - multiple 41
 - named 120
- constructs
 - valid 201
- continue 89
- conventions
 - variable 33
- conversion, datatype 4, 36
- COPY files 155, 156
- ct_bind routine 93
- ct_fetch routine 93
- current row 51, 54
- cursor names
 - and scoping rules 13
- cursor position 54
- cursors 51, 55, 125, 126, 128, 158, 159
 - and scoping 51
 - closing 55, 116
 - declaring 52
 - deleting current row 55
 - deleting rows 131
 - dynamic 124, 161
 - example 56
 - opening 53
 - persistent binding 100
 - retrieving data 53, 54
 - updating current row 55
 - updating rows 167

D

- data Definition Language(DDL) 67
- data Manipulation Language (DML) 48, 67
- database
 - pubs2 4
- databases
 - accessing 39
 - selecting rows from 164
- datatype conversions 4
 - input variables 37

- result variables 36, 37
- datatypes 35
 - and declaring variables 35
 - C and SQL 35
 - converting 36
 - list of 36
 - list of equivalent 36
- DDL
 - Data Definition Language 67
- ddl in tran 116
- deallocate descriptor 123
- deallocate prepare 124
- deallocated cursors
 - persistent binding 101
- declare cursor 51, 52, 125, 126, 128
 - and stored procedures 61
 - dynamic 125
 - persistent binding 100
 - static 126
 - stored procedure 127
- declare section 21, 22
- default server
 - connecting to 40
- default transaction mode 62
- delete 55
 - searched 131
 - with cursors 54
- delete positioned cursor 129
- delete where current of 72
- describe input 133
- describe output 135
- directories
 - and searches 16
- disconnect 44, 137
- DML
 - Data Manipulation Language 67
- DML (Data Manipulation Language) 48
- documentation
 - online 57
- double-dimensional array 25
- DSQUERY environment variable 120
- dynamic cursors
 - persistent binding 100
- dynamic parameter marker 66
- dynamic parameter markers 69, 145, 159, 162
- dynamic SQL 2, 65, 124, 146, 147, 161

- and stored procedures 66
- method 1 67, 68
- method 2 69, 71
- method 3 72, 74
- method 4 74, 79
- overview 65
- prepare and execute 145, 162
- prepare and fetch 162
- stored procedures 66

dynamic SQL protocol 66

dynamic SQL statement 66

E

efficiency 93

embedded SQL ix, 1, 2

- advantages 2
- definition 1
- rules 11
- sample program 10

embedded SQL program

- creating 5

embedded SQL statements

- syntax-checking 91

environment variables 120

SYBASE 120

error

- testing 4

error handler

- writing 90

error_hndl 90

error-handling

- and warning-handling routines 90
- routines 90

errors

- failure to detect 91
- precompiler-detected 91
- SQLSTATE 20
- testing for 86
- trapping 87, 89

error-testing 86

exec 139

exec sql 142

exec statements

- binding 98

execute 144

execute immediate 67, 146

- dynamic SQL 84

extended transaction 63

external 35

external configuration file 175

F

features and enhancements 2, 3

- compatibility 7

fetch 53, 54, 147

fetch into 34

fetch within a loop 54

files

- directory 16
- isql 60
- listing 88
- multiple 6
- precompiler-generated 6

FIPS flagger 3

G

get descriptor 150

get diagnostics 61, 152

- and batches 61
- using 89

go to 89

H

handlers

- error and warning 90

host input variables 29

host output variables 30

host status variables 30

host variable

- and datatypes 37
- character string 34

host variables 2, 31

- assigning data to 53
- declaring 21

- in fetch 54
- naming 33
- persistent binding 94
- scope 108
- using 27
- with indicator variables, using 31

I

- identifiers
 - Embedded SQL 12
- implementation Limit 199
- include 16, 155, 156
 - filename 153
- include file directory 16
- include sqlca 155, 156
- indicator 49
- indicator arrays 49
- indicator variable with host variables
 - example 31
- indicator variables
 - and colons 31
 - declaring 21
 - using 31, 34
 - with input variables 32
 - with output and result variables 32
- input variables 29
 - converting datatypes for 37
 - host 29
- insert statements
 - and binding 95
- interactive SQL 60
- interfaces file 120
- into 48, 59
- invalid statements
 - print 48
 - readtext 48
 - writetext 48
- ISO
 - and dynamic SQL 65
- isql 6
- isql file 6, 60

K

- keywords
 - and variable names 33
 - Embedded SQL 12

L

- label
 - variable 39
- labels 171
- library
 - Client-Library 7
- listing file 6, 88
- localization 2
- logical names 120

M

- markers, dynamic parameter 145, 159, 162
- multiple arrays 34
- multiple connections 41
- multiple source files 6, 7
- multiple SQLCAs 16

N

- named connections 120
- naming conventions
 - variables 13
- nesting
 - stored procedure 60
- null
 - input value 32
- null password
 - specifying 120

O

- ocs.cfg file 175
- online sample programs 57
- open 53, 158

- dynamic cursor 158
- static cursor 159
- open cursor statement
 - persistent binding 100
- output 59
- output file 60
- output variables 30

P

- p precompiler option 97
- parse 6, 91
- password 40
 - null 120
- performance
 - persistent binding 93
- persistent binding 109
 - commands that cannot use 99
 - cursors 100
 - guidelines 105
 - non-cursor statements 99
 - programs that benefit 96
 - scope 97
 - subscripted arrays 106
- persistent command structure 97
- placement
 - Embedded SQL statements 11
- precompiler
 - and dynamic SQL statements 84
 - diagnostics 91
 - functionality 6, 7
- precompiler command line options 6
- precompiler options
 - binding 96, 97
- precompiler-detected errors 90
- prepare 161
- prepare and execute 69, 70, 145
 - dynamic SQL 84
- prepare and fetch
 - dynamic SQL 84
- prepare and fetch with System Descriptor
 - dynamic SQL 84
- procedure_name 58
- product family x
- program

- creating 5
- pubs2 database 4

Q

- question mark
 - and dynamic parameter marker 66
- quotation marks
 - in Embedded SQL 12

R

- rebind/norebind clause 106
- related documents x
- reserved words
 - and variable names 33
 - Embedded SQL 12
- result variables 29
 - converting datatypes for 37
 - host 29
- return code 15, 18
 - SQLCODE 19
 - testing 4
- rollback
 - and SQL Server triggers 63
 - in a trigger 62
- rollback transaction 163
- rollback work 63
- routines
 - error- and warning-handling 90
- rows
 - current 54
 - deleting 129
- rules
 - Embedded SQL 11
- run-time library, Client-Library 2

S

- sample programs
 - online 57
- scope
 - host variables 108

- p and -b precompiler options 98
- scoping 13, 16
 - and cursors 51
 - cursor 51
 - rules 13, 35
 - SQLCA, SQLCODE, and SQLSTATE 15
- select 13, 164
 - with cursors 125, 126, 128, 147
 - returning multiple rows 50, 54
 - returning single rows 49
 - syntax 49
- select clause 61
- select statements
 - binding 98
- server 40
 - connecting to 39
- set connection 41, 165
- set descriptor 166
- source files
 - multiple 6, 7
- SQL descriptors
 - and persistent binding 99
- SQL Server
 - connecting to 39
 - multiple connections to 41, 42
- SQL2 standard
 - and dynamic SQL 65
- SQLCA
 - declaring 16
 - multiple 16
 - table 17
- sqlca 18
- SQLCA variables 16, 17
 - accessing 17
 - setting 15
- sqlca variables
 - accessing 17
 - list of 17
 - SQL Server-related 17
- SQLCODE
 - and fetch 149
 - and multiple row selects 48
 - as a stand-alone 18
 - values 19
 - values table 19
 - within SQLCA 18
- sqlcode 86, 87
 - in error-testing 86
 - return values 86
- SQLCODE values
 - table 19
- SQLCODE variables
 - setting 15
- SQLDAs
 - and persistent binding 99
- SQLSTATE
 - Codes and error messages 20
 - using 20
- SQLSTATE variables
 - setting 15
- sqlwarn 86
 - flags 86
- stack variables 108
- statement batches 13
- statement labels
 - with whenever 171
- statements
 - dynamic SQL 77, 84
 - Embedded SQL 11
 - list of 201
- static cursors
 - persistent binding 100
- status variables 29, 30
 - host 29
- status_variable 58
- stop 89
- stored procedures 2, 3, 7, 47, 58
 - and declare cursor 61
 - and dynamic SQL 66
 - and parameters 58
 - and return status variables 58
 - definition 47
 - dynamic SQL 66
 - executing 58
 - types of 58
- subscripted arrays
 - persistent binding 106
- SYBASE environment variable 120
- syntax checking
 - of Embedded SQL statements 91
- syntax conventions
 - document xii

system variables 17, 18, 21

T

tables

deleting rows from 129

target file 6

testing conditions

whenever 88

transaction

extended 63

transaction Mode

ANSI 63

transaction mode

default 62

Transact-SQL 62

transactions 62, 118

ANSI 62

ISO 62

restricted statements 63

rolling back 163

Transact-SQL

invalid keywords in Embedded SQL 4, 48

keywords in Embedded SQL 12

support 3

using with Embedded SQL 47

Transact-SQL statements 129, 139, 164, 167

triggers 3, 62

typedefs 24, 199

U

update 55, 167

protocol 55

with cursors 54

user 39

V

value

stored procedures 58

variables 21

and datatypes 35, 37

assigning data to 54

declaring 21, 22, 35

host 4, 32

host input 29

host result 29

host status 30

indicator 21

input 21, 29

naming conventions 12, 33

precompiler 13

status 30

system 17, 21

using 28

variables in declare section

example 22

W

warning- and error-handling routines 85, 90

warning handler

writing 90

warning_hndl 90

warning-handling routines 90

warnings

testing for 86, 87

whenever 86, 87, 88

and scoping rules 13

canceling 171

scope of 171

testing conditions 87

whenever...continue 88

where current of 131, 149

